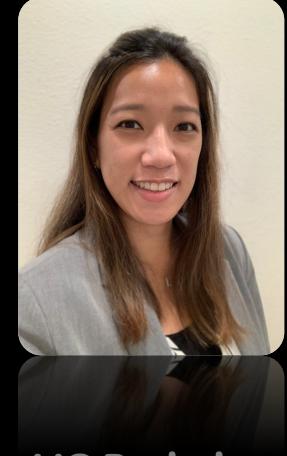


UC Berkeley
Teaching Professor
Dan Garcia

CS61C

Great Ideas in Computer Architecture (a.k.a. Machine Structures)



UC Berkeley
Teaching Professor
Lisa Yan

RISC-V Instruction Formats, Part I

Computing in the News

MIT researchers uncover 'unpatchable' flaw in Apple M1 chips

Carly Page @carlypage_ / 4:10 AM PDT • June 10, 2022

Comment



The vulnerability lies in a hardware-level security mechanism utilized in Apple M1 chips called *pointer authentication codes, or PAC*. This feature makes it much harder for an attacker to inject malicious code into a device's memory ...

Researchers from MIT's [CSAIL], however, have created a novel hardware attack, which combines memory corruption and *speculative execution* attacks to sidestep the security feature. The attack shows that pointer authentication can be defeated without leaving a trace, and *as it utilizes a hardware mechanism, no software patch can fix it*.

<https://techcrunch.com/2022/06/10/apple-m1-unpatchable-flaw/>

11-RISC-V Instruction Formats, Part I (2)

Garcia, Yan



Machine Language

- Machine Language
- R-Format Layout
- I-Format Layout
- I-Format: Load
- S-Format Layout (Store)

Great Idea #1: Abstraction (Levels of Representation/Interpretation)



High Level Language
Program (e.g., C)

```
temp = v[k];
v[k] = v[k+1];
v[k+1] = temp;
```

| *Compiler*

Assembly Language
Program (e.g., RISC-V)

lw	x3, 0(x10)
lw	x4, 4(x10)
sw	x4, 0(x10)
sw	x3, 4(x10)

Anything can be represented
as a number,
i.e., data or instructions

| *Assembler*

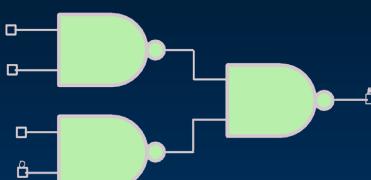
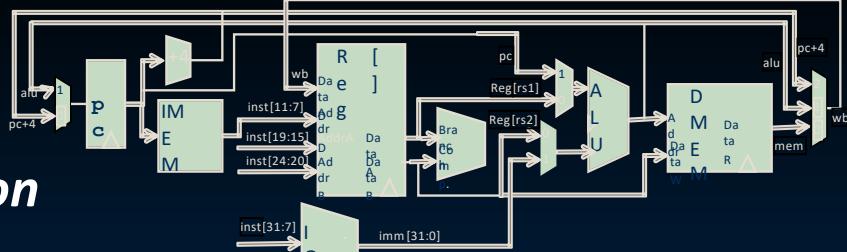
Machine Language
Program (RISC-V)

1000 1101 1110 0010 0000 0000 0000 0000 0000
1000 1110 0001 0000 0000 0000 0000 0000 0100
1010 1110 0001 0010 0000 0000 0000 0000 0000
1010 1101 1110 0010 0000 0000 0000 0000 0100

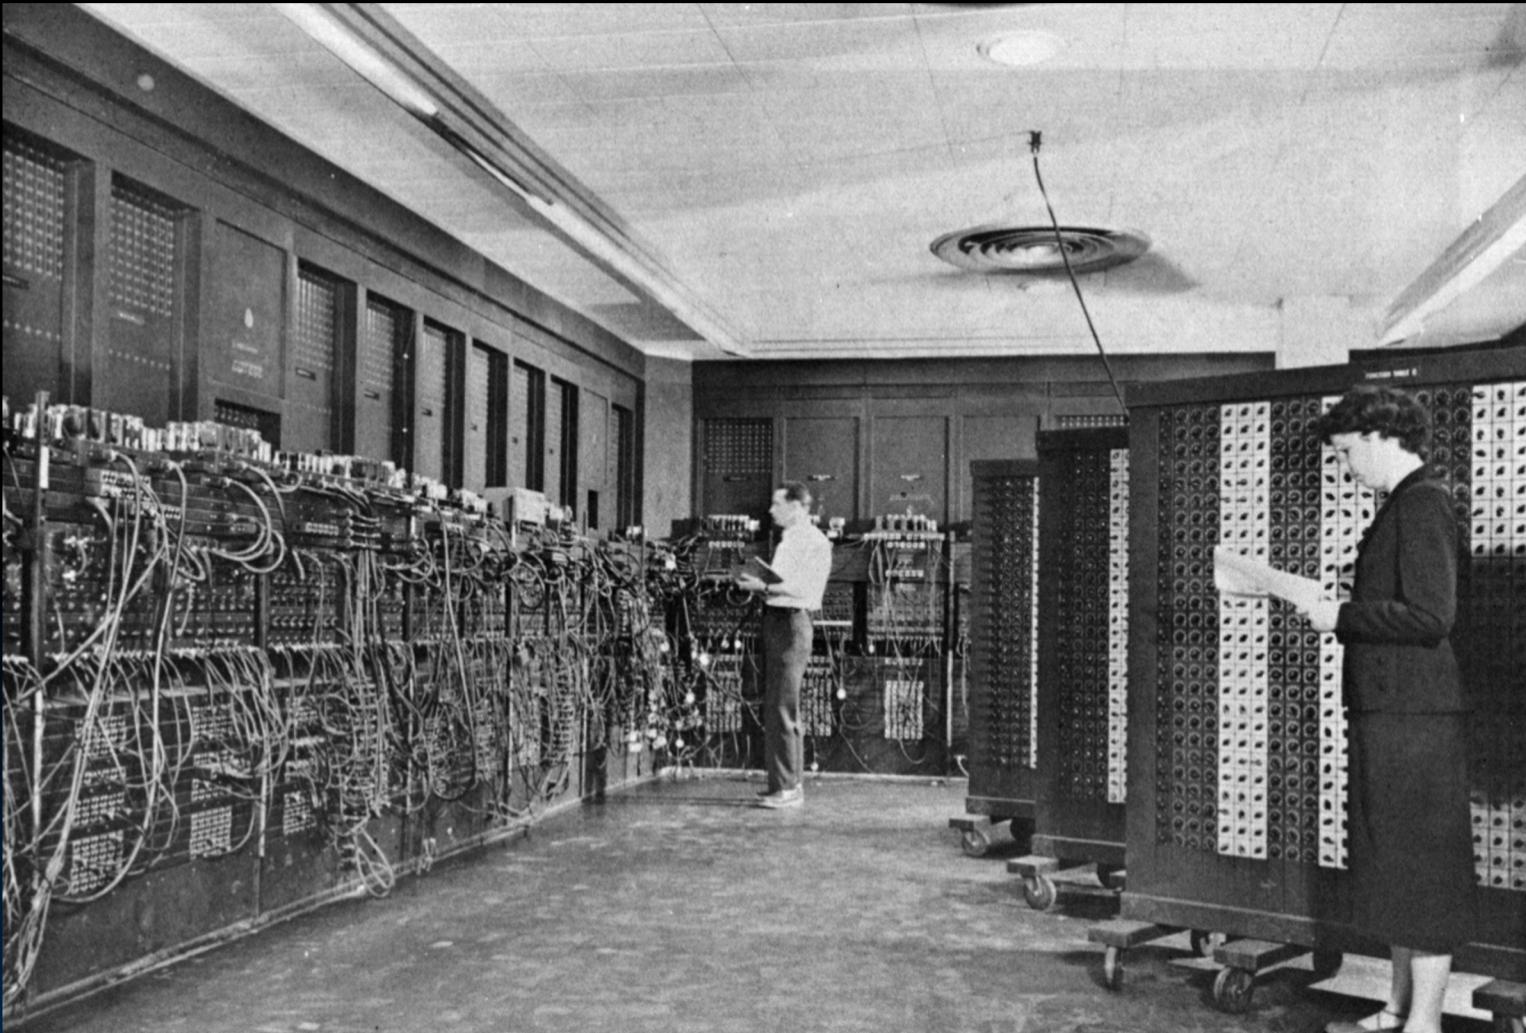
Hardware Architecture Description
(e.g., block diagrams)

| *Architecture Implementation*

Logic Circuit Description
(Circuit Schematic Diagrams)



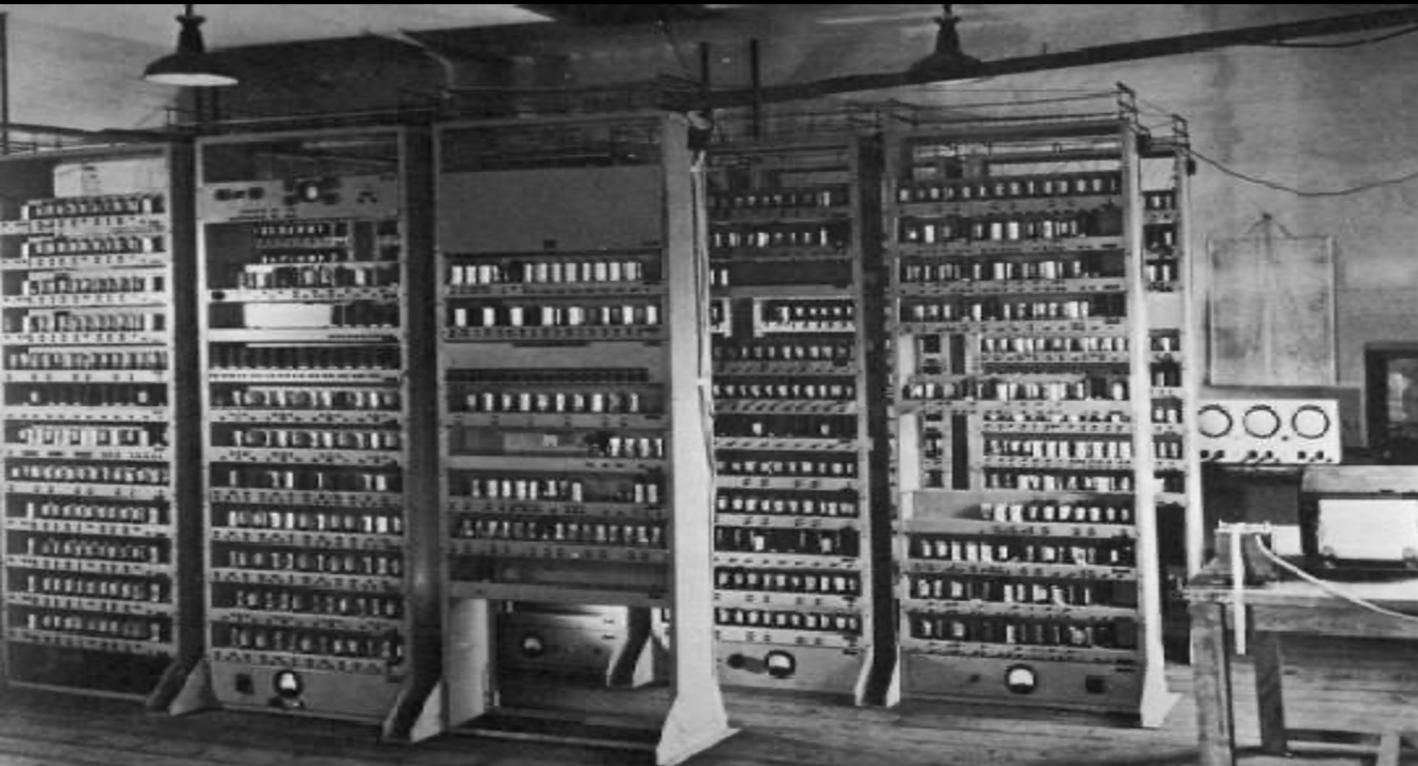
Quick CS History: ENIAC (UPenn, 1946)



Betty Snyder, one of the six original ENIAC "computers," i.e., programmers
11-RISC-V Instruction Formats, Part I (5)

- First Electronic General-Purpose Computer
- Blazingly fast!
 - Multiply 10×10 decimal digits in 2.8 ms
- ...but needed 2-3 days to set up new program.
 - Programming meant patch cords and switches

Quick CS History: EDSAC (Cambridge, 1949)



- First General Stored-Program Electronic Computer
- 17-bit words (35-bit double word), binary 2's complement
- Memory: 512 words
- Multiplication: 6ms
- Subroutines

Quick CS History: IBM 701 (1952)

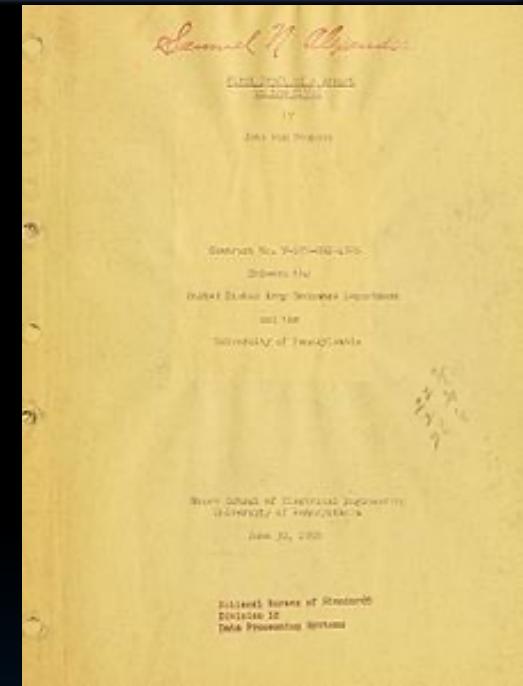


- IBM's first commercial scientific computer
- 36-bit sign-magnitude words
- Memory: 2048 words with vacuum tubes
- Multiplication: <0.5 ms
- Two programmer-accessible registers: Accumulator, Multiplier/Quotient

https://www.ibm.com/ibm/history/exhibits/701/701_intro.html

History's Big Idea: Stored Program Computer

- Instructions are represented as bit patterns (i.e., numbers).
- Therefore, entire programs can be stored to memory to be read or written, just like data.
 - This means reprogramming can happen quickly (seconds); don't have to rewire computer (days).
- Stored-Program Computer AKA the “von Neumann” computer is named after the widely distributed tech report
 - Based on discussions of EDVAC (1944-): Eckert, Mauchly
 - Anticipated earlier by Turing and Zuse



First Draft of a Report on the EDVAC
By John von Neumann
Contract No. W-670-ORD-4926
Between the
United States Army Ordnance Department
and the
University of Pennsylvania
Moore School of Electrical Engineering
University of Pennsylvania
June 30, 1945



Stored Program Computer: Implications

1. Everything (instructions, data words) has a memory address.

- Both branches and jumps use instruction memory addresses. 指令内存地址
- C pointers are just memory addresses and can point to anything.
 - Unconstrained use of addresses lead to nasty bugs! Up to C programmer to avoid errors
- The Program Counter (PC) register keeps address of instruction to execute.
 - PC is effectively a pointer to memory; Intel calls it Instruction Pointer (PC).

2. Programs are distributed in *binary form*, i.e., as assembled machine code.

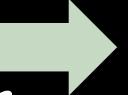
- Programs are each bound to a specific instruction set.
 - Different instruction sets, e.g., for phone vs. PCs.
- Many instruction sets are backwards-compatible and evolve over time.
- E.g., latest PCs with x86 ISA today can still run programs from Intel 8088 (1981)!

Program Counter
程序计数器

Instructions Are 32 Bits Wide

32 bits wide

- ISA defines instructions for CPU down to the bit level:

`add x18, x19, x10`  *assembly code*  *machine code* 

00000000101010011000100100110011

- Remember: Computer only understands 1s and 0s.
Text like “`add x18, x19, x10`” is meaningless to hardware.
- Most data we work with is in words (32-bit chunks).  *4 bytes*
 - 32b registers; lw, sw access memory one word at a time.
- Similarly, RISC-V instructions are fixed size, 32-bit words.
 - RISC-V seeks simplicity: since data is in words, make instructions be fixed-size 32-bit words also.
 - Same 32-bit instructions used for RV32, RV64, RV128

- RISC-V's 32b instruction words are divided into fields.
 - Each field tells processor something about the instruction.
- The RISC-V ISA defines six basic types of instruction formats.
 - RISC-V seeks simplicity: Avoid defining different fields for each instruction; instead, use same format for similar instructions.

R-Format	Register-register arithmetic operations
I-Format	Register-immediate arithmetic operations; Loads
S-Format	Stores
B-Format	Branches (minor variant of S-format)
U-Format	20-bit upper immediate instructions
J-Format	Jumps (minor variant of U-format)

六種基本類型

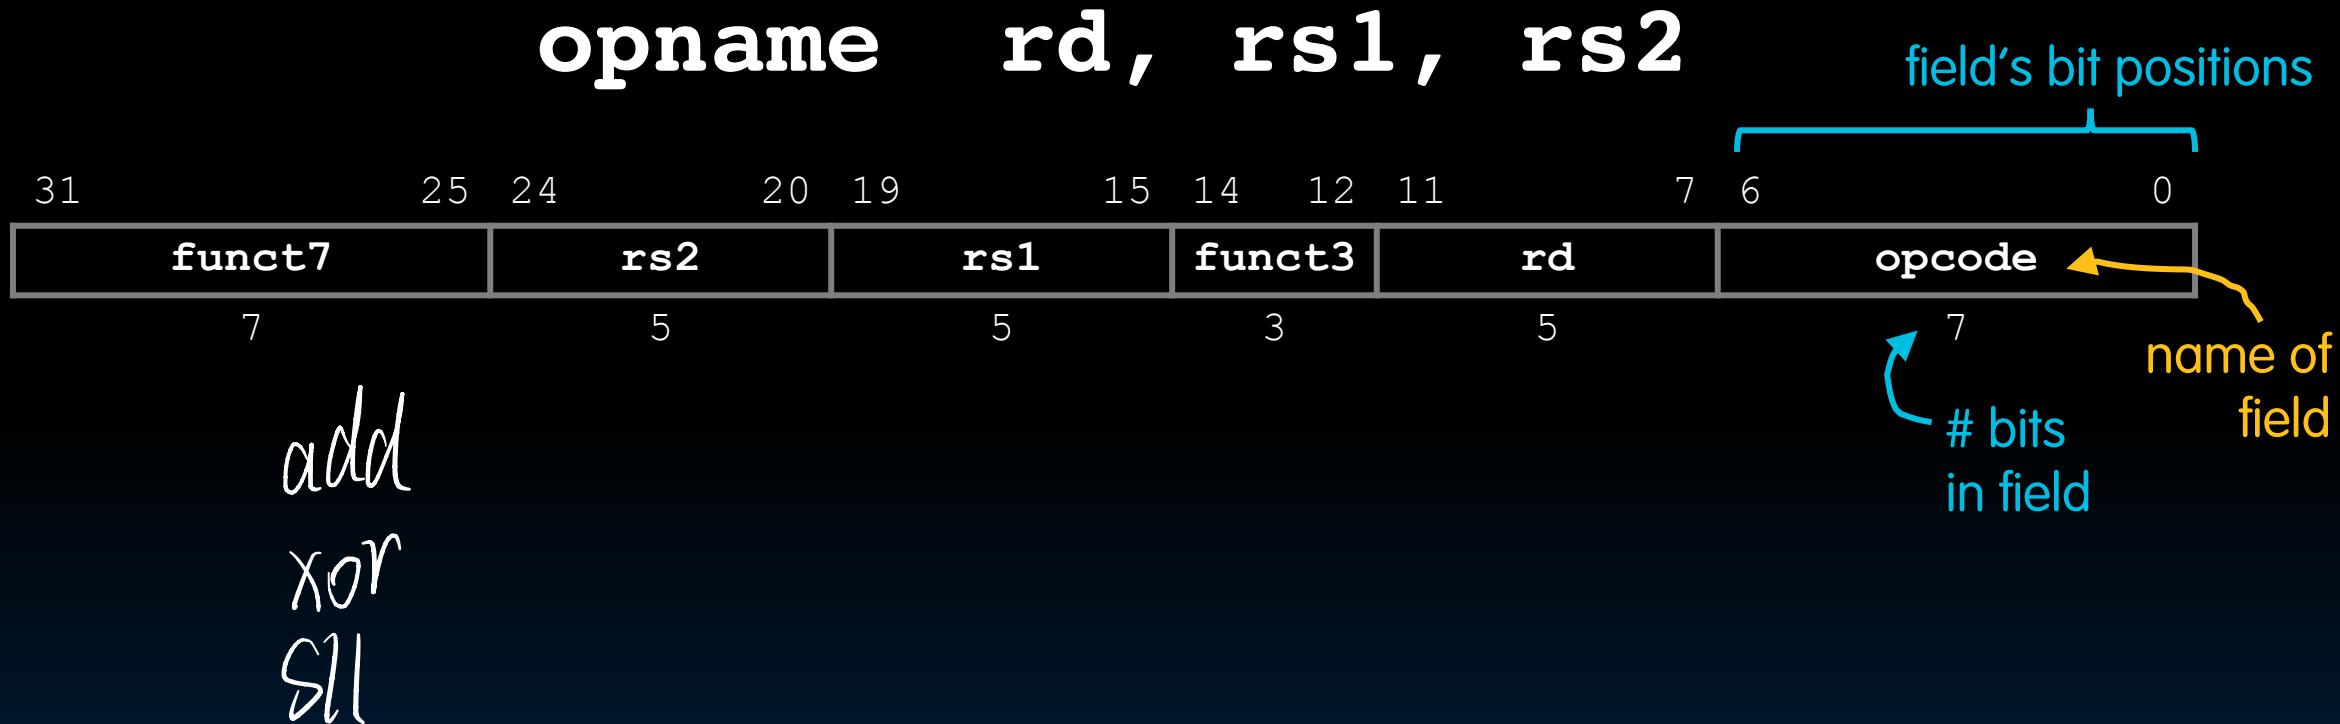
today

R-Format Layout

- Machine Language
- R-Format Layout
- I-Format Layout
- I-Format: Load
- S-Format Layout (Store)

R-Format Instruction Layout

- Register-Register Arithmetic Instructions (add, xor, sll, etc.)



R-Format Instruction Layout

- Register-Register Arithmetic Instructions (add, xor, sll, etc.)

opname rd, rs1, rs2



funct7, funct3 combined with opcode describes what operation to perform.

- Why not just a 17-bit field for simplicity?

We'll answer this later...

funct7, funct3 combined with opcode

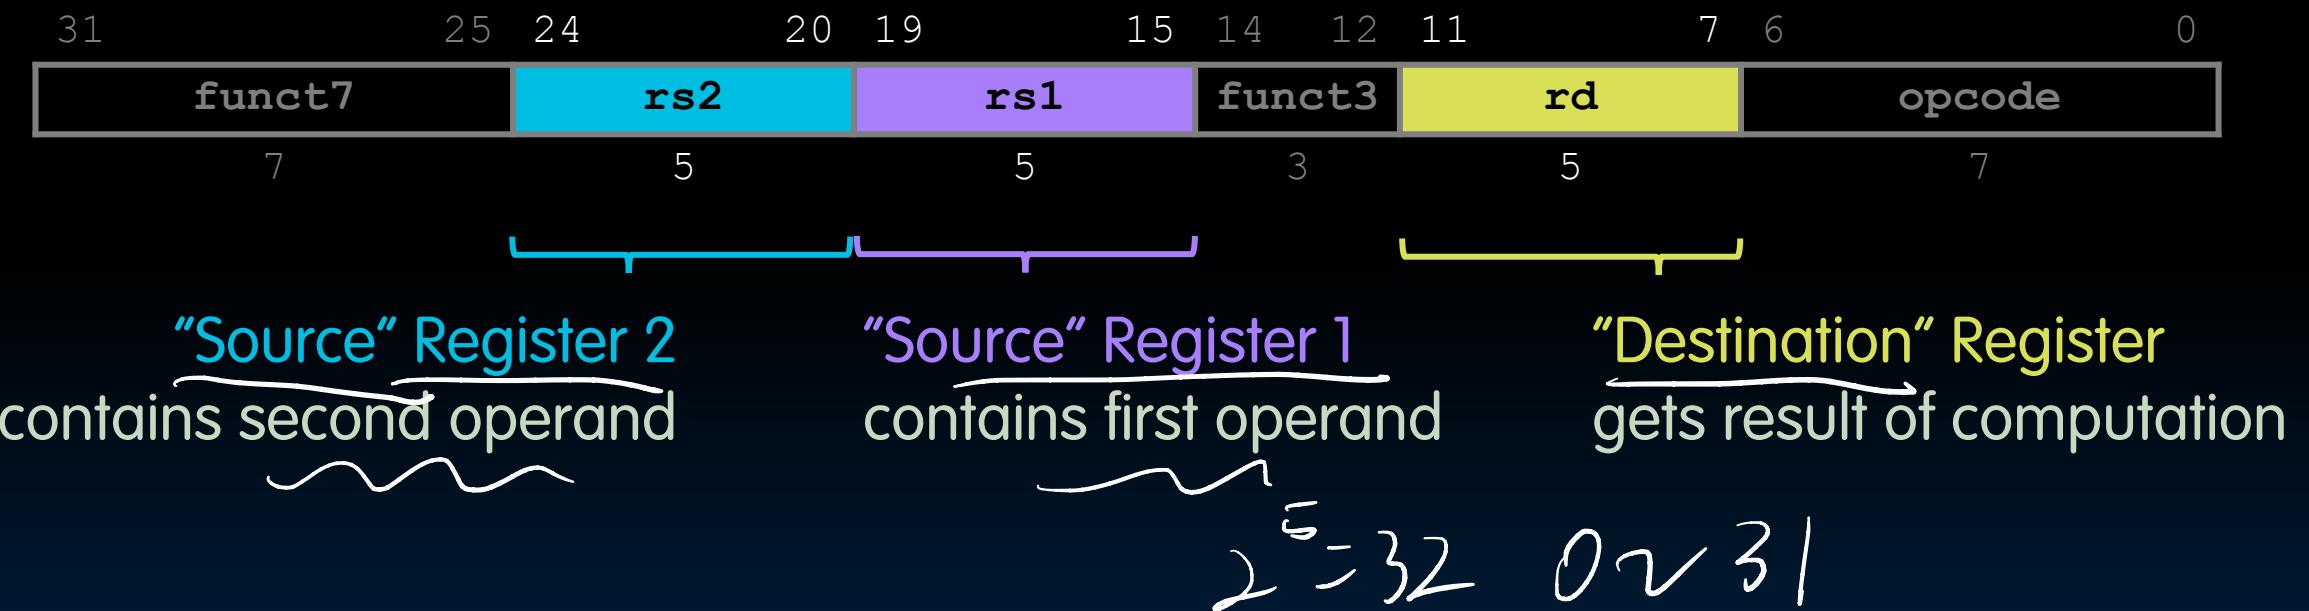
opcode partially specifies which instruction it is.

All R-Format instructions have opcode 0110011.

R-Format Instruction Layout

- Register-Register Arithmetic Instructions (add, xor, sll, etc.)

opname **rd**, **rs1**, **rs2**

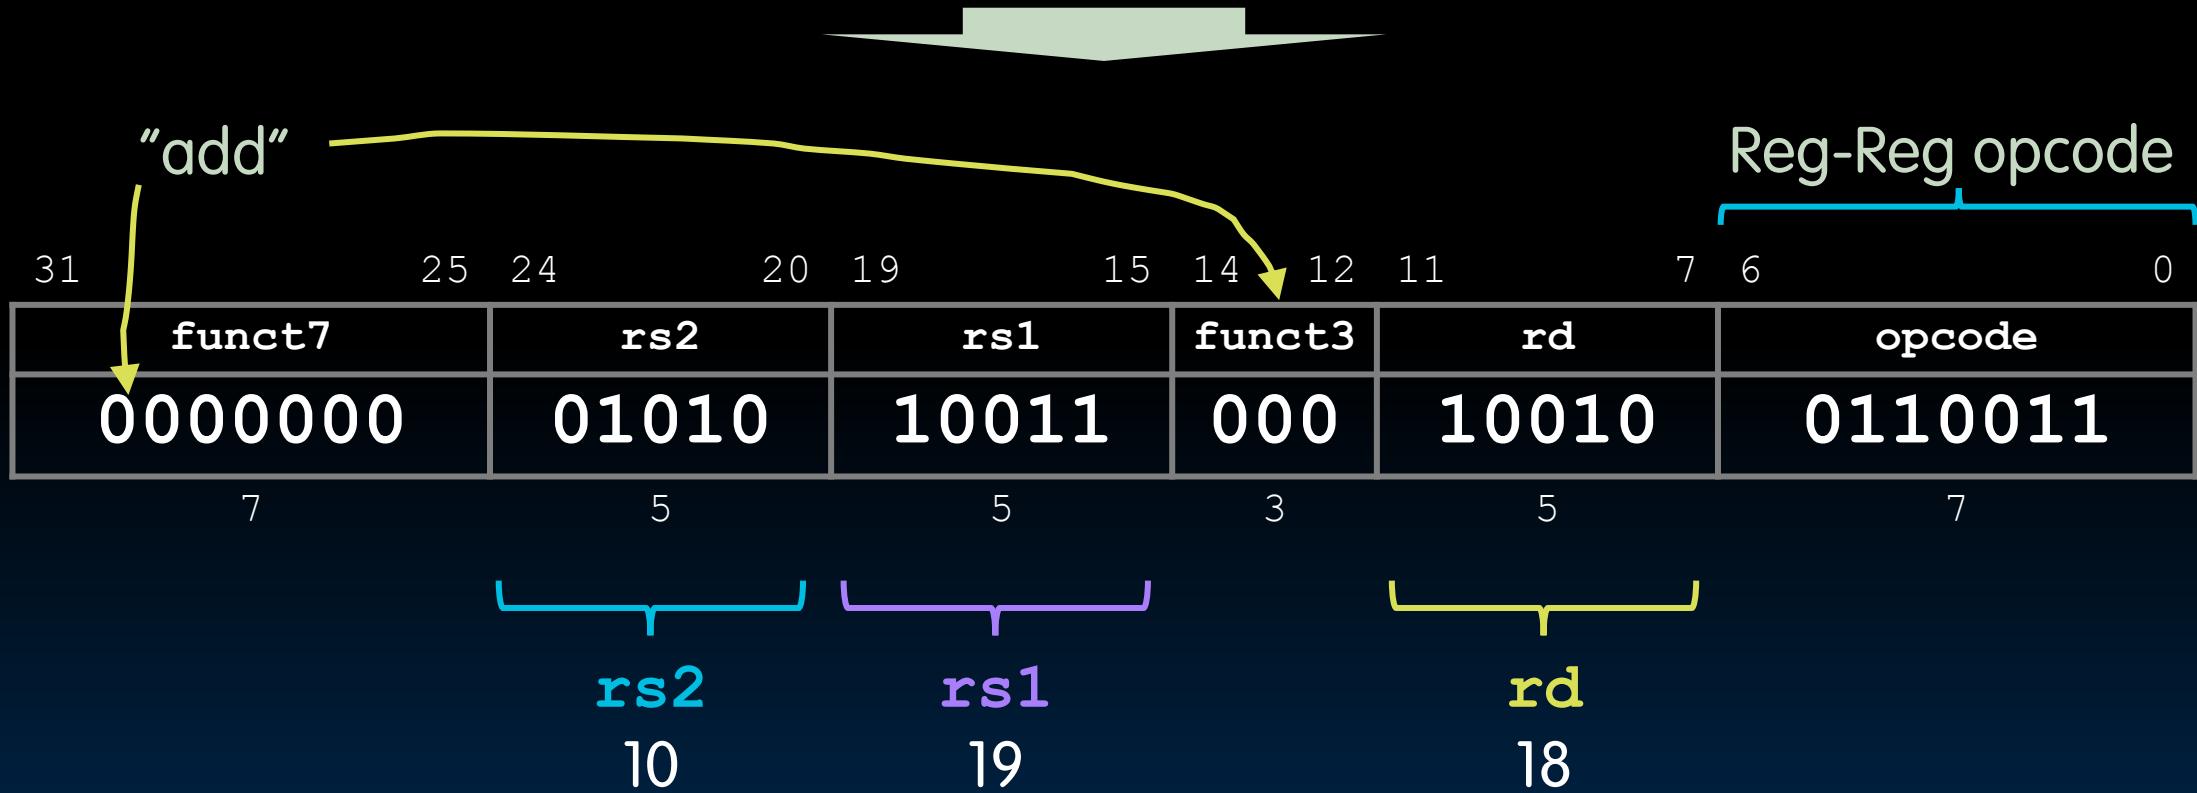


Register field (**rs1**, **rs2**, **rd**) holds a 5-bit unsigned integer [0-31] corresponding to a register number (**x0-x31**).

the number of register

R-Format Example

add x18, x19, x10



All ten RV32 R-Format Instructions

2's comp
negation
of rs2

			Eight funct3 fields for ten instructions			
			funct3	f	opcode	
funct7	rs2	rs1	000	rd	0110011	add
0000000	rs2	rs1	000	rd	0110011	sub
0100000	rs2	rs1	001	rd	0110011	sll
0000000	rs2	rs1	010	rd	0110011	slt
0000000	rs2	rs1	011	rd	0110011	sltu
0000000	rs2	rs1	100	rd	0110011	xor
0000000	rs2	rs1	101	rd	0110011	srl
0100000	rs2	rs1	101	rd	0110011	sra
0000000	rs2	rs1	110	rd	0110011	or
0000000	rs2	rs1	111	rd	0110011	and

} set less
than
(see refcard)

sign
extend

Different funct7 & funct3 encodings select different operations.

Your Turn

0000 0000 0010 00011 000 00101 0110011

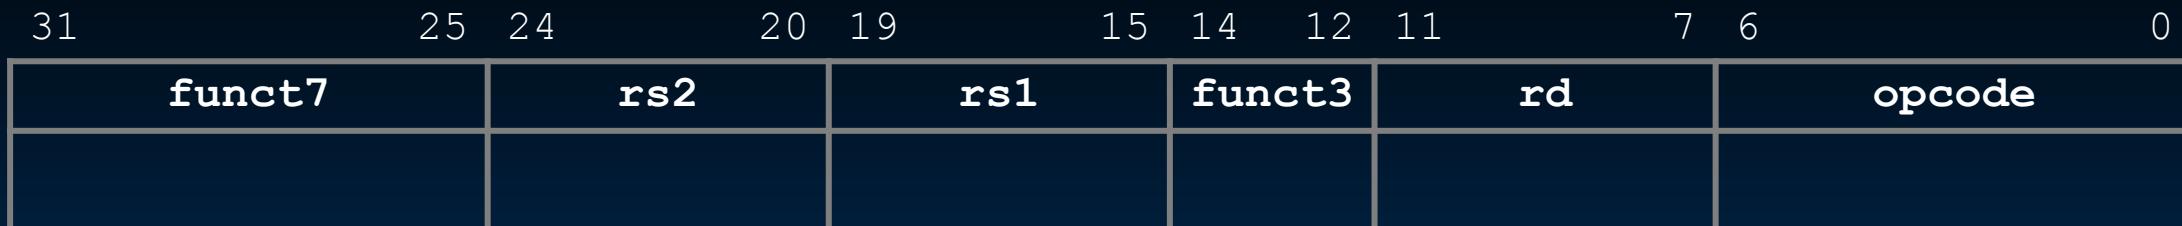
How do we encode

add x4, x3, x2 ? 0 2 | 8233

Lookup table:

funct7	rs2	rs1	funct3	rd	opcode	
0000000	rs2	rs1	000	rd	0110011	add
0100000	rs2	rs1	000	rd	0110011	sub
0000000	rs2	rs1	100	rd	0110011	xor
0000000	rs2	rs1	111	rd	0110011	and

- A. 4021 8233_{hex}
- B. 0021 82b3_{hex}
- C. 4021 82b3_{hex}
- D. 0021 8233_{hex}
- E. 0021 8234_{hex}
- F. Something else



pollev.com/yaml

🌐 When poll is active, respond at **pollev.com/yanl**

SMS Text **YANL** to **22333** once to join

How do we encode add x4, x3, x2?



Your Turn

How do we encode

add x4, x3, x2 ?

Lookup table:

funct7	rs2	rs1	funct3	rd	opcode	
0000000	rs2	rs1	000	xd	0110011	add
0100000	rs2	rs1	000	rd	0110011	sub
0000000	rs2	rs1	100	rd	0110011	xor
0000000	rs2	rs1	111	rd	0110011	and

31	25	24	20	19	15	14	12	11	7	6	0
funct7	rs2	rs1		funct3	rd	opcode					

4021 8233

0021 82b3

4021 82b3

0021 8233

0021 8234

Something else

Total Results: 0

Powered by Poll Everywhere

Start the presentation to see live content. For screen share software, share the entire screen. Get help at pollev.com/app

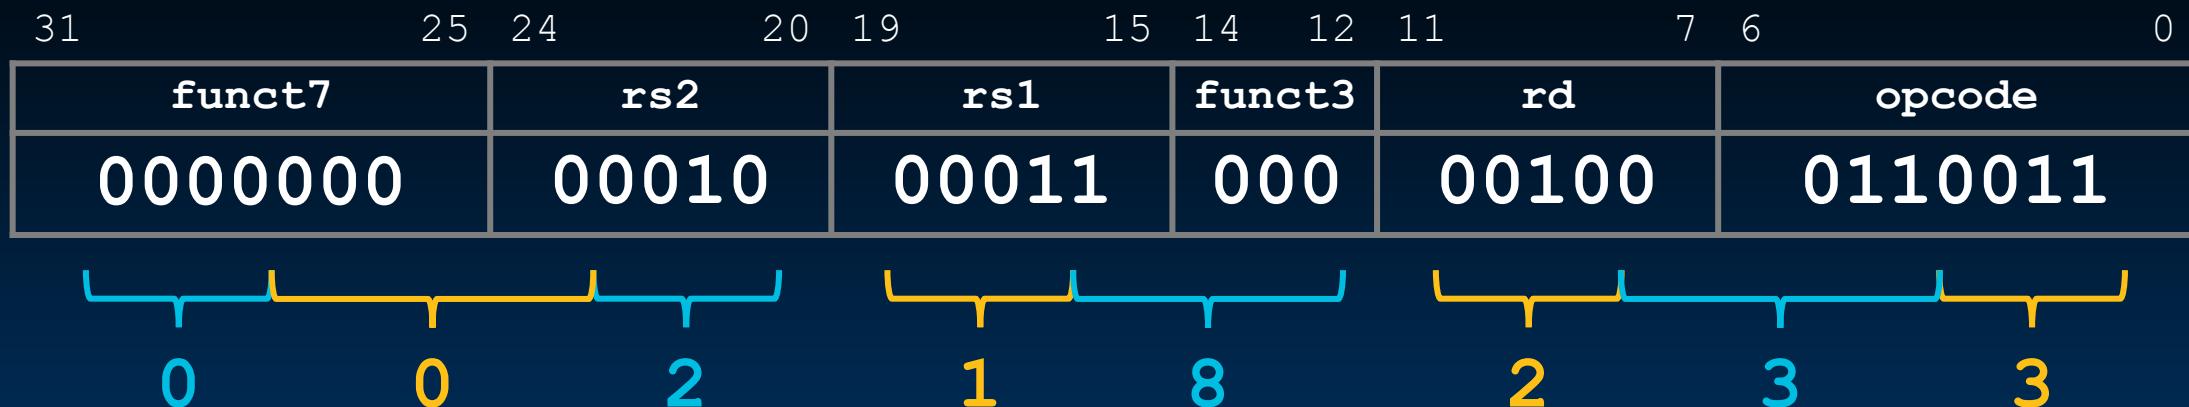
How do we encode

add x4, x3, x2 ?

Lookup table:

funct7	rs2	rs1	funct3	rd	opcode	
0000000	rs2	rs1	000	rd	0110011	add
0100000	rs2	rs1	000	rd	0110011	sub
0000000	rs2	rs1	100	rd	0110011	xor
0000000	rs2	rs1	111	rd	0110011	and

- A. 4021 8233_{hex}
- B. 0021 82b3_{hex}
- C. 4021 82b3_{hex}
- D. 0021 8233_{hex}
- E. 0021 8234_{hex}
- F. Something else



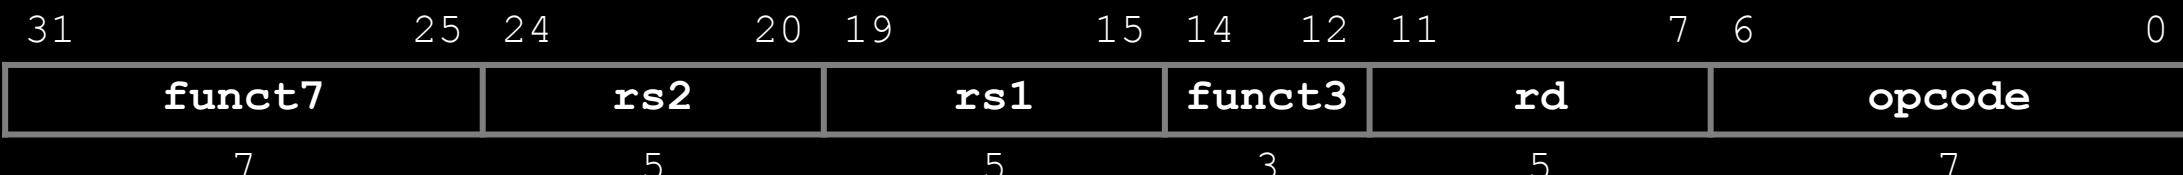
I-Format Layout

- Machine Language
- R-Format Layout
- I-Format Layout
- Load (I-Format)
- S-Format Layout (Store)

Immediate Fields Need to be Wider

- **R-Format:**

add rd, rs1, rs2



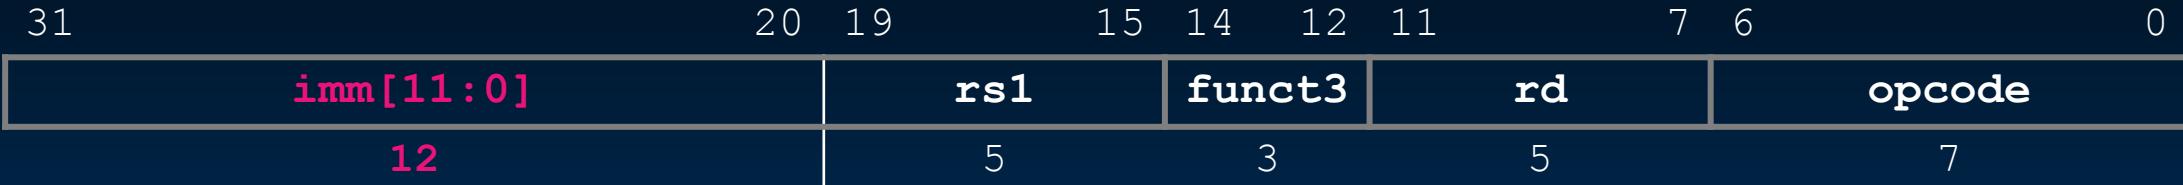
- **What about instructions with immediates?**

addi rd, rs1, imm

- If we used R-Format, we'd only be able to represent 32 values in the 5-bit **rs2** field.
- Immediates may be (and are often) much bigger!

- **Enter the I-Format:**

$$2^{12} = 4096$$



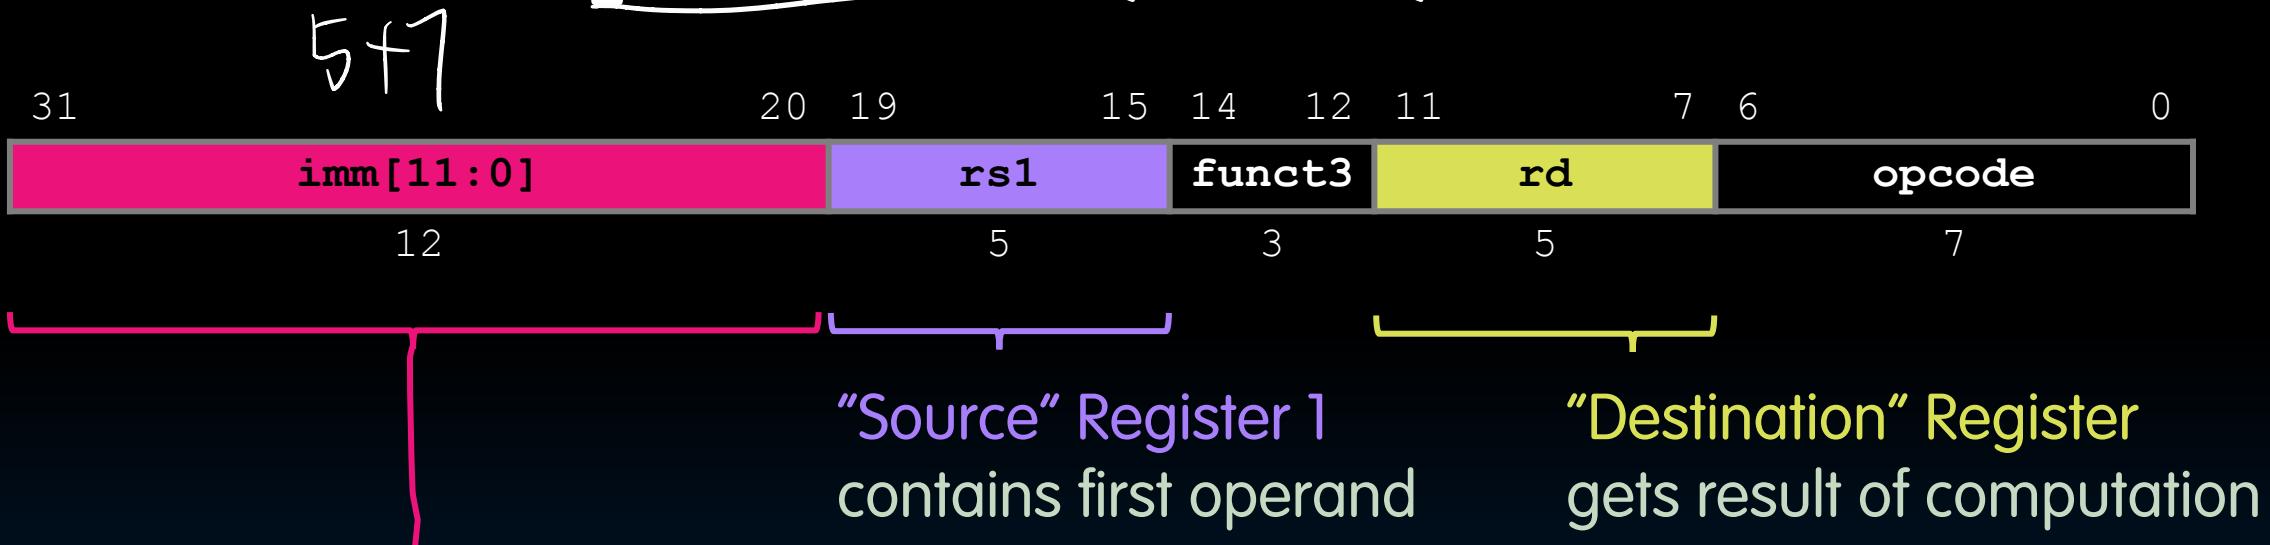
I-Format is *mostly* consistent with R-Format. Simplify how the CPU processes instructions!

I-Format Instruction Layout

addi, xor

- Register-Immediate Arithmetic Instructions (addi, xor, etc.)

opname rd, rs1, imm

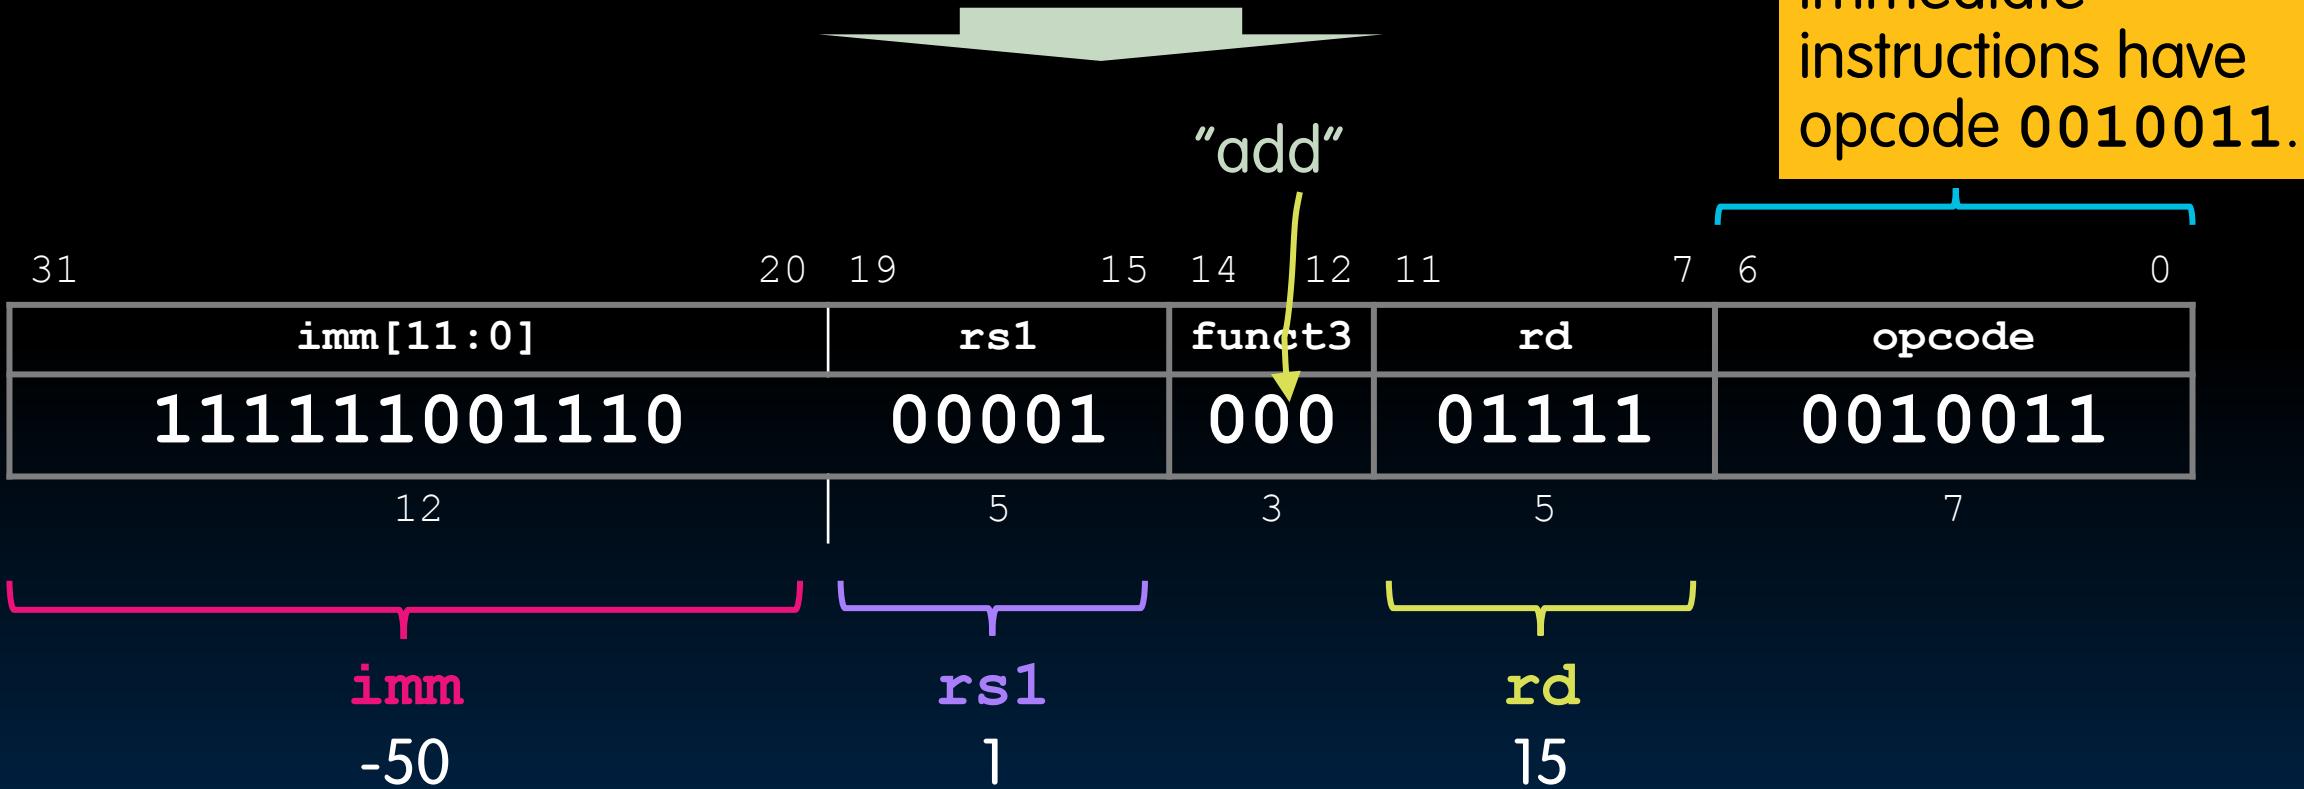


Imm[11:0] holds 12-bit-wide immediate values:

- Values in range $[-2048_{\text{ten}}, +2047_{\text{ten}}]$
- CPU sign-extends to 32 bits before use in an arithmetic operation
- How to handle immediates > 12 bits? (more next time)

I-Format Example

addi x15, x1, -50



All nine RV32 I-Format Arithmetic Instructions

Same funct3 fields as corresponding R-format operation (remember, no **subi**)

		funct3		opcode	
imm[11:0]	rs1	000	rd	0010011	addi
imm[11:0]	rs1	010	rd	0010011	slti
imm[11:0]	rs1	011	rd	0010011	sltiu
imm[11:0]	rs1	100	rd	0010011	xori
imm[11:0]	rs1	110	rd	0010011	ori
imm[11:0]	rs1	111	rd	0010011	andi
0000000	shamt	001	rd	0010011	slli
0000000	shamt	101	rd	0010011	srlti
0100000	shamt	101	rd	0010011	srai



"Shift by Immediate" instructions encode the shift amount in the lower-order 5 bits.

- We can only (meaningfully) shift 32-b word by 0-31 positions.
- One higher-order immediate bit used for sign extend (**srlti** vs. **srai**). Same bit position as in R-Format!

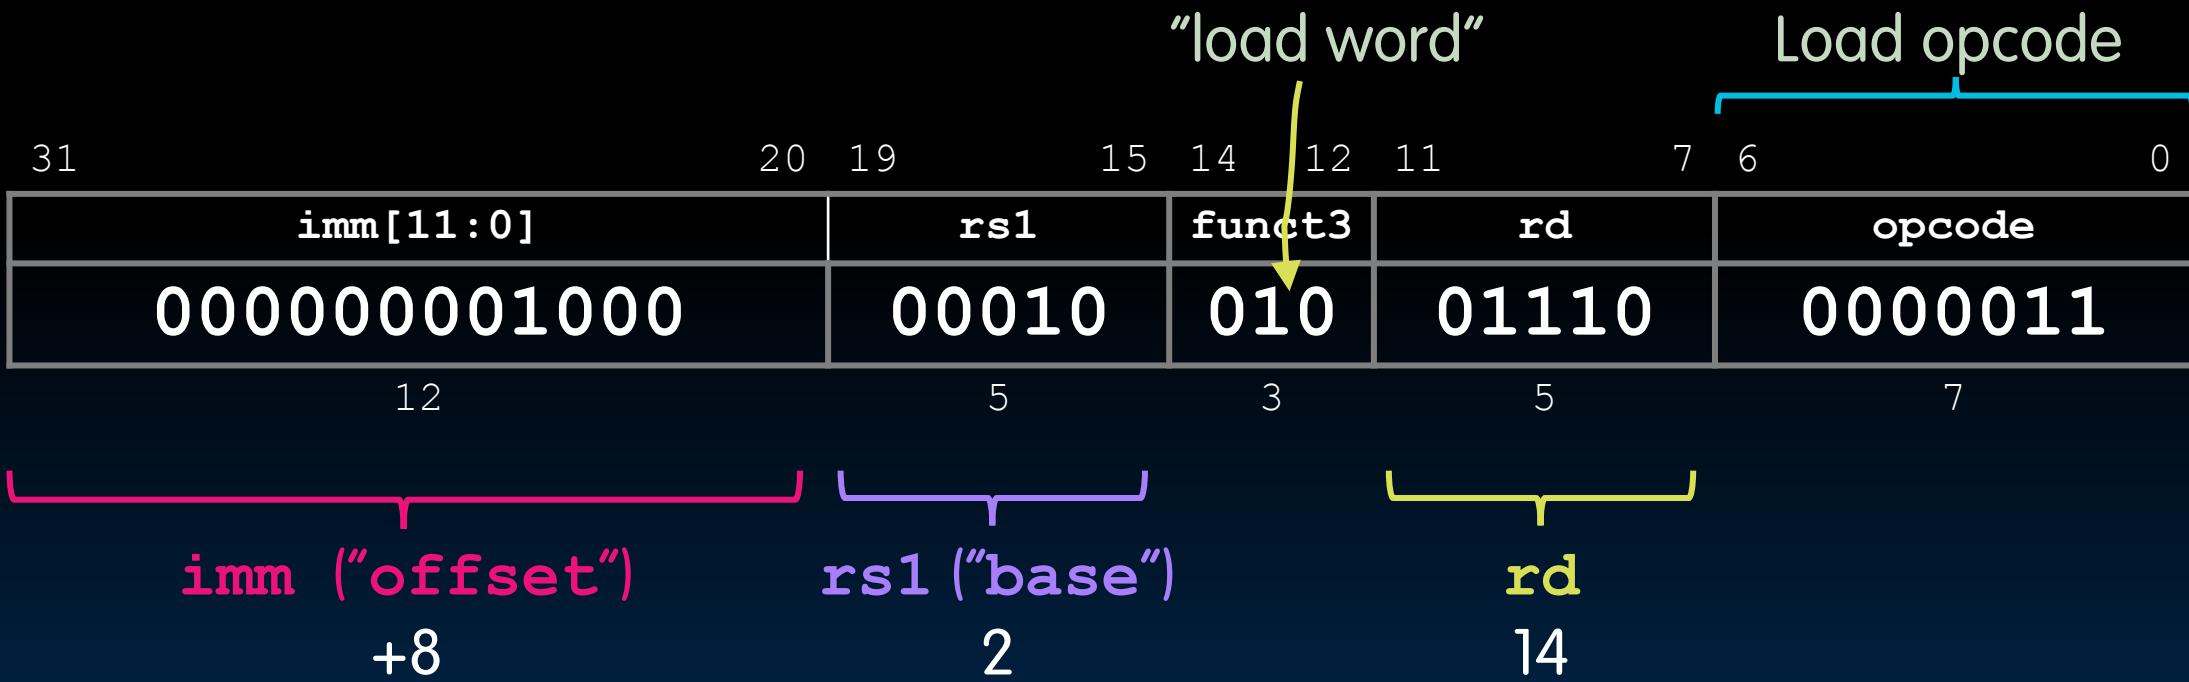
I-Format: Load

- Machine Language
- R-Format Layout
- I-Format Layout
- I-Format: Load
- S-Format Layout (Store)

Load Instructions Are Also I-Format: Example

Load word

lw x14, 8(x2)



Load address = (Base Register) + (Immediate Offset)

基址寄存器 + 偏移量

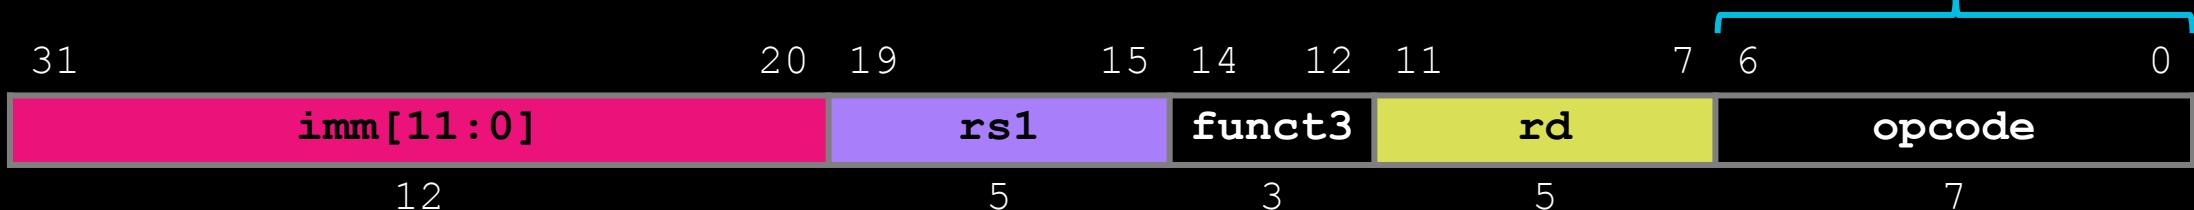
Garcia, Yan

I-Format Instruction Layout: Loads

- Load instructions use I-Format:

loadop rd, imm(rs1)

All load instructions have opcode 0000011.



Immediate “**offset**”
added to base address
to form memory address

offset

“**base**” Register
Base address of
load

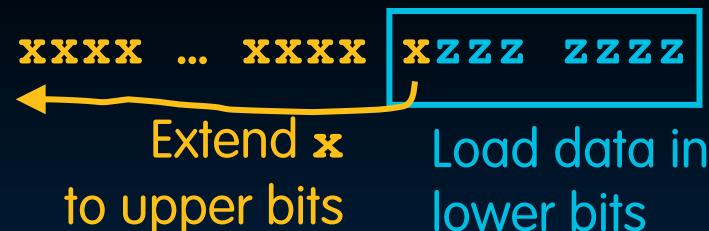
“**Destination**” Register
receives value loaded
from memory

All five RV32 Load Instructions

Encodes data size and “signedness” of load operation

		funct3		opcode	
imm[11:0]	rs1	000	rd	0000011	lb
imm[11:0]	rs1	001	rd	0000011	lh
imm[11:0]	rs1	010	rd	0000011	lw
imm[11:0]	rs1	100	rd	0000011	lbu
A bits imm[11:0]	rs1	101	rd	0000011	lhu

- **lb: “load byte,” lh: “load halfword” (16 bits)**
 - *Sign extend* to fill upper bits of destination 32-bit register.
- **lbu: “load unsigned byte,” lhu: “load unsigned halfword”**
 - *Zero extend* to fill upper bits of destination 32-bit register.
- **Note no lwu instruction in RISC-V.**
 - Simplicity: No need to sign/zero extend when copying 32 bits into 32-bit register.



S-Format Layout (Store)

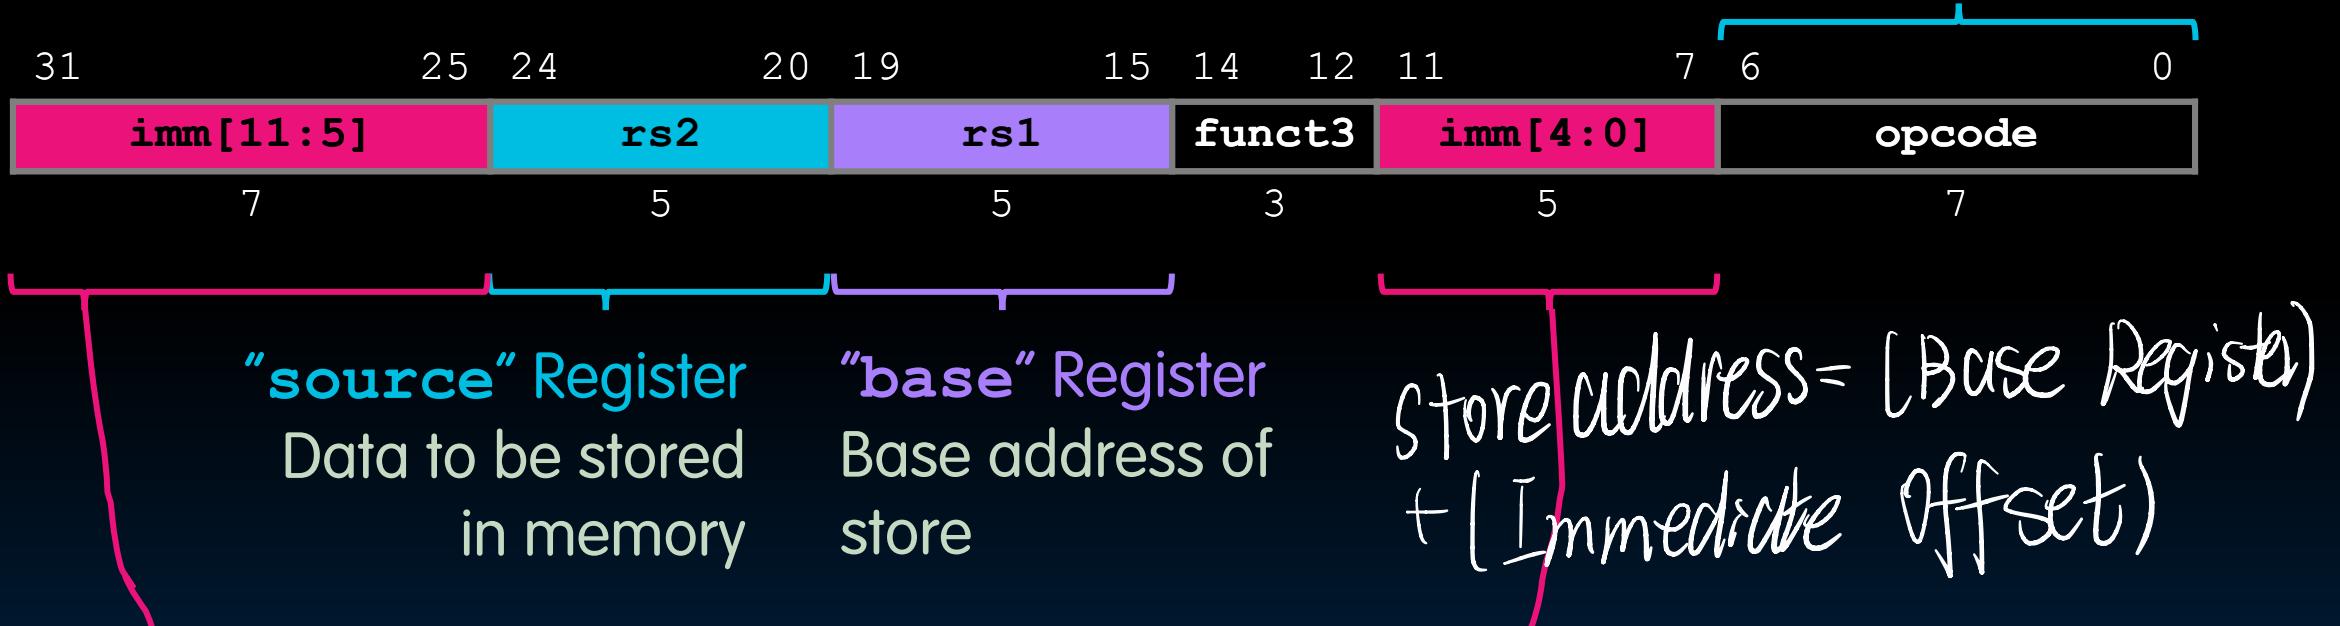
- Machine Language
- R-Format Layout
- I-Format Layout
- I-Format: Load
- S-Format Layout (Store)

S-Format Instruction Layout

- Store instructions have their S-Format.

storeop rs2, imm(rs1)

All store instructions have opcode 0100011.



Immediate "offset" added to base address to form memory address.

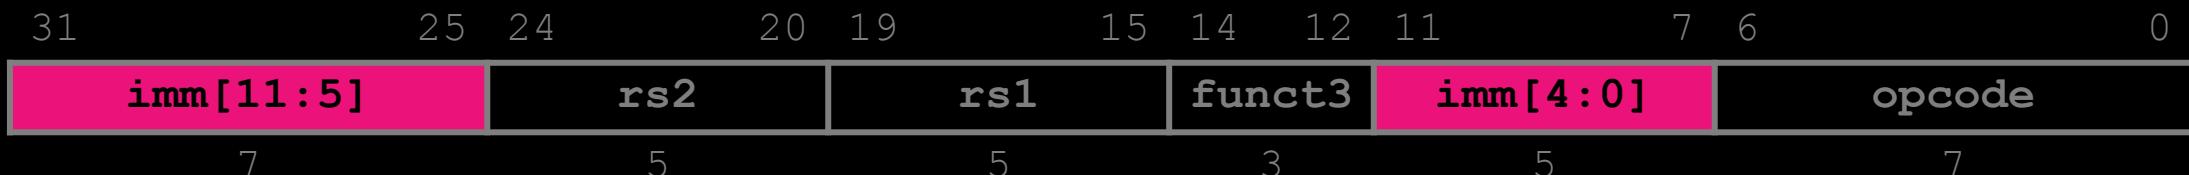
Store address = (Base Register) + (Immediate Offset)

- The immediate's higher 7 bits and lower 5 bits are in separate fields!

S-Format Simplifies Hardware Design

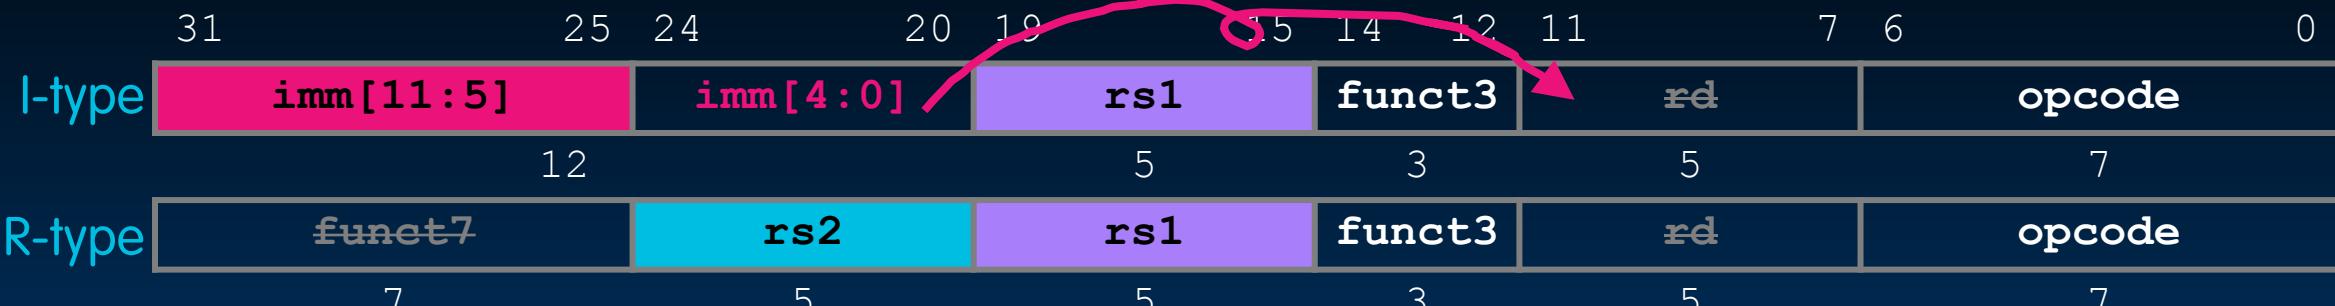
- Why split up the immediate field?

storeop rs2, imm(rs1)



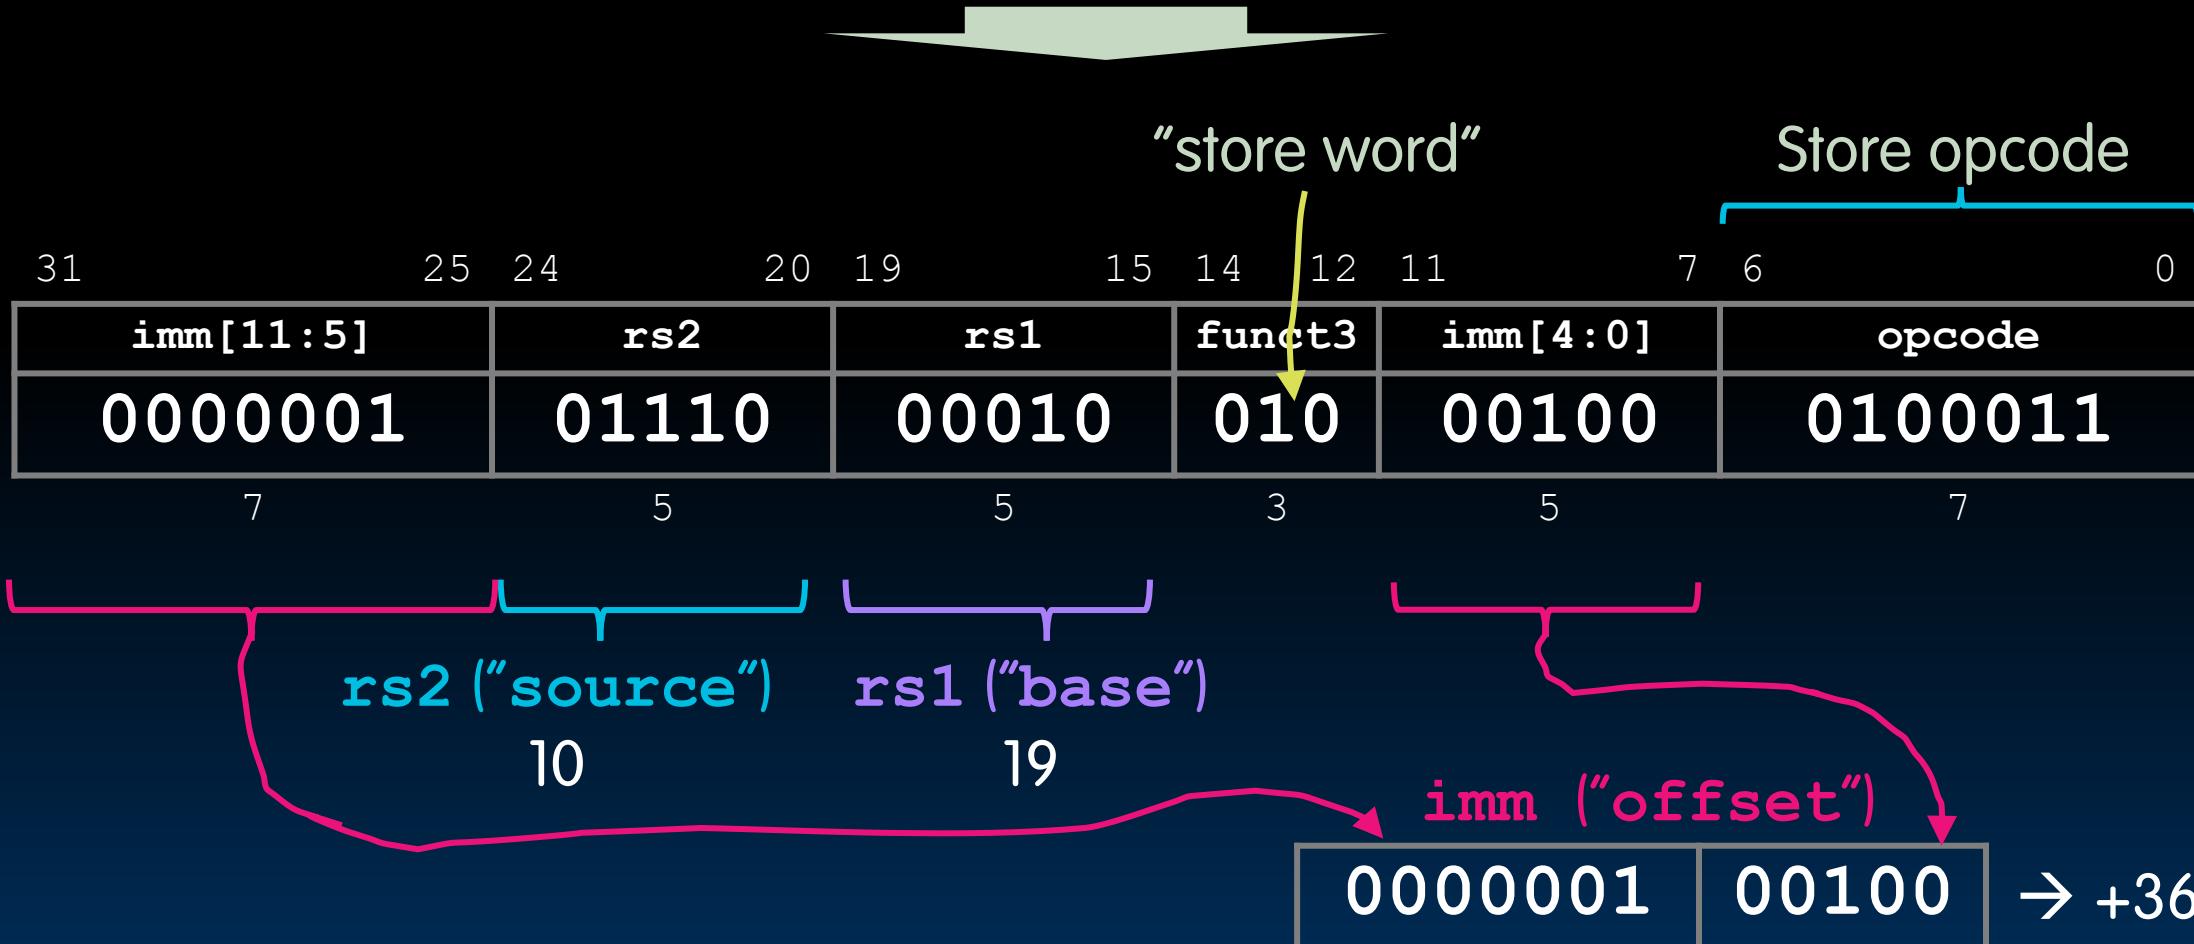
RISC-V design prioritizes keeping *register fields* in the same places. Immediates are less critical to hardware.

- Store needs immediate, two read registers, no destination register!
- Move lower 5 bits of immediate to rd field location in other formats:



S-Format Example

sw x14, 36(x2)



All three RV32 Store Instructions

Encodes data size of store operation						
			funct3	opcode		
imm[11:5]	rs2	rs1	000	imm[4:0]	0100011	sb
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011	sh
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011	sw

No sign/zero extending for store!
We only write to memory
the data width specified.

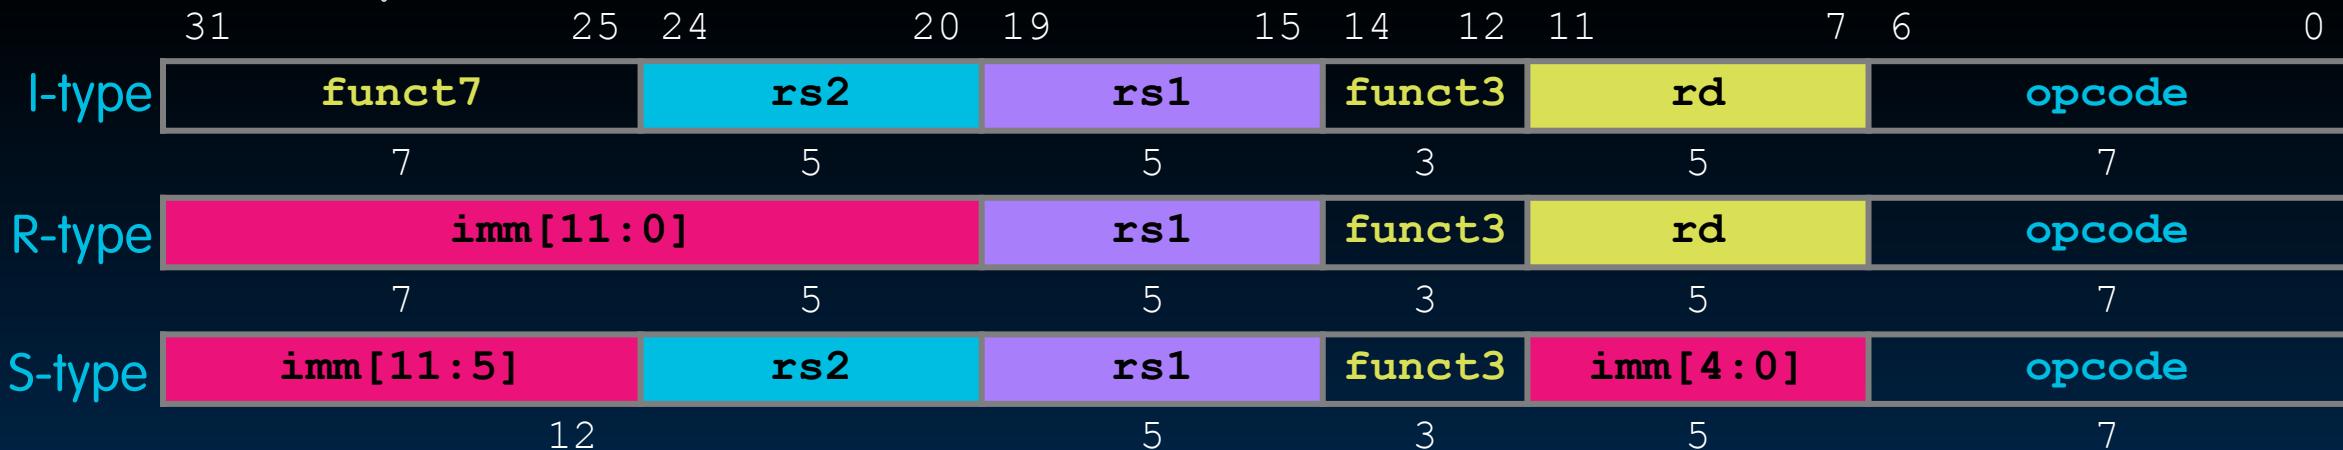


Same data width fields as
load instructions.

"And in Conclusion..."

- Simplification works for RISC-V: Instructions are same size as data word (32 bits) so that they use the same memory.
- Computer program stored as a series of these 32-bit numbers.
- RISC-V machine language instruction formats:

- Today:



- Next time: B-type, U-type, J-type