



UC Berkeley
Teaching Professor
Dan Garcia

CS61C

Great Ideas in **Computer Architecture** (a.k.a. Machine Structures)



UC Berkeley
Teaching Professor
Lisa Yan

Parallelism

From last time...

- We've discussed memory caching in detail. Caching in general shows up over and over in computer systems
 - Filesystem cache, Web page cache, Game databases / tablebases, Software memoization, Others?
- Big idea: if something is expensive but we want to do it repeatedly, do it once and cache the result.
- Cache design choices:
 - Size of cache: speed v. capacity
 - Block size (i.e., cache aspect ratio)
 - Write Policy (Write through v. write back)
 - Associativity choice of N (direct-mapped v. set v. fully associative)
 - Block replacement policy
 - 2nd level cache?
 - 3rd level cache?
- Use performance model to pick between choices, depending on programs, technology, budget, ...

New-School Machine Structures

Software
Parallel Requests
Assigned to computer
e.g., Search "Cats"

Parallel Threads
Assigned to core e.g., Lookup, Ads

Parallel Instructions
>1 instruction @ one time
e.g., 5 pipelined instructions

Parallel Data
>1 data item @ one time
e.g., Add of 4 pairs of words

Hardware descriptions
All gates work in parallel at same time

Harness
Parallelism &
Achieve High
Performance

Warehouse
Scale
Computer



Hardware



Computer

Core

Memory

Input/Output

Exec. Unit(s)

Functional
Block(s)

Main Memory

Logic Gates

$$\begin{array}{c} A_0 + B_0 \\ \quad A_1 + B_1 \end{array}$$

Parallelism (3)

Perspective...

"The real problem is that programmers have spent far too much time worrying about efficiency in the wrong places and at the wrong times; premature optimization is the root of all evil (or at least most of it) in programming."

– Donald Knuth
"The Art of Computer Programming"




Photo credit: Hector Garcia-Molina


Flynn's Taxonomy

Software vs. Hardware Parallelism

		Software	
		Sequential	Concurrent
Hardware	Serial	Matrix Multiply written in MatLab running on an Intel Pentium 4	Windows Vista Operating System running on an Intel Pentium 4
	Parallel	Matrix Multiply written in MATLAB running on an Intel Core i7	Windows Vista Operating System running on an Intel Core i7

- Choice of hardware and software parallelism are independent
 - Concurrent software can also run on serial hardware
 - Sequential software can also run on parallel hardware
- *Flynn's Taxonomy* is for parallel hardware

Single Instruction/Single Data Stream (SISD)




- Sequential computer that exploits no parallelism in either the instruction or data streams
- Examples of SISD architecture are traditional uniprocessor machines

This is what we did up to now in 61C

Single Instruction/Multiple Data Stream (SIMD)


SIMD



- Computer that applies a single instruction stream to multiple data streams for operations that may be naturally parallelized (e.g. SIMD instruction extensions or Graphics Processing Unit)

This lecture


Multiple Instruction/Multiple Data Stream (MIMD)



- Multiple autonomous processors simultaneously executing different instructions on different data
- MIMD architectures include multicore and Warehouse Scale Computers

Later in this module

Multiple Instruction/Single Data Stream (MISD)



- Exploits multiple instruction streams against a single data stream for data operations that can be naturally parallelized (e.g. certain kinds of array processors)
- MISD no longer commonly encountered, mainly of historical interest only

This has few applications. Not covered in 61C.

Flynn's Taxonomy



Prof. Michael J. Flynn

Data Streams		
	Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4
	Multiple	MISD: No examples today
		MIMD: Intel Xeon e5345 (Clovertown)

- SIMD and MIMD most commonly encountered today
- Most common parallel processing programming style:
Single Program Multiple Data ("SPMD")
 - Single program that runs on all processors of an MIMD
 - Cross-processor execution coordination through conditional expressions (will see later in Thread Level Parallelism)
- SIMD: specialized function units (hardware), for handling lock-step calculations involving arrays
 - Scientific computing, machine learning, signal processing, multimedia (audio/video processing)

SIMD Architectures


SIMD Architectures

- Data-Level Parallelism (DLP): Executing one operation on multiple data streams
- Example: Multiplying a coefficient vector by a data vector (e.g. in filtering)

$$y[i] := c[i] \times x[i], \quad 0 \leq i < n$$

Many programming languages support this paradigm (e.g., Snap!)


numbers from 1 to 4 × numbers from 1 to 4



- Sources of performance improvement:
 - One instruction is fetched & decoded for entire operation
 - Multiplications are known to be independent
 - Pipelining/concurrency in memory access as well




First SIMD Extensions: MIT Lincoln Labs TX-2, 1957



AE = "Arithmetic Element"

"It should be emphasized that the programmer can realize several arithmetic elements simultaneously.... The coupling (18,18) gives two complete, independent 18-bit arithmetic elements which are separately but simultaneously controlled by the instruction being executed. the (9,9,9,9) coupling gives four, 9-bit arithmetic elements.


“Advanced Digital Media Boost”



- To improve performance, Intel's SIMD instructions
 - Fetch one instruction, do the work of multiple instructions
 - MMX (MultiMedia eXtension, Pentium II processor family)
 - SSE (Streaming SIMD Extension, Pentium III and beyond)

Intel x86 SIMD Evolution

- Started with multimedia extensions (MMX)
 - New instructions every few years
 - New and wider registers
 - More parallelism




XMM Registers in SSE

- Architecture extended with eight 128-bit data registers
 - 64-bit address architecture: available as 16 64-bit registers
 - (XMM8 – XMM15)
 - e.g. 128-bit packed single-precision floating-point data type (doublewords), allows four single-precision operations to be performed simultaneously



Intel Architecture SSE2+128-Bit SIMD Data Types

Fundamental 128-Bit Packed SIMD Data Types




- Note: in Intel Architecture (unlike RISC V) a word is 16 bits
 - Single precision FP: Double word (32 bits)
 - Double precision FP: Quad word (64 bits)


SIMD Registers in AVX512

AVX512 state

Kmask k0..k7



zmm0..zmm31



ymm0..ymm15

xmm0..xmm7

High amounts of compute need large amounts of state to compensate for memory BW
AVX512 has 8x state compared to SSE (commensurate with its 8x flops level)

Intel confidential — presented under NDA only — under embargo until 6:01 a.m. PDT, June 19, 2017



Check Out a Typical Laptop (lscpu)

```
Model: 126
Model name: Intel(R) Core(TM) i7-1065G7 CPU @ 1.30GHz
Stepping: 5
CPU MHz: 1497.605
BogoMIPS: 2995.21
Hypervisor vendor: Microsoft
Virtualization type: full
L1d cache: 192 Kib
L1i cache: 128 Kib
L2 cache: 2 MiB
L3 cache: 8 MiB
Vulnerability Itlb multihit: KVM: Vulnerable
Vulnerability L1tf: Not affected
Vulnerability Mds: Not affected
Vulnerability Meltdown: Not affected
Vulnerability Spec store bypass: Mitigation; Speculative Store Bypass disabled via prctl and seccomp
Vulnerability Spectre v1: Mitigation; usercopy/swapgs barriers and __user pointer sanitization
Vulnerability Spectre v2: Mitigation; Enhanced IBRS, IBPB conditional, RSB filling
Vulnerability Srbds: Not affected
Vulnerability Tsx async abort: Not affected
Flags: fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush mmx fxsr sse sse2 ss ht syscall nx pdpe1gb rdtscp lm constant_tsc rep_good nopl xtopology cpuid pnpi pclmulqdq ssse3 fma cx16 pcid sse4_1 sse4_2 movbe popcnt aes xsave avx f16c rdrand hypervisor lahf_lm abm 3dnowprefetch invpcid_single ssbd ibrs ibpb stibp ibrs_enhanced fsgsbase bmi1 avx2 smep bmi2 erms invpcid avx512f avx512dq rdseed adx smap avx512ifma clflushopt avx512cd sha_ni avx512bw avx512vl xsaveopt xsavec xgetbv1 xsaves avx512vbmi umip avx512_vbmi2 gfni vaes vpclmulqdq avx512_vnni avx512_bitalg avx512_vpocntdq rdpid flush_l1d arch_capabilities
```

SIMD Array Processing

Example: SIMD Array Processing

for each f in array:

 f = sqrt(f)

} pseudocode

for each f in array {

 load f to the floating-point register

 calculate the square root

 write the result from the register to memory

}

} SISD

for every 4 members in array {

 load 4 members to the SSE register

 calculate 4 square roots in one operation

 write the result from the register to memory

}

} SIMD

Example: Add Single-Precision FP Vectors

Computation to be performed:

```
vec_res.x = v1.x + v2.x;  
vec_res.y = v1.y + v2.y;  
vec_res.z = v1.z + v2.z;  
vec_res.w = v1.w + v2.w;
```

move from mem to XMM register
memory aligned, packed single precision

add from mem to XMM register
packed single precision

SSE Instruction Sequence:

```
movaps address-of-v1, %xmm0  
// v1.w | v1.z | v1.y | v1.x -> xmm0  
addps address-of-v2, %xmm0  
// v1.w+v2.w | v1.z+v2.z | v1.y+v2.y | v1.x+v2.x -> xmm0  
movaps %xmm0, address-of-vec_res
```



Intel SSE Intrinsics

- Intrinsics are C functions and procedures for putting in assembly language, including SSE instructions
 - With intrinsics, can program using these instructions indirectly
 - One-to-one correspondence between SSE instructions and intrinsics


Intrinsics:

- Vector data type:
 `_m128d`
- Load and store operations:
 `_mm_load_pd`
 `_mm_store_pd`
- Arithmetic:
 `_mm_add_pd`
 `_mm_mul_pd`

Corresponding SSE instructions:

- MOVAPD/aligned, packed double
- MOVAPD/aligned, packed double
- ADDPD/add, packed double
- MULPD/multiple, packed double

Back to RISC-V: Vector Extensions (Draft)




- To improve RISC-V performance, add SIMD instructions (and hardware) – V extension
 - Fetch one instruction, do the work of multiple instructions
 - OP denotes a vector instruction, prefix v – vector register
 - vadd vd, vs1, vs2** (adds two vectors stored in vector registers)
 - Assume vectors are 512-bits wide



L28 Flynn Taxonomy ... how much faster is matrix multiply coded in C using AVX extensions vs. Python

1x
10x
100x
1,000x
10,000x
100,000x
1,000,000x





“And in Conclusion...”

- Flynn Taxonomy of Parallel Architectures
 - SIMD: Single Instruction Multiple Data
 - MIMD: Multiple Instruction Multiple Data
 - SISD: Single Instruction Single Data
 - MISD: Multiple Instruction Single Data (unused)
- Intel AVX SIMD Instructions
 - One instruction fetch that operates on multiple operands simultaneously
 - 512/256/128/64-bit AVX registers
 - Use C intrinsics
- Mantra
 - “...premature optimization is the root of all evil (or at least most of it) in programming.” – Donald Knuth

Great Matrix
Multiply example
in Bonus slides


https://youtu.be/_z_0l6FmnbU

<https://youtu.be/V-uBL49SFK0>

Matrix Multiply Example

Application: Machine Learning

- Inference in machine learning applications




Reference Problem: Matrix Multiplication

- Matrix multiplication
 - Basic operation in many engineering, data, and imaging processing tasks
 - Image filtering, noise reduction, machine learning...
 - Many closely related operations
- **dgemm**
 - double-precision floating-point matrix multiplication
 - In FORTRAN

Matrices



- Square matrix of dimension NxN



Matrix Multiplication

$$C = A * B$$

$$C_{ij} = \sum_k (A_{ik} * B_{kj})$$



Example: 2 x 2 Matrix Multiply

Matrix Multiply:

$$C_{i,j} = (A \times B)_{i,j} = \sum_{k=1}^2 A_{i,k} \times B_{k,j}$$

$$\begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix} \times \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix} = \begin{bmatrix} C_{1,1} = A_{1,1}B_{1,1} + A_{1,2}B_{2,1} & C_{1,2} = A_{1,1}B_{1,2} + A_{1,2}B_{2,2} \\ C_{2,1} = A_{2,1}B_{1,1} + A_{2,2}B_{2,1} & C_{2,2} = A_{2,1}B_{1,2} + A_{2,2}B_{2,2} \end{bmatrix}$$
$$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \times \begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix} = \begin{bmatrix} C_{1,1} = 1*1 + 0*2 = 1 & C_{1,2} = 1*3 + 0*4 = 3 \\ C_{2,1} = 0*1 + 1*2 = 2 & C_{2,2} = 0*3 + 1*4 = 4 \end{bmatrix}$$

Reference: Python

- Matrix multiplication in Python


```
def dgemm(N, a, b, c):
    for i in range(N):
        for j in range(N):
            c[i+j*N] = 0
            for k in range(N):
                c[i+j*N] += a[i+k*N] * b[k+j*N]
```

N	Python [MFLOPs]
32	5.4
160	5.5
480	5.4
960	5.3

- 1 MFLOP = 1 Million floating-point operations per second (fadd, fmul)
- **dgemm (N ...)** takes $2 \cdot N^3$ FLOPs

- $c = a * b$
- a, b, c are $N \times N$ matrices

```
// Scalar; P&H p. 226
void dgemm_scalar(int N, double *a, double *b, double *c) {
    for (int i=0; i<N; i++)
        for (int j=0; j<N; j++) {
            double cij = 0;
            for (int k=0; k<N; k++)
                // a[i][k] * b[k][j]
                cij += a[i+k*N] * b[k+j*N];
            // c[i][j]
            c[i+j*N] = cij;
        }
}
```



Timing Program Execution

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

int main(void) {
    // start time
    // Note: clock() measures execution time, not real time
    //       big difference in shared computer environments
    //       and with heavy system load
    clock_t start = clock();

    // task to time goes here:
    // dgemm(N, ...);

    // "stop" the timer
    clock_t end = clock();

    // compute execution time in seconds
    double delta_time = (double)(end-start)/CLOCKS_PER_SEC;
}
```



N	Python [GFLOPS]	C [GFLOPS]
32	0.0054	1.30
160	0.0055	1.30
480	0.0054	1.32
960	0.0053	0.91

Which class gives you this kind of power?
We could stop here ... but why? Let's do better!

Example: 2×2 Matrix Multiply

Matrix Multiply:

$$C_{i,j} = (A \times B)_{i,j} = \sum_{k=1}^2 A_{i,k} \times B_{k,j}$$

$$\begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix} \times \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix} = \begin{bmatrix} C_{1,1} = A_{1,1}B_{1,1} + A_{1,2}B_{2,1} & C_{1,2} = A_{1,1}B_{1,2} + A_{1,2}B_{2,2} \\ C_{2,1} = A_{2,1}B_{1,1} + A_{2,2}B_{2,1} & C_{2,2} = A_{2,1}B_{1,2} + A_{2,2}B_{2,2} \end{bmatrix}$$
$$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \times \begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix} = \begin{bmatrix} C_{1,1} = 1*1 + 0*2 = 1 & C_{1,2} = 1*3 + 0*4 = 3 \\ C_{2,1} = 0*1 + 1*2 = 2 & C_{2,2} = 0*3 + 1*4 = 4 \end{bmatrix}$$

Example: 2×2 Matrix Multiply

- Initialization

C_1	0	:	0
C_2	0	:	0

Example: 2×2 Matrix Multiply

- Initialization

$$\begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix} \times \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix} = \begin{bmatrix} C_{1,1} = A_{1,1}B_{1,1} + A_{1,2}B_{2,1} & C_{1,2} = A_{1,1}B_{1,2} + A_{1,2}B_{2,2} \\ C_{2,1} = A_{2,1}B_{1,1} + A_{2,2}B_{2,1} & C_{2,2} = A_{2,1}B_{1,2} + A_{2,2}B_{2,2} \end{bmatrix}$$

$$\begin{array}{c} C_1 \quad \boxed{0} \quad | \quad 0 \\ C_2 \quad \boxed{0} \quad | \quad 0 \end{array}$$

- $i = 1$

$$A \quad \boxed{A_{1,1}} \quad | \quad A_{2,1}$$

_mm_load_pd: Stored in memory in Column order

$$\begin{array}{c} B_1 \quad \boxed{B_{1,1}} \quad | \quad B_{1,1} \\ B_2 \quad \boxed{B_{1,2}} \quad | \quad B_{1,2} \end{array}$$

_mm_load1_pd: SSE instruction that loads a double word and stores it in the high and low double words of the XMM register

Example: 2×2 Matrix Multiply

$$\begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix} \times \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix} = \begin{bmatrix} C_{1,1} = A_{1,1}B_{1,1} + A_{1,2}B_{2,1} & C_{1,2} = A_{1,1}B_{1,2} + A_{1,2}B_{2,2} \\ C_{2,1} = A_{2,1}B_{1,1} + A_{2,2}B_{2,1} & C_{2,2} = A_{2,1}B_{1,2} + A_{2,2}B_{2,2} \end{bmatrix}$$

$$\begin{bmatrix} C_1 & 0+A_{1,1}B_{1,1} & 0+A_{2,1}B_{1,1} \\ C_2 & 0+A_{1,1}B_{1,2} & 0+A_{2,1}B_{1,2} \end{bmatrix}$$

- I = 1, intermediate result

$$A \quad \begin{bmatrix} A_{1,1} & | & A_{2,1} \end{bmatrix}$$

`c1 = _mm_add_pd(c1,_mm_mul_pd(a,b1));`
`c2 = _mm_add_pd(c2,_mm_mul_pd(a,b2));`
 SSE instructions first do parallel multiplies and then parallel adds in XMM registers

`_mm_load_pd`: Stored in memory in Column order

$$\begin{bmatrix} B_1 & \begin{bmatrix} B_{1,1} & | & B_{1,1} \end{bmatrix} \\ B_2 & \begin{bmatrix} B_{1,2} & | & B_{1,2} \end{bmatrix} \end{bmatrix}$$

`_mm_load1_pd`: SSE instruction that loads a double word and stores it in the high and low double words of the XMM register

Example: 2×2 Matrix Multiply

$$\begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix} \times \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix} = \begin{bmatrix} C_{1,1} = A_{1,1}B_{1,1} + A_{1,2}B_{2,1} & C_{1,2} = A_{1,1}B_{1,2} + A_{1,2}B_{2,2} \\ C_{2,1} = A_{2,1}B_{1,1} + A_{2,2}B_{2,1} & C_{2,2} = A_{2,1}B_{1,2} + A_{2,2}B_{2,2} \end{bmatrix}$$

$$\begin{array}{c|c|c} C_1 & 0+A_{1,1}B_{1,1} & 0+A_{2,1}B_{1,1} \\ \hline C_2 & 0+A_{1,1}B_{1,2} & 0+A_{2,1}B_{1,2} \end{array}$$

- I = 2, intermediate result

$$A \quad \boxed{A_{1,2}} \quad \boxed{A_{2,2}}$$

`c1 = _mm_add_pd(c1,_mm_mul_pd(a,b1));`
`c2 = _mm_add_pd(c2,_mm_mul_pd(a,b2));`
 SSE instructions first do parallel multiplies and then parallel adds in XMM registers

`_mm_load_pd`: Stored in memory in Column order

$$\begin{array}{c|c|c} B_1 & \boxed{B_{2,1}} & \boxed{B_{2,1}} \\ \hline B_2 & \boxed{B_{2,2}} & \boxed{B_{2,2}} \end{array}$$

`_mm_load1_pd`: SSE instruction that loads a double word and stores it in the high and low double words of the XMM register

Example: 2×2 Matrix Multiply

$C_{1,1}$ $C_{1,2}$

$$\begin{matrix} C_1 & \boxed{A_{1,1}B_{1,1} + A_{1,2}B_{2,1}} & \boxed{A_{2,1}B_{1,1} + A_{2,2}B_{2,1}} \\ C_2 & \boxed{A_{1,1}B_{1,2} + A_{1,2}B_{2,2}} & \boxed{A_{2,1}B_{1,2} + A_{2,2}B_{2,2}} \end{matrix}$$

$C_{2,1}$ $C_{2,2}$

- I = 2, intermediate result

$$A \quad \boxed{A_{1,2}} \quad | \quad \boxed{A_{2,2}}$$

$$\begin{matrix} B_1 & \boxed{B_{2,1}} & | & \boxed{B_{2,1}} \\ B_2 & \boxed{B_{2,2}} & | & \boxed{B_{2,2}} \end{matrix}$$

$$\begin{bmatrix} A_{1,1} \\ A_{2,1} \end{bmatrix} \times \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix} = \begin{bmatrix} C_{1,1} = A_{1,1}B_{1,1} + A_{1,2}B_{2,1} & C_{1,2} = A_{1,1}B_{1,2} + A_{1,2}B_{2,2} \\ C_{2,1} = A_{2,1}B_{1,1} + A_{2,2}B_{2,1} & C_{2,2} = A_{2,1}B_{1,2} + A_{2,2}B_{2,2} \end{bmatrix}$$

`c1 = _mm_add_pd(c1,_mm_mul_pd(a,b1));`
`c2 = _mm_add_pd(c2,_mm_mul_pd(a,b2));`
 SSE instructions first do parallel multiplies and then parallel adds in XMM registers

`_mm_load_pd`: Stored in memory in Column order

`_mm_load1_pd`: SSE instruction that loads a double word and stores it in the high and low double words of the XMM register



Example: 2×2 Matrix Multiply

```
#include <stdio.h>
// header file for SSE compiler intrinsics
#include <emmintrin.h>


// NOTE: vector registers will be represented in
// comments as v1 = [ a | b ]
// where v1 is a variable of type __m128d and
// a, b are doubles

int main(void) {
    // allocate A,B,C aligned on 16-byte boundaries
    double A[4] __attribute__ ((aligned (16)));
    double B[4] __attribute__ ((aligned (16)));
    double C[4] __attribute__ ((aligned (16)));
    int lda = 2;
    int i = 0;
    // declare several 128-bit vector variables
    __m128d c1,c2,a,b1,b2;
```

```
// Initialize A, B, C for example
/* A = (note column order!)
   1 0
   0 1
*/
A[0] = 1.0; A[1] = 0.0; A[2] = 0.0;
A[3] = 1.0;

/* B = (note column order!)
   1 3
   2 4
*/
B[0] = 1.0; B[1] = 2.0; B[2] = 3.0;
B[3] = 4.0;

/* C = (note column order!)
   0 0
   0 0
*/
C[0] = 0.0; C[1] = 0.0; C[2] = 0.0;
C[3] = 0.0;
```



Example: 2×2 Matrix Multiply

```
// used aligned loads to set
// c1 = [c_11 | c_21]
c1 = _mm_load_pd(C+0*lda);
// c2 = [c_12 | c_22]
c2 = _mm_load_pd(C+1*lda);

for (i = 0; i < 2; i++) {
/* a =
   i = 0: [a_11 | a_21]
   i = 1: [a_12 | a_22]
*/
a = _mm_load_pd(A+i*lda);
/* b1 =
   i = 0: [b_11 | b_11]
   i = 1: [b_21 | b_21]
*/
b1 = _mm_load1_pd(B+i+0*lda);
/* b2 =
   i = 0: [b_12 | b_12]
   i = 1: [b_22 | b_22]
*/
b2 = _mm_load1_pd(B+i+1*lda);
```

```
/* c1 =
   i = 0: [c_11 + a_11*b_11 | c_21 + a_21*b_11]
   i = 1: [c_11 + a_21*b_21 | c_21 + a_22*b_21]
*/
c1 = _mm_add_pd(c1, _mm_mul_pd(a,b1));
/* c2 =
   i = 0: [c_12 + a_11*b_12 | c_22 + a_21*b_12]
   i = 1: [c_12 + a_21*b_22 | c_22 + a_22*b_22]
*/
c2 = _mm_add_pd(c2, _mm_mul_pd(a,b2));
}

// store c1,c2 back into C for completion
_mm_store_pd(C+0*lda,c1);
_mm_store_pd(C+1*lda,c2);

// print C
printf("%g,%g\n%g,%g\n",C[0],C[2],C[1],C[3]);
return 0;
```

C vs. Python

240x!

4x!

N	Python [GFLOPS]	C [GFLOPS]	AVX [GFLOPS]
32	0.0054	1.30	4.56
160	0.0055	1.30	5.47
480	0.0054	1.32	5.27
960	0.0053	0.91	3.64

Theoretical Intel i7-5557U performance is ~25 GFLOPS

$$3.1\text{GHz} \times 2 \text{ instructions/cycle} \times 4 \text{ mults/inst} = 24.8\text{GFLOPS}$$