



UC Berkeley  
Teaching Professor  
Dan Garcia

# CS61C

## Great Ideas in Computer Architecture (a.k.a. Machine Structures)



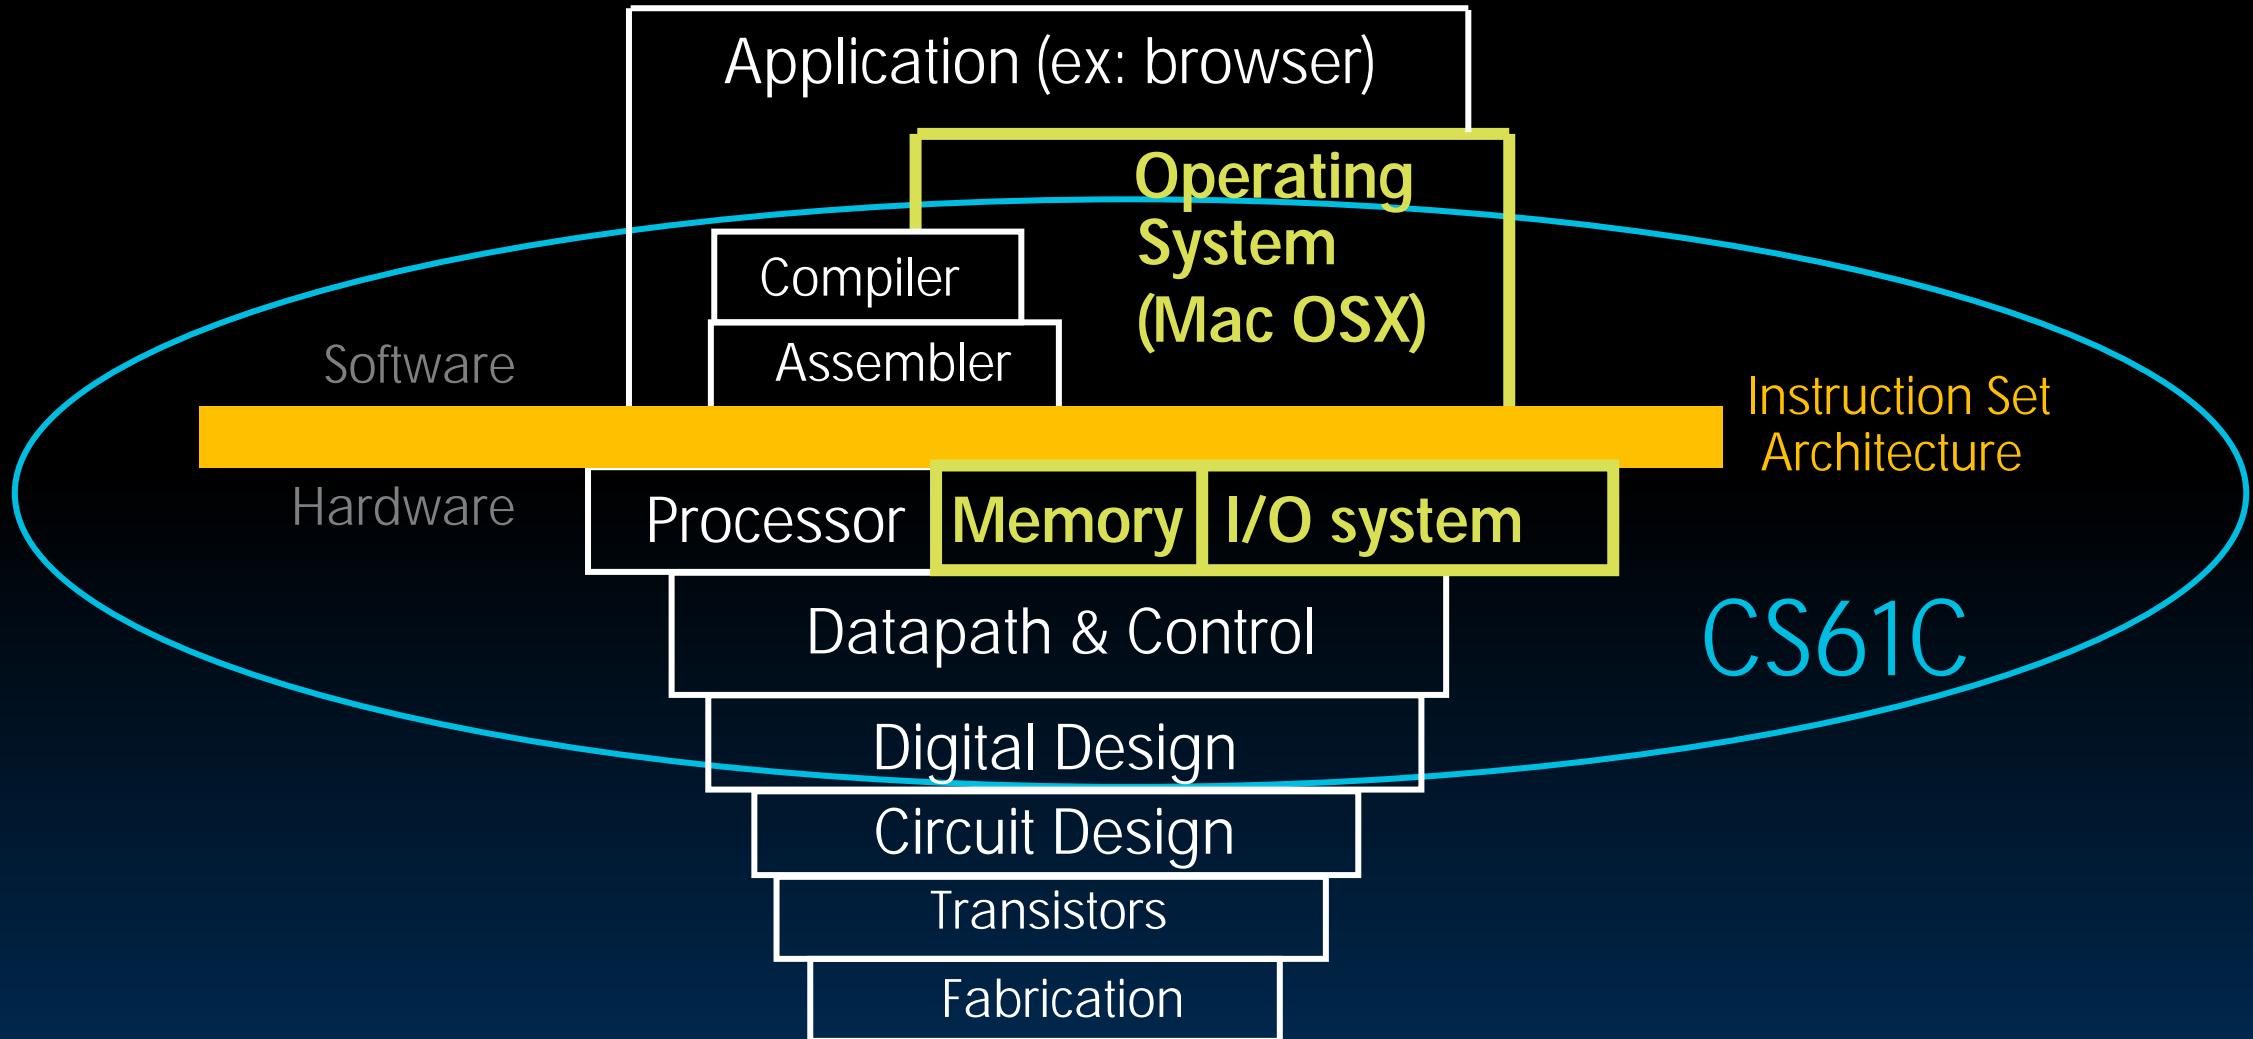
UC Berkeley  
Teaching Professor  
Lisa Yan

## Virtual Memory I

# The Computer

- The Computer
- OS Basics: Context Switching
- Physical Memory and Disk Storage
- Virtual Memory and Virtual Addresses
- Paged Memory
- Page Table Details I

# Old-school Machine Structures



# New-school Machine Structures

## Parallel Programs

Parallel Requests

Assigned to computer  
e.g., Search "Cats"

Parallel Threads

Assigned to core e.g., Lookup, Ads

Parallel Instructions

>1 instruction @ one time  
e.g., 5 pipelined instructions

Parallel Data

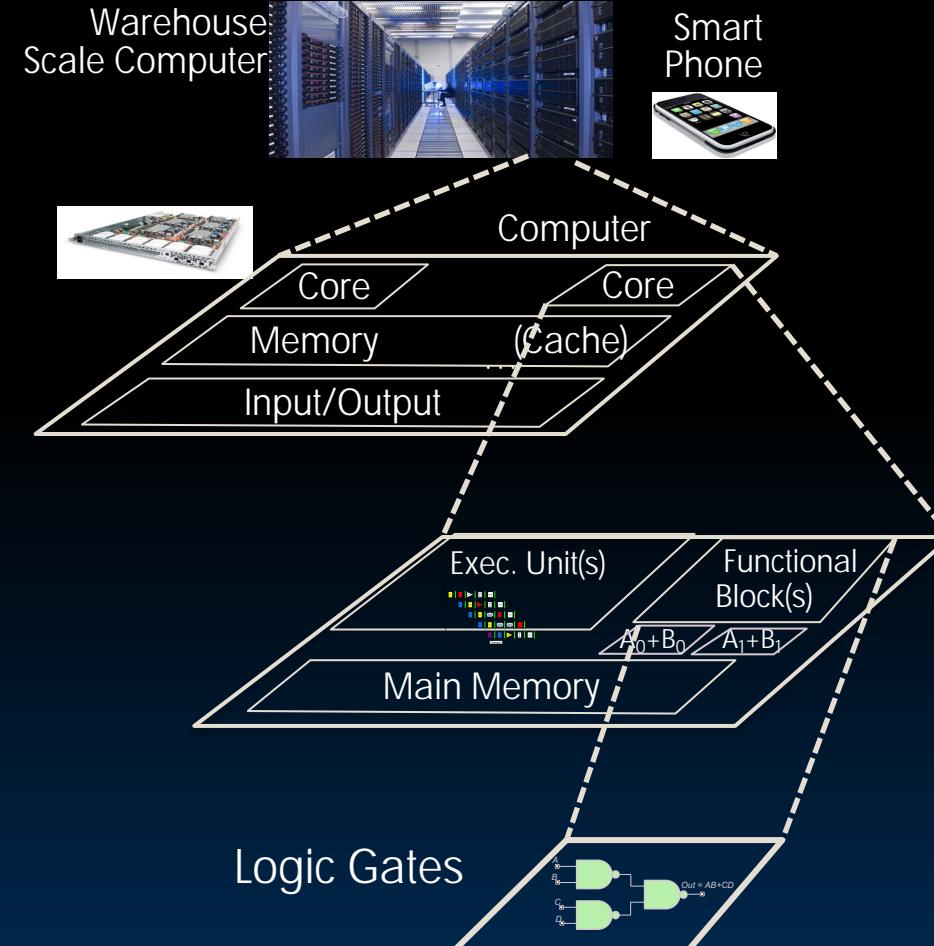
>1 data item @ one time  
e.g., Add of 4 pairs of words

Hardware descriptions

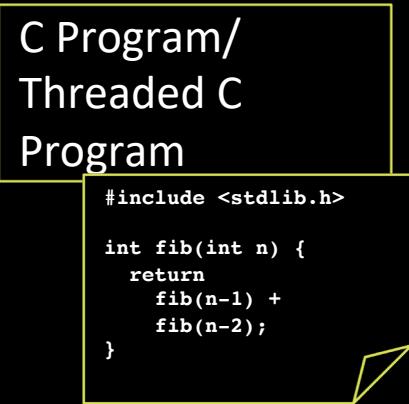
All gates work in parallel at same time

Harness Parallelism &  
Achieve High Performance

Hardware

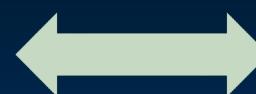
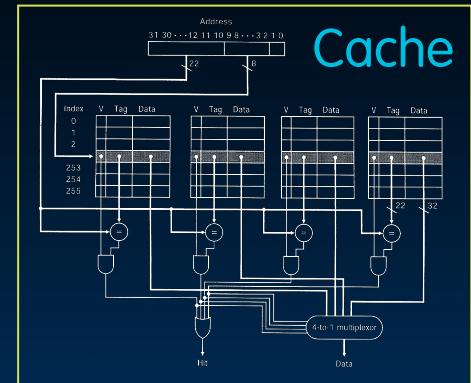
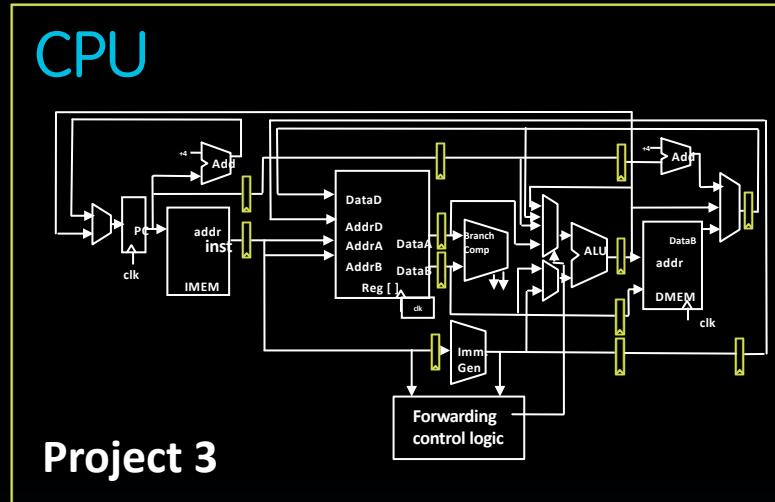


# CS61C so far...



```
.foo
lw t0, 4(a0)
addi t1, t0, 3
beq t1, t2, foo
nop
```

Project 2

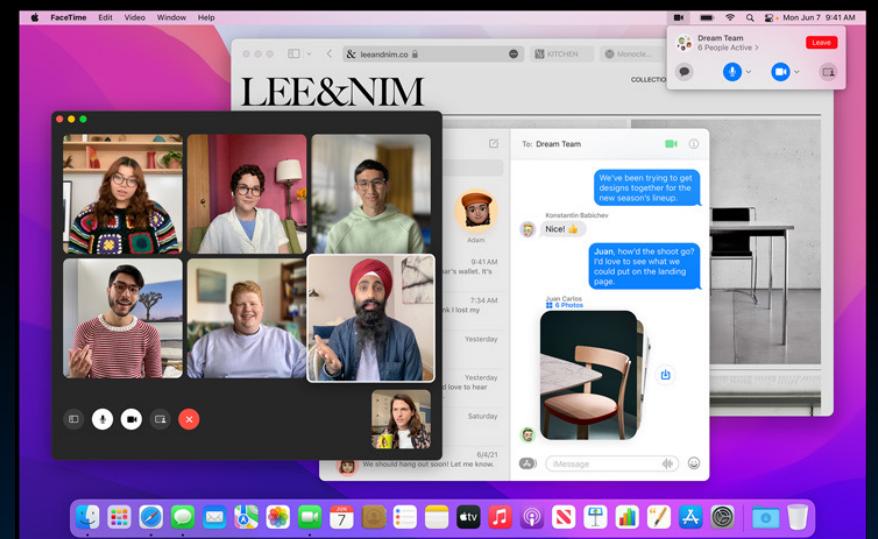
Project 1/  
Project 4

# But Wait...

- When I run Venus, it only executes one program, then stops.
- My laptop looks and acts differently!



1. Multiple *I/O devices* (input-output)



2. Multiple programs running simultaneously "on" a software program called the *Operating System (OS)*

Garcia, Yan

# Raspberry Pi (\$35) and I/O ports

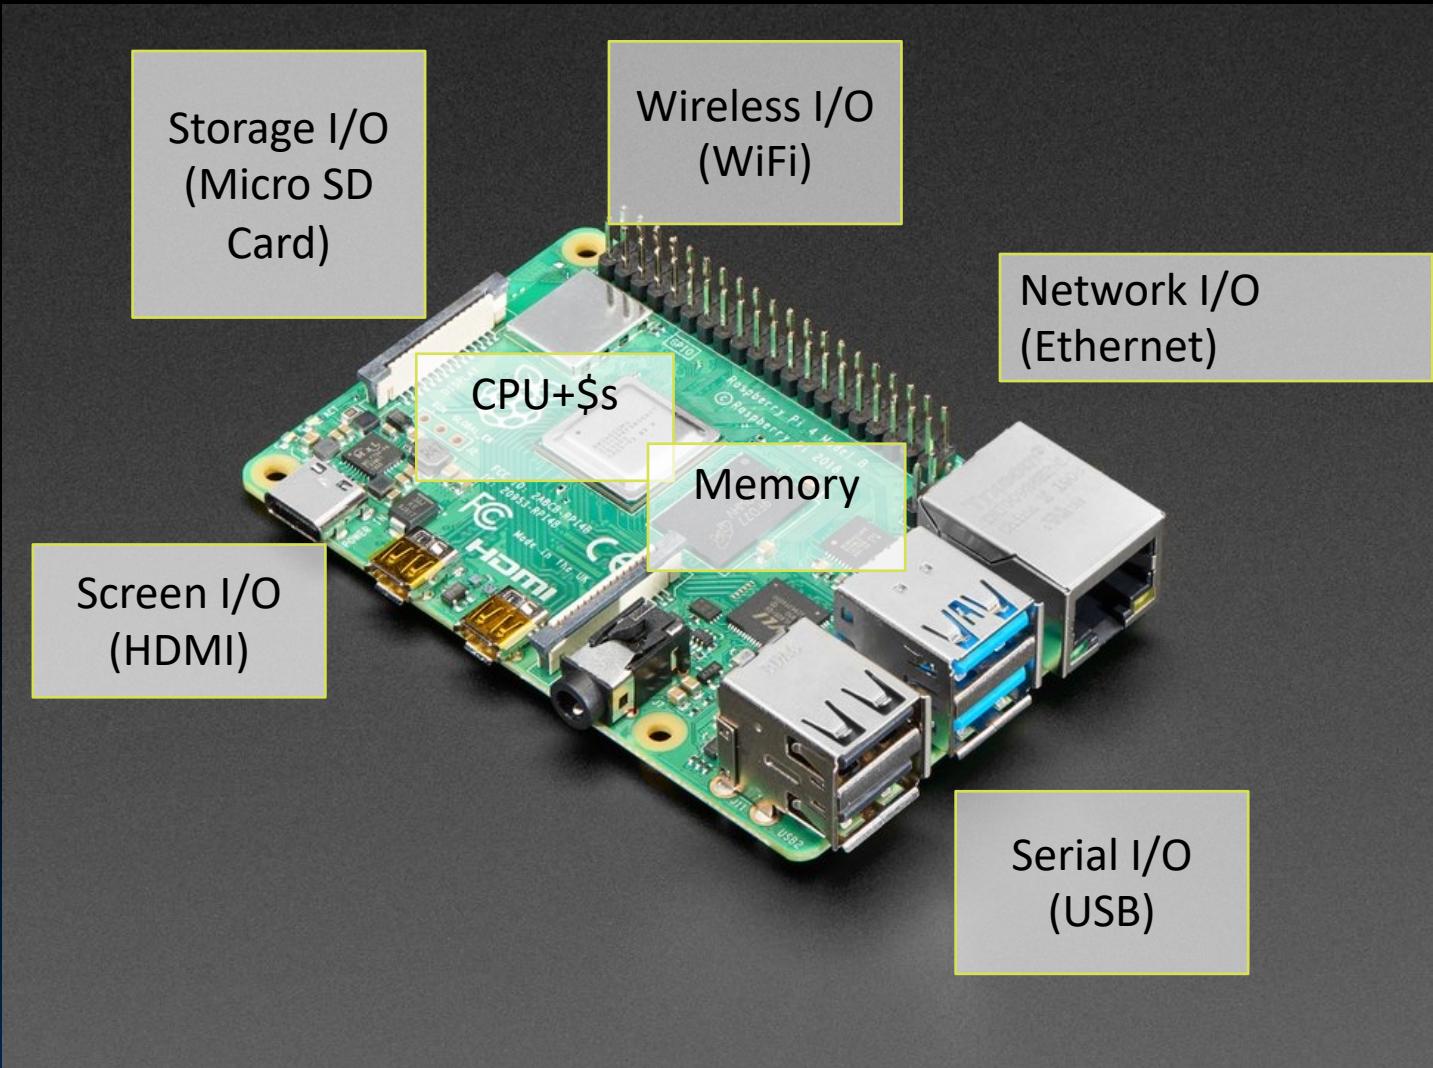
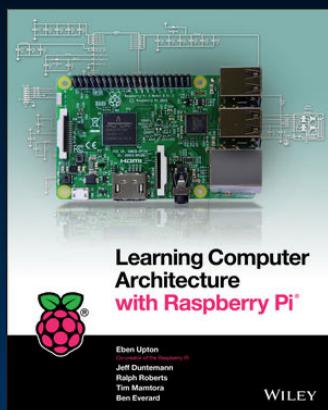
The Raspberry Pi is a low-cost computer.

- Motherboard with CPU, I/O, caches, etc., soldered on
  - ("\$" stands for cache)

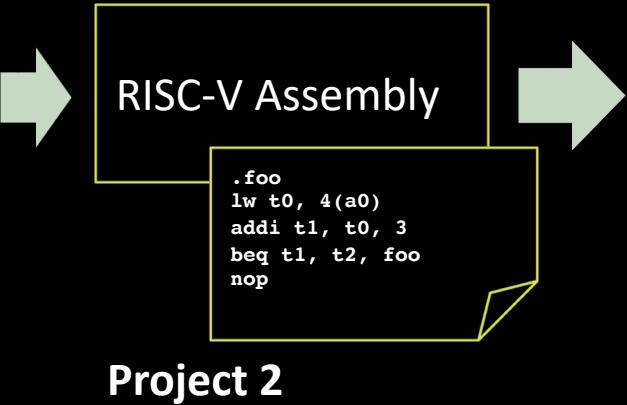
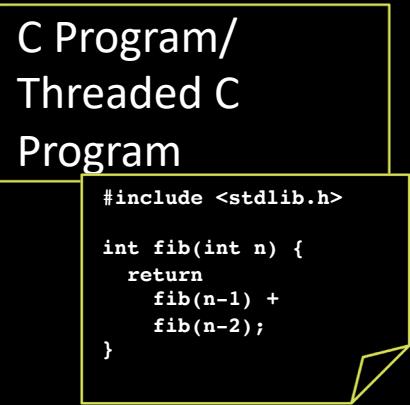
## CS61C w/Raspberry Pi?

- Free to download with UC Berkeley login!

<https://onlinelibrary.wiley.com/doi/book/10.1002/9781119415534>

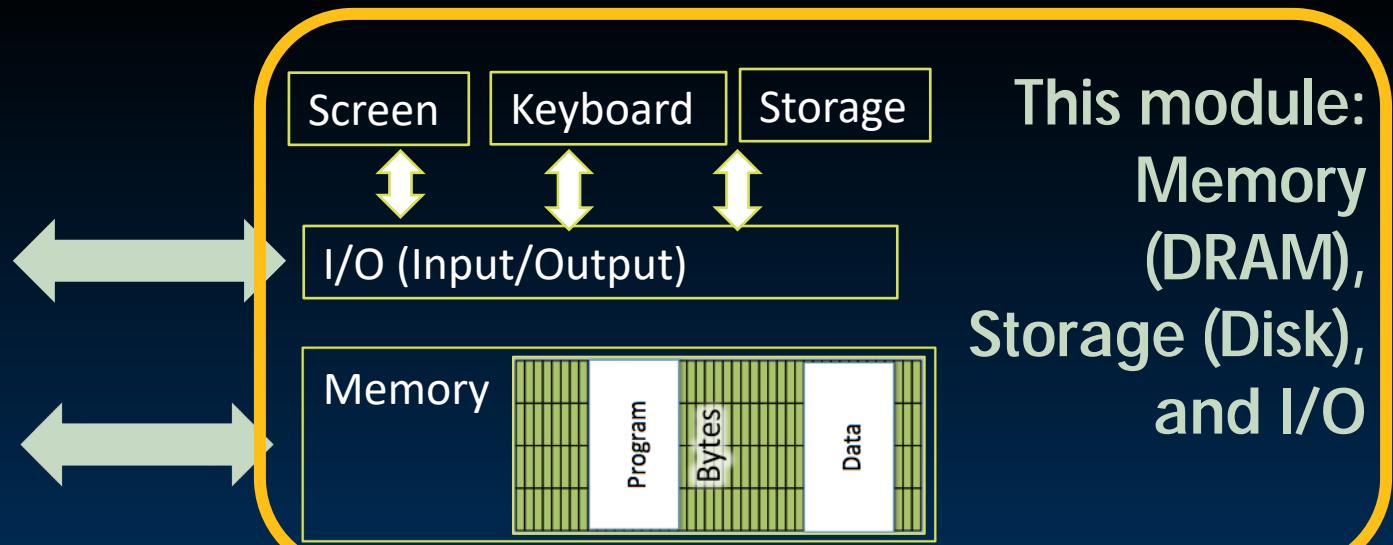
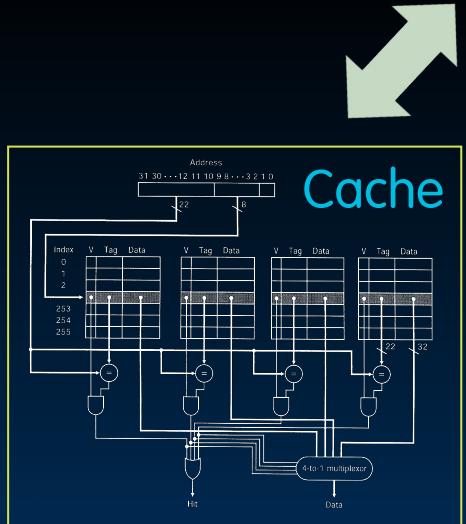
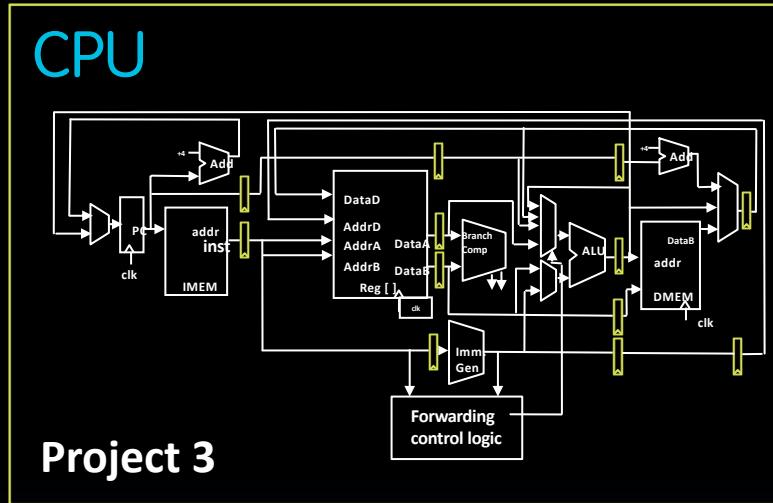


# The Next Step of CS61C



Project 2

Project 1/  
Project 4

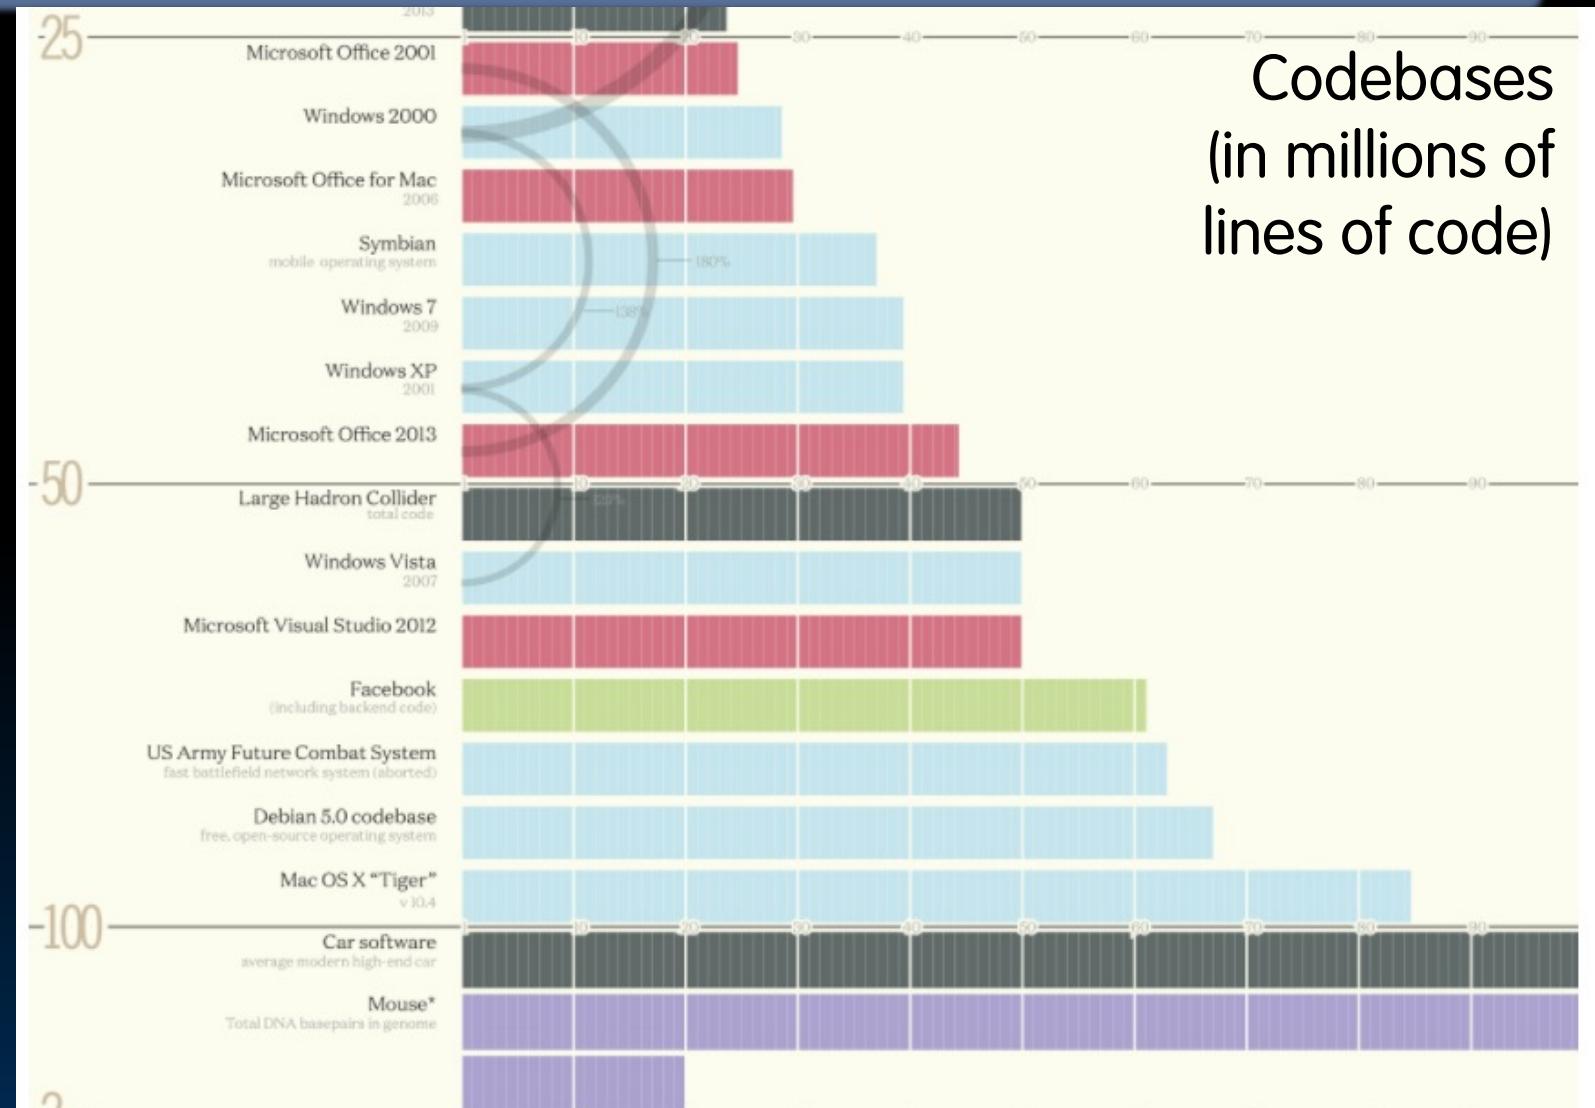


# OS Basics: Context Switching

- The Computer
- OS Basics: Context Switching
- Physical Memory and Disk Storage
- Virtual Memory and Virtual Addresses
- Paged Memory
- Page Table Details I

# The OS is “Just Software”

- The biggest piece of software on your machine?
- How many lines of code? These are guesstimates:



CC BY-NC 3.0

David McCandless © 2015

<http://www.informationisbeautiful.net/visualizations/million-lines-of-code/>

# Linux Kernel Over Time

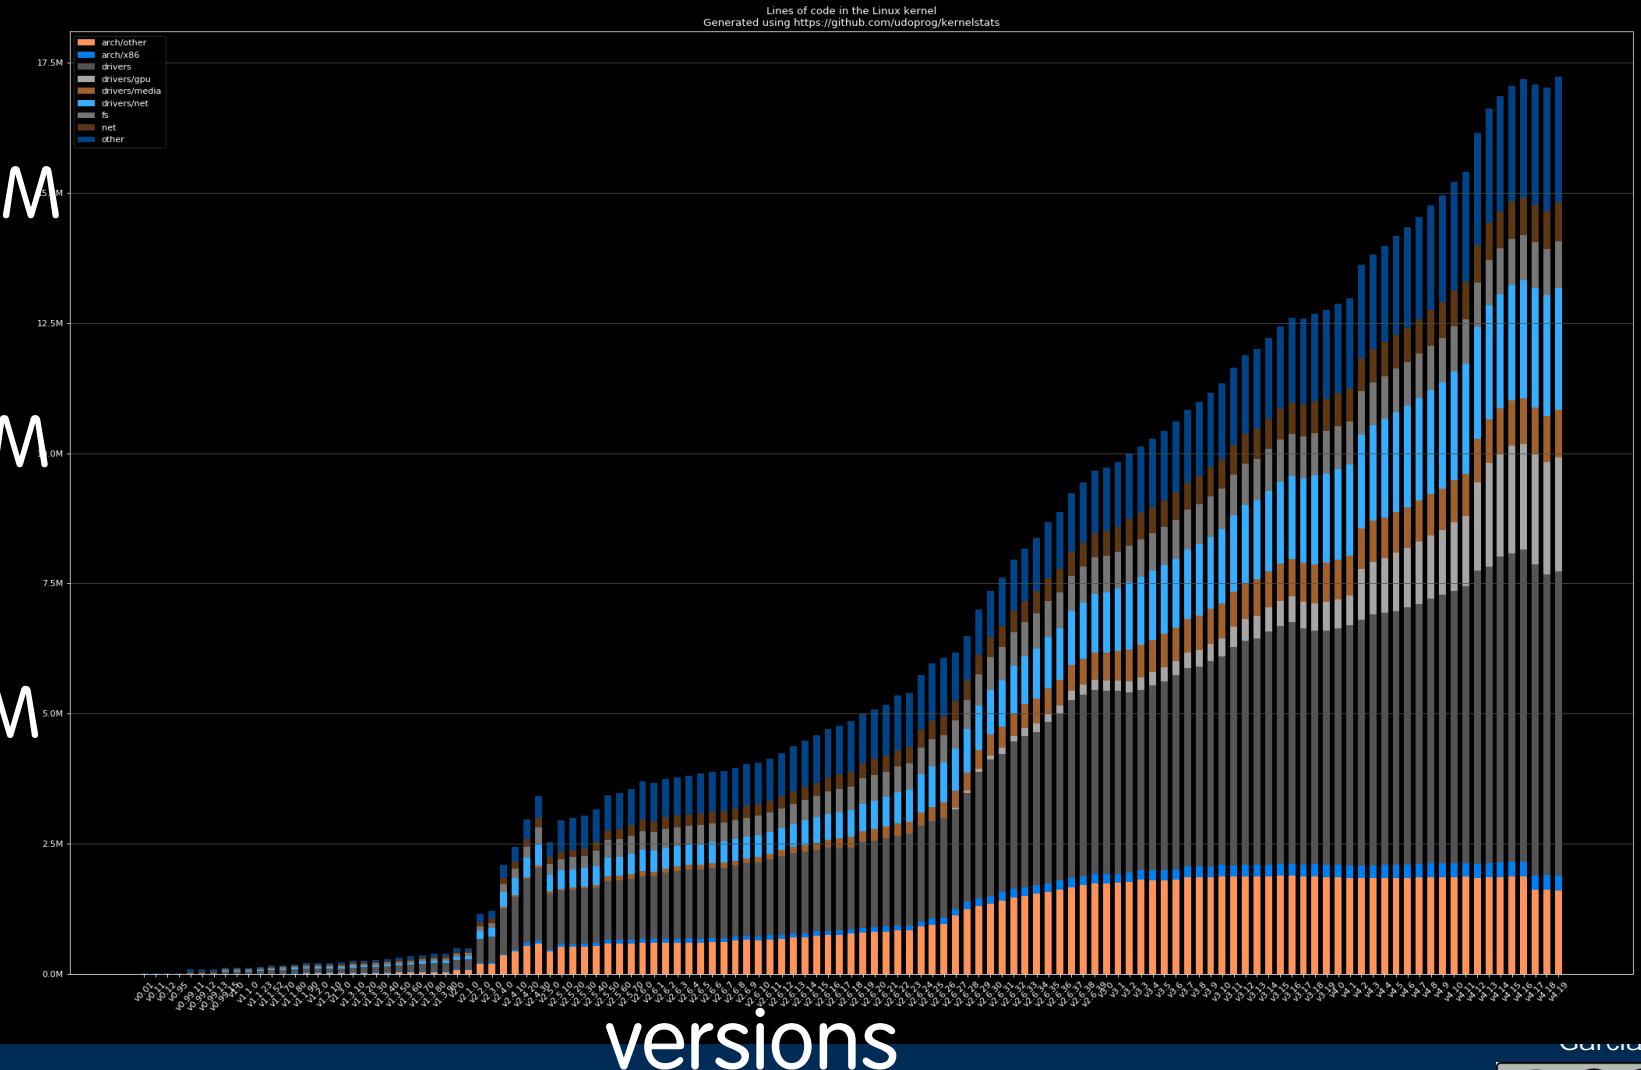
In CS61C,  
OS  $\approx$  kernel.

- Other components, e.g., User Interface, not covered.

Take CS162  
for more!

Millions of lines

Lines of code in Linux kernel



# What does the core of the OS do?

- **OS is the (first) thing that runs when computer starts.**
  - Starts services (100+).
    - File system, Network stack (Ethernet, WiFi, Bluetooth, ...), TTY (keyboard), etc.
- **Provides interaction with the outside world:**
  - Finds and controls all devices in the machine in a general way:
    - Relies on hardware specific "device drivers"
- **Loads, runs and manages programs:**
  - *Isolation*: Each program runs (i.e., appears to run) in its own little world.
  - Resource-sharing: Multiple programs share the same resources:
    - Memory
    - I/O devices: disk, keyboard, display, network, etc.
  - Time-sharing: Processor (CPU) runs multiple processes.

# What does the core of the OS do?

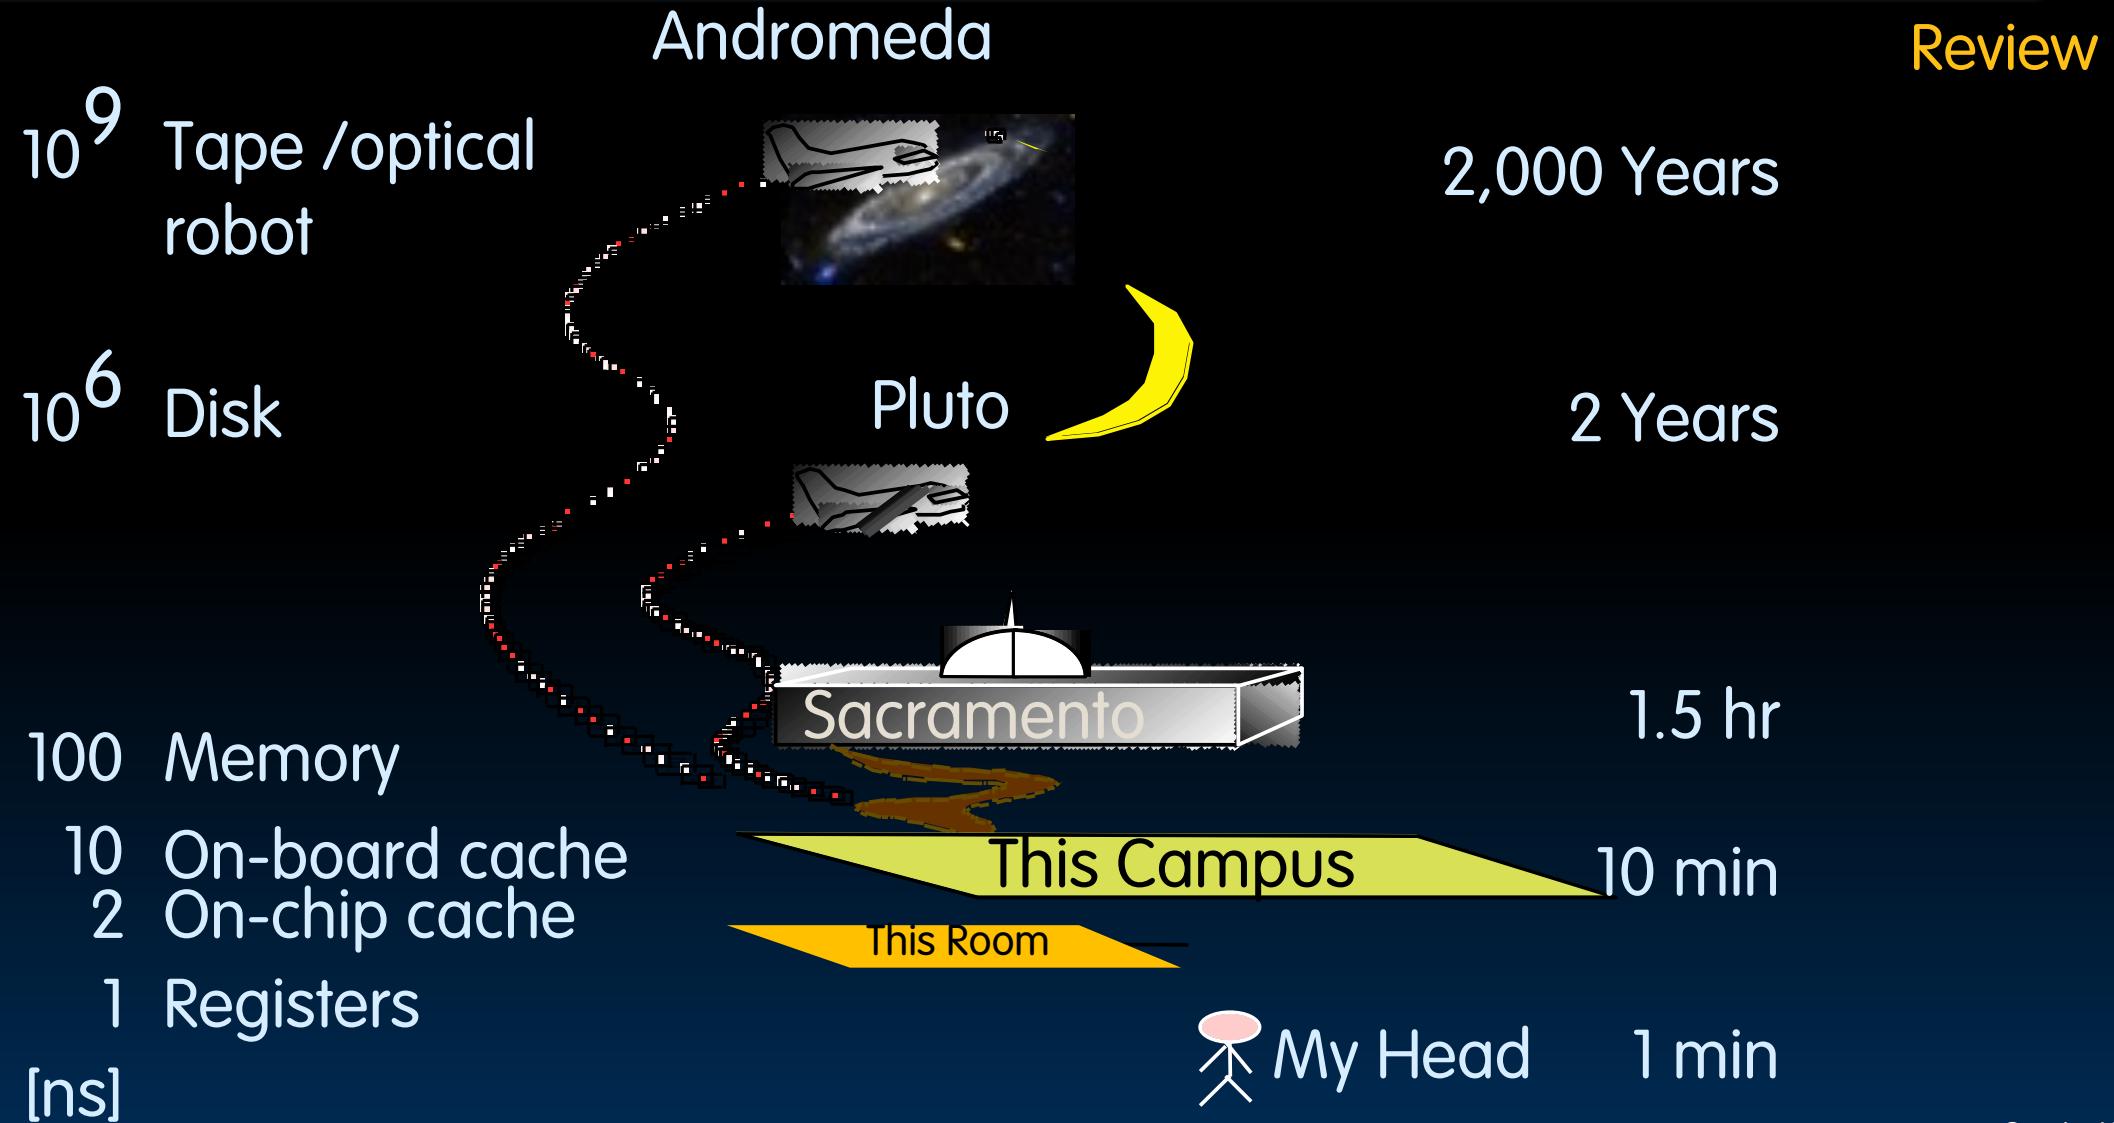
- OS is the (first) thing that runs when computer starts.
  - Starts services (100+).
    - File system, Network stack (Ethernet, WiFi, Bluetooth, ...), TTY (keyboard), etc.
- Provides interaction with the outside world:
  - Finds and controls all devices in the machine in a general way:
    - Relies on hardware specific "device drivers"
- Loads, runs and manages programs:
  - *Isolation*: Each program runs (i.e., appears to run) in its own little world.
  - Resource-sharing: Multiple programs share the same resources:
    - Memory
    - I/O devices: disk, keyboard, display, network, etc.
  - Time-sharing: Processor (CPU) runs multiple processes.



Over the next few lectures, we will discuss this entire slide. For now, we focus on resource-sharing *memory*.

- The OS manages multiprogramming, which is running multiple applications (processes) “simultaneously” on one CPU.
  - (vs. multiprocessing: running processes simultaneously on different CPUs. The OS also manages this.)
- This is achieved via OS context switches, i.e., switches between processes very quickly (on the human time scale):
  - Save current process state (program counter, registers, etc.)
  - Load next process state to execute next instruction on CPU
  - Do *not* switch out data between main memory and disk! Too costly...

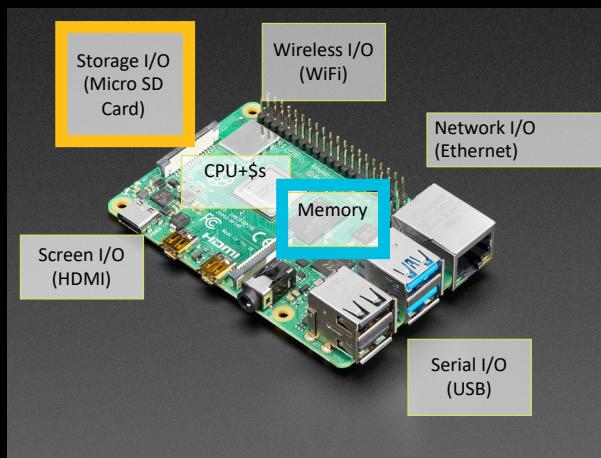
# Storage Latency Analogy: How Far Away Is the Data?



# Physical Memory and Disk Storage

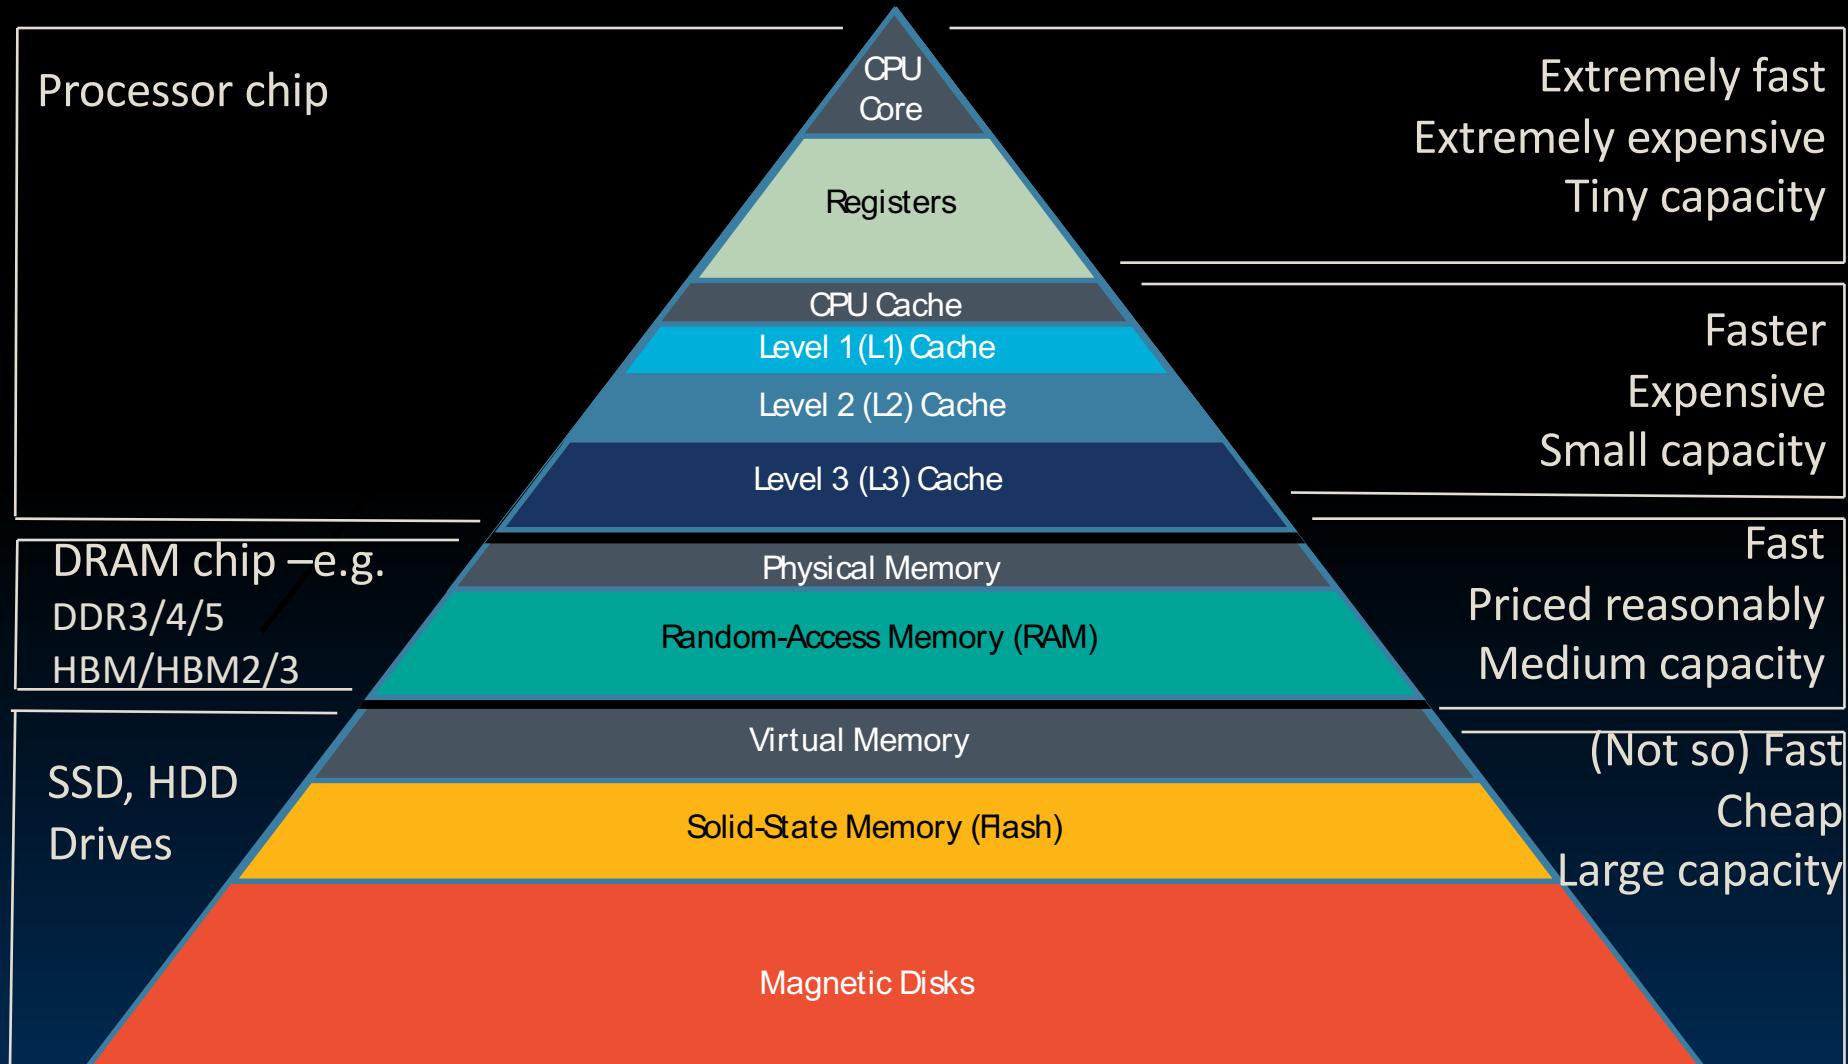
- The Computer
- OS Basics: Context Switching
- Physical Memory and Disk Storage
- Virtual Memory and Virtual Addresses
- Paged Memory
- Page Table Details I

# Main Memory and Secondary Memory



Main memory

Secondary  
Memory



Garcia, Yan

# Main Memory is DRAM

- **Dynamic Random Access Memory:**

- Latency to access first word: ~10ns (~30-40 processor cycles), each successive (0.5ns – 1ns)
- Each access brings 64 bits, supports ‘bursts’
- \$3/GiB

- **Data is impermanent:**

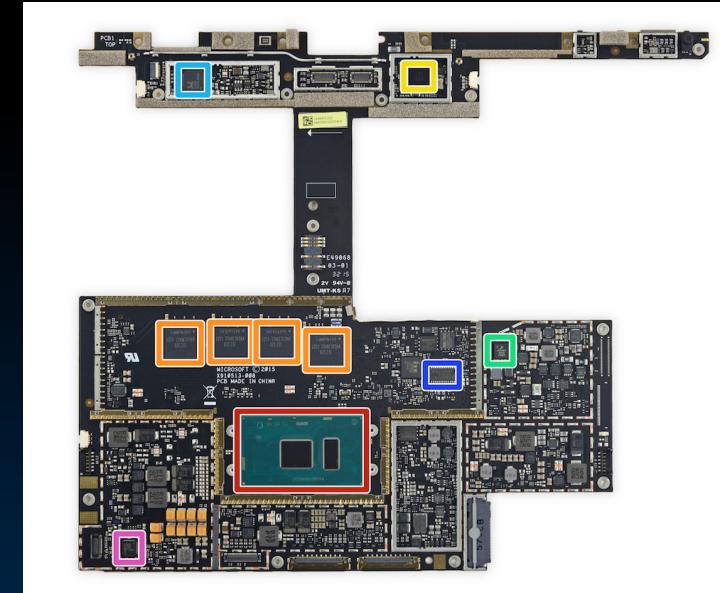
- *Dynamic*: capacitors store bits, so needs periodic refresh to maintain charge
- *Volatile*: when power is removed, loses data.

- **Contrast with SRAM (for caches):**

- Static (no capacitors) but still volatile
- Faster (0.5 ns)/more expensive/lower density



Desktop/server DIMM



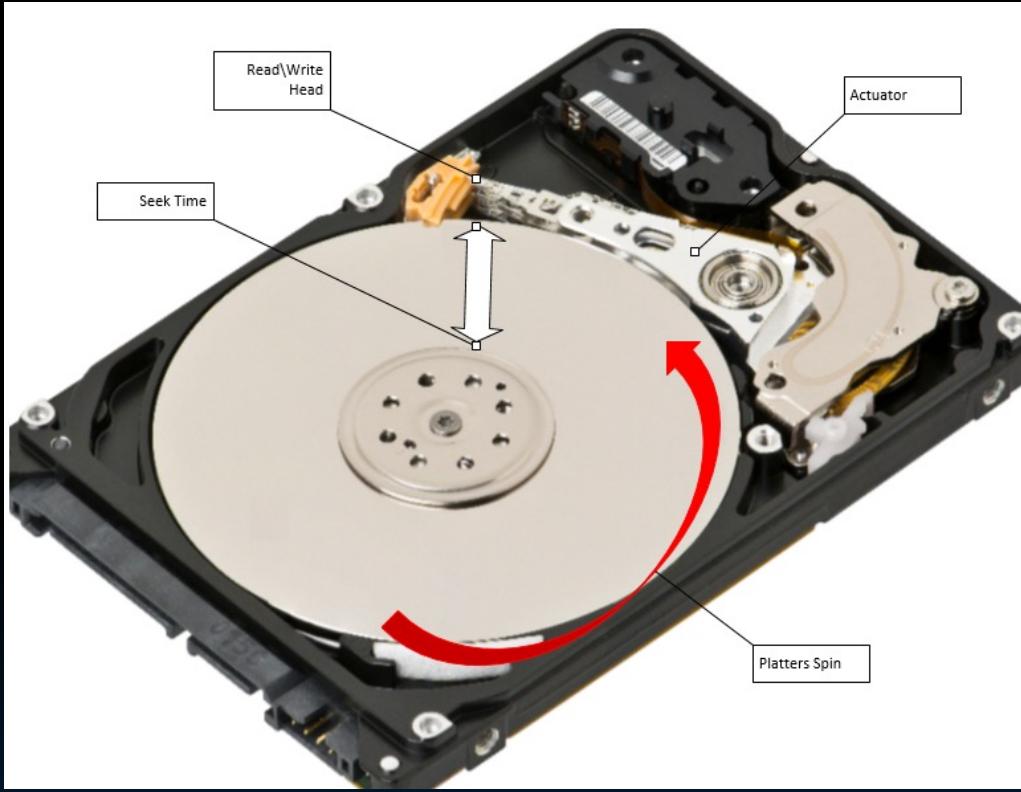
MS Surface Book

# Storage / "Disk" / Secondary Memory

- Attached as a peripheral I/O device. Non-volatile
- **Solid-State Drive (SSD)**
  - Access:  $40\text{-}100\mu\text{s}$   
( $\sim 100\text{k}$  proc. cycles)
  - $\$0.05\text{-}0.5/\text{GB}$
  - Usually flash memory
- **Hard Disk Drive (HDD)**
  - Access:  $<5\text{-}10\text{ms}$   
( $10\text{-}20\text{M}$  proc. cycles)
  - $\$0.01\text{-}0.1/\text{GB}$
  - Mechanical



# Aside ... How do HDDs work?



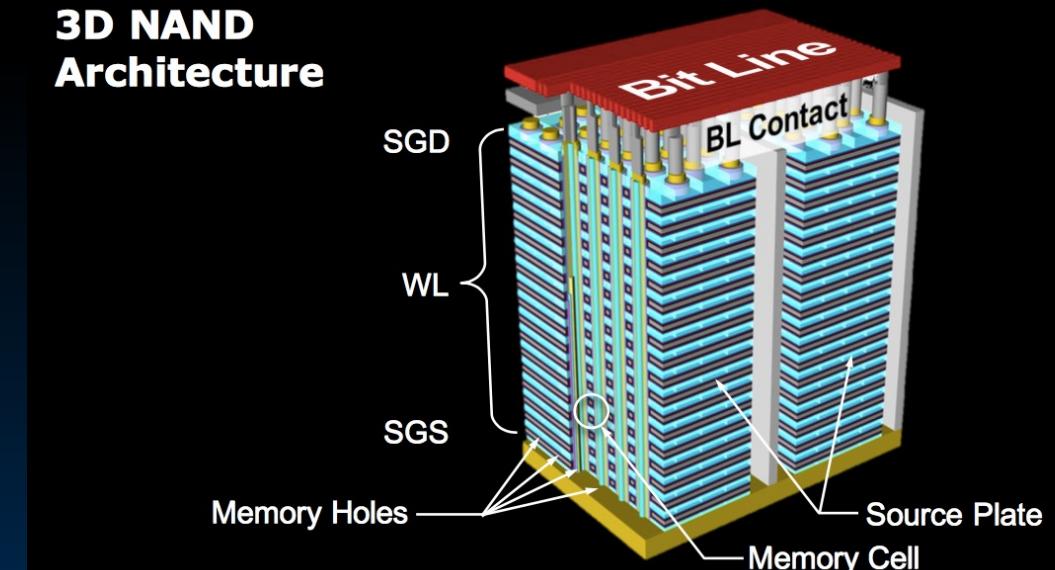
- 10,000 rpm (revolutions per minute)
  - 6 ms per revolution
  - Average random access time: 3 ms (~ $10^7$  processor cycles)

# Nick Parlante's explanation: [anford.edu/people/nick/how-hard-drive-works/](http://anford.edu/people/nick/how-hard-drive-works/)



# Aside 2...What about SSDs?

- **Made with transistors (nothing mechanical/rotating)**
- **Operates like a “Ginormous” register file.**
  - Furthermore, does not “forget” when power is off (non-volatile)
- **Fast access to all locations, regardless of address**
  - Still much slower than register, DRAM.
  - Read/write blocks, not bytes.
  - Potential reliability issues...
- **Some unusual requirements:**
  - Can’t erase single bits
    - only entire blocks



Flash memory is generally a 3D array of bit cells (up to 256 layers!)

# Virtual Memory and Virtual Addresses

- The Computer
- OS Basics: Context Switching
- Physical Memory and Disk Storage
- Virtual Memory and Virtual Addresses
- Paged Memory
- Page Table Details I

# The Case for Virtual Memory (1/2)

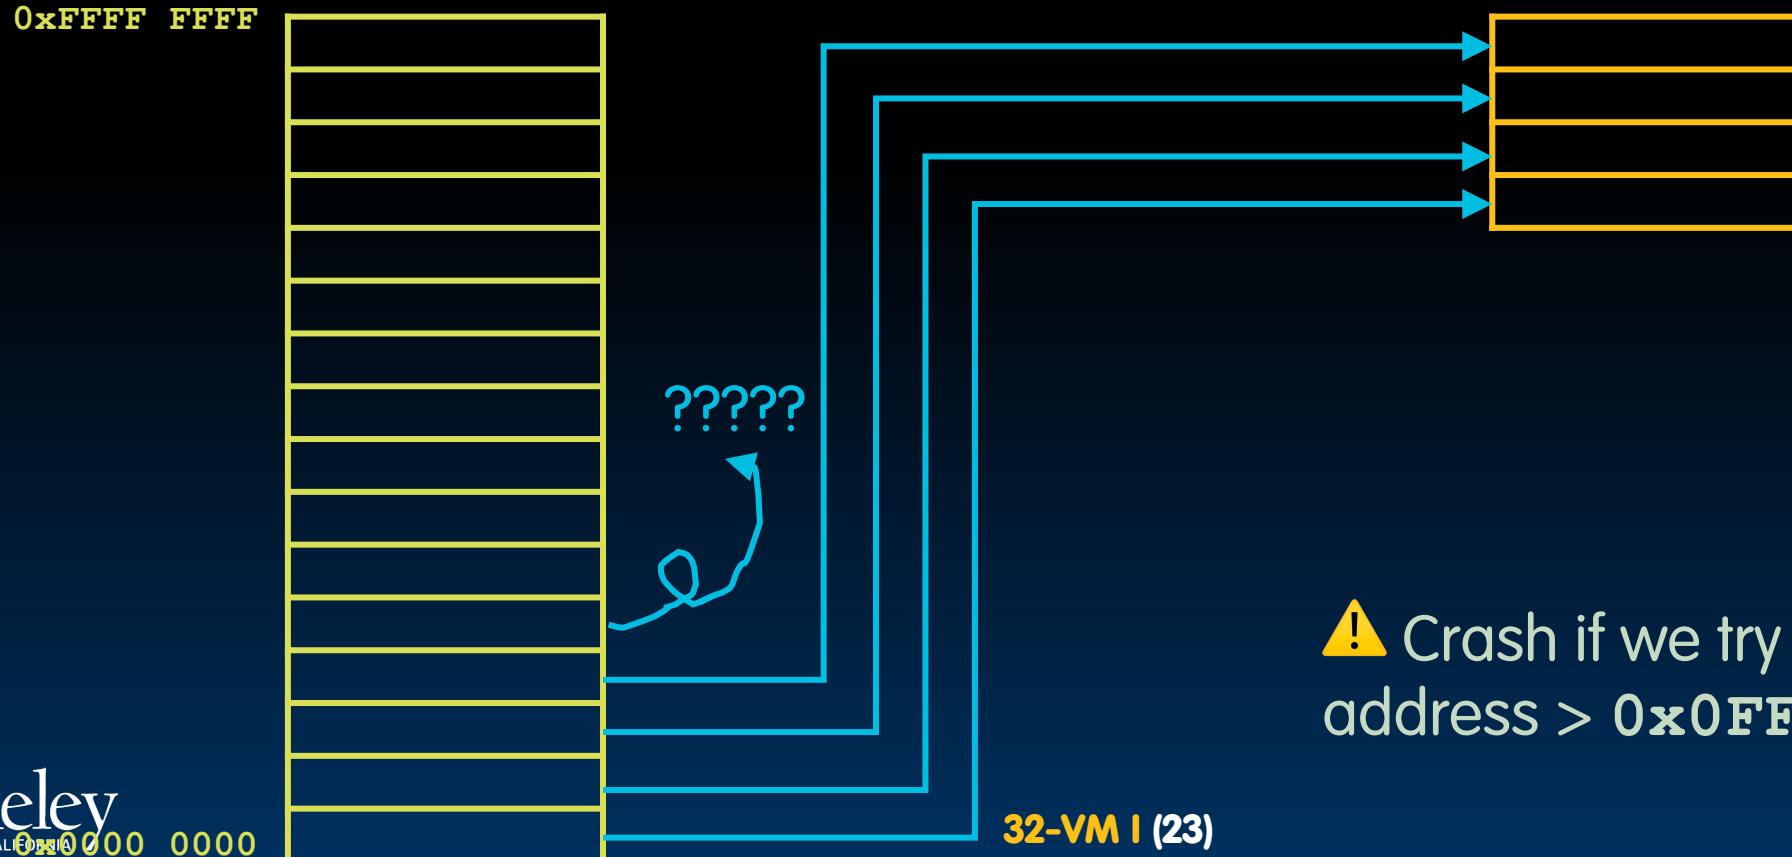
# 1. What if main memory is smaller than the program address space?

RV32I provides a 32-bit address space.

→  $2^{32}$  B = 4 GiB addressable memory

# Suppose RAM is 1GiB.

→  $2^{30}$  B addressable memory.



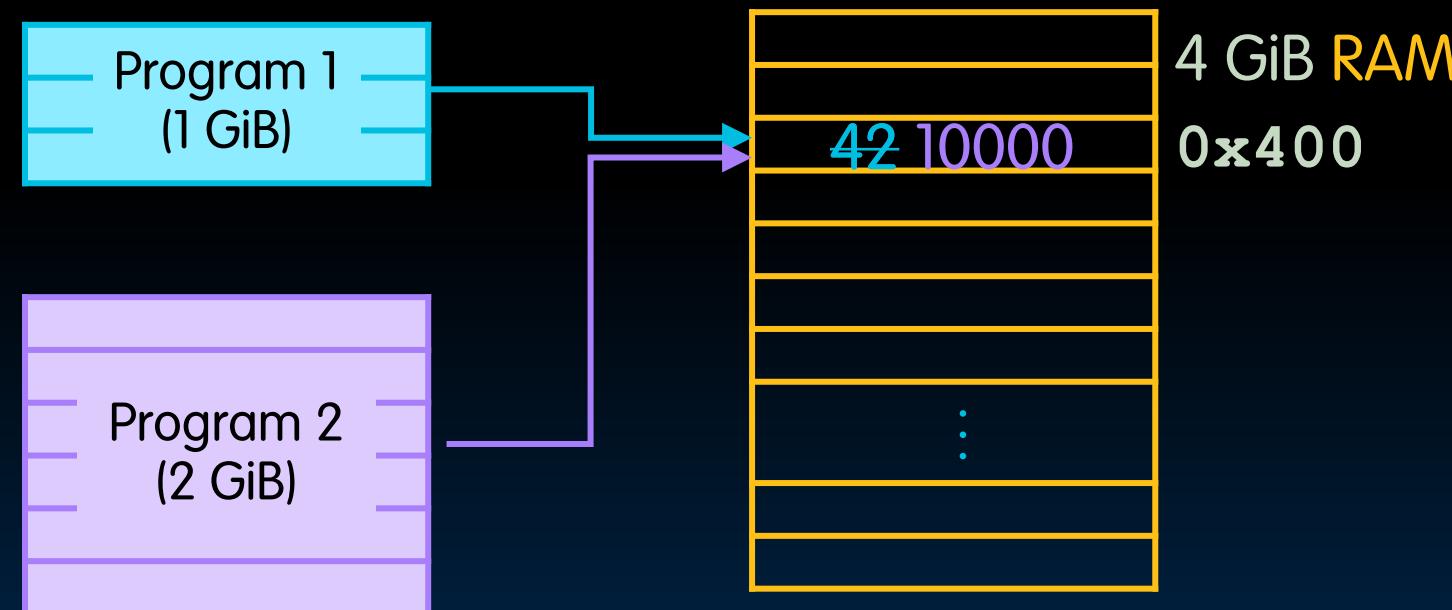
**⚠ Crash if we try to access an address > 0x0FFF - FFFF!**

# The Case for Virtual Memory (2/2)

1. What if main memory is smaller than the program address space?
2. What if two programs access the same memory address?

Program 1 stores your bank account balance at address **0x400**

Program 2 stores your video game score at address **0x400**



⚠ If all processes can access any 32-bit memory address, they can corrupt/crash others.  
▫ Need *protection* (i.e., isolation) between processes.

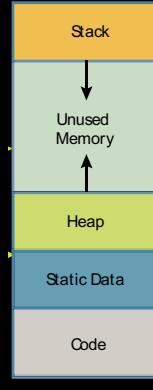
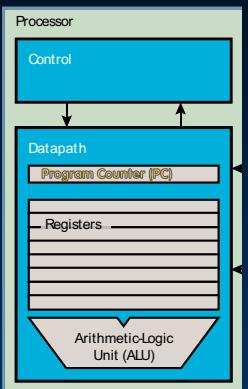
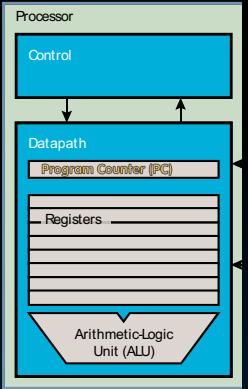
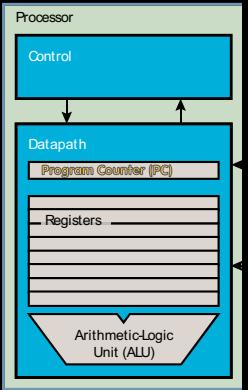
- Virtual memory is the next level in the memory hierarchy:
  - Give each process the *illusion* of a full memory address space that it has completely for itself.
  - Under the hood: working set of *pages* reside in main memory; other pages are in disk.
- Benefits:
  - *Demand paging* provides the ability to run programs larger than the primary memory (DRAM).
  - OS can share memory and *protect* programs from each other.
  - Hides differences between machine configurations.
- Today, more important for *protection* than space management.
  - (Historically, virtual memory predates caches.)

# Virtual vs. Physical Addresses

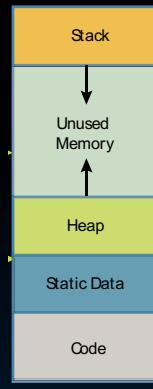
- Address Space: set of addresses for all available memory locations.
  - Now, two kinds of memory addresses!
- Virtual Address Space
  - Set of addresses that the user program knows about
- Physical Address Space
  - Set of addresses that map to actual physical locations in memory
  - Hidden from user applications
- For each program, a memory manager maps (translates) between these two address spaces.

# Virtual Address Space Illusion

Different processes run simultaneously

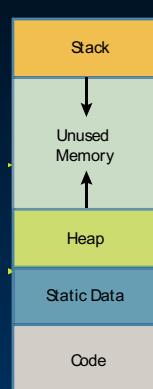


ffff ffff<sub>hex</sub>



0000 0000<sub>hex</sub>

ffff ffff<sub>hex</sub>

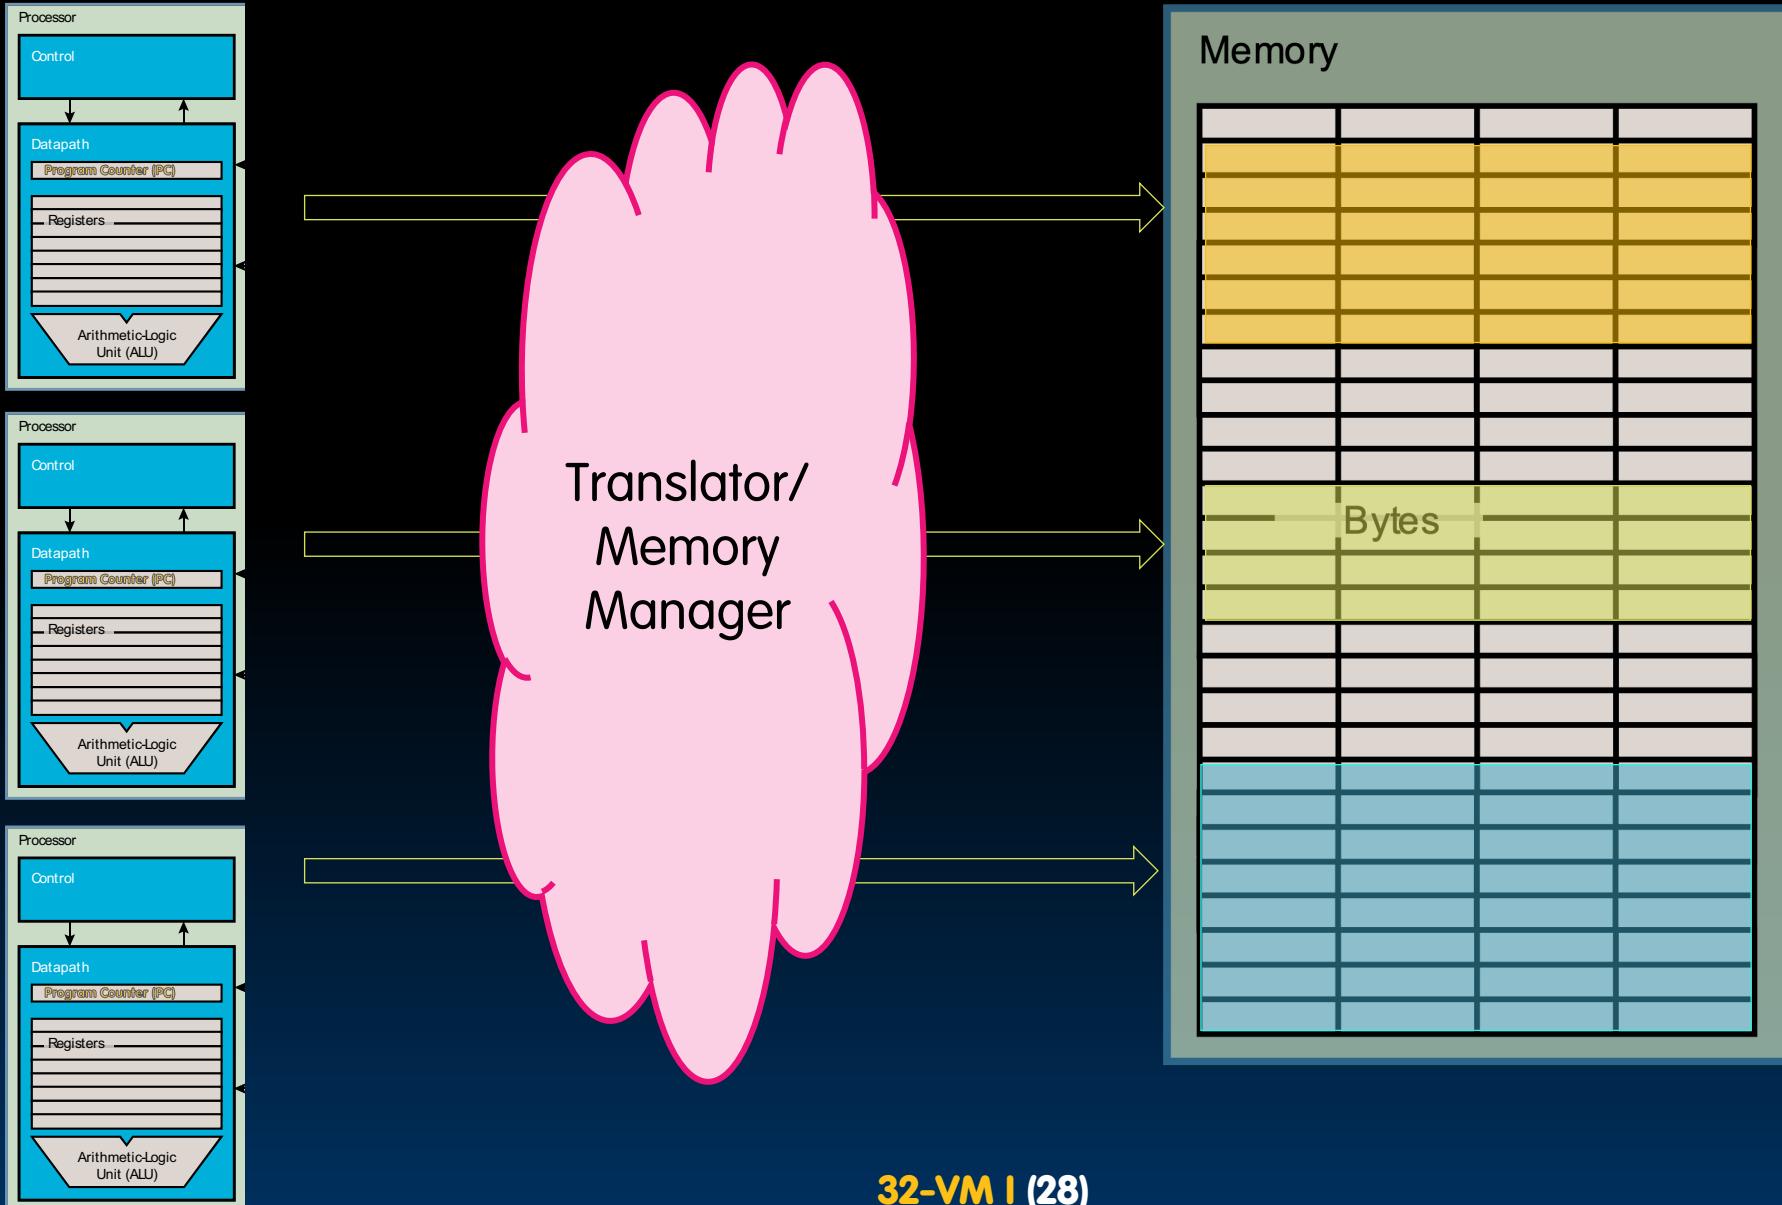


0000 0000<sub>hex</sub>

ffff ffff<sub>hex</sub>

**Processes use virtual addresses.**  
Many processes, all using same (conflicting) addresses

# Conceptual Memory Manager in OS



Memory  
uses  
physical  
addresses.



# CS61C Hive Machines

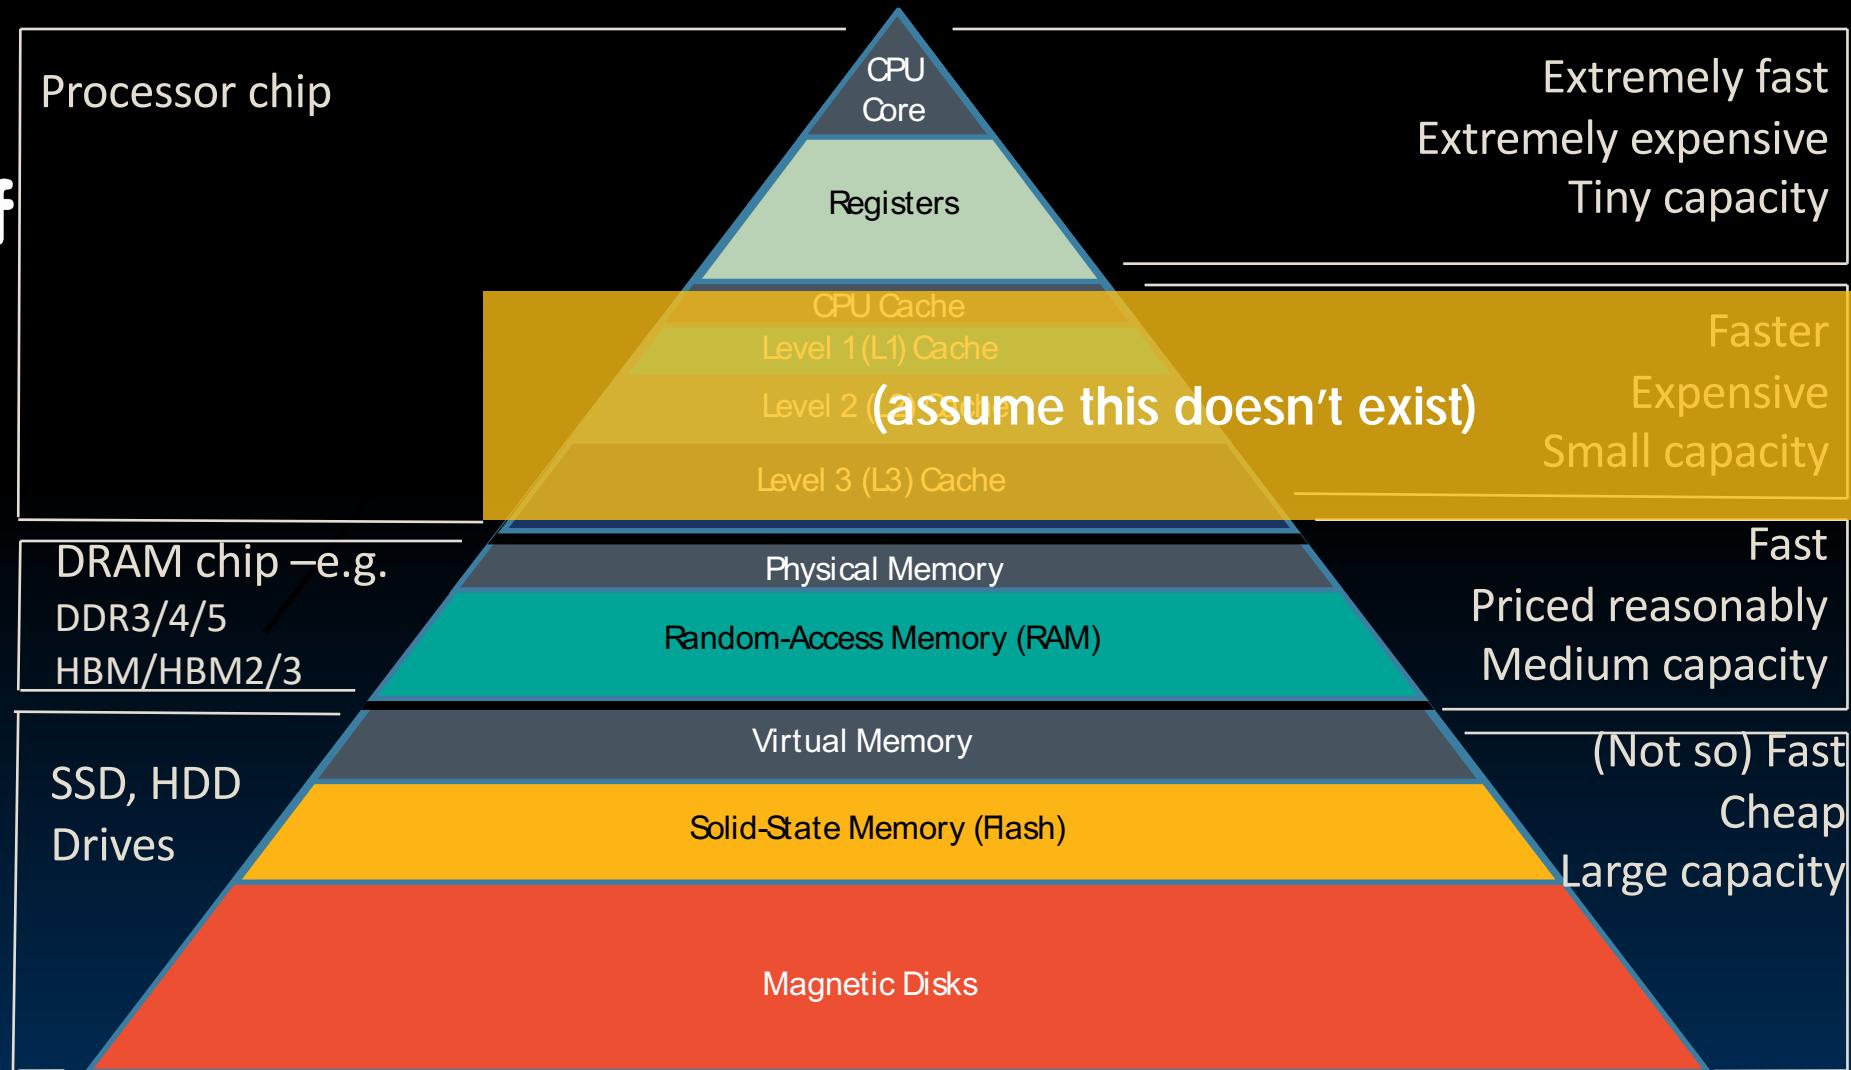
```
(00:18:45 Mon Nov 07 2022 cs61c-tab@hive2 Linux x86_64)
[~ $ cat /proc/cpuinfo
processor      : 0
vendor_id     : GenuineIntel
cpu family    : 6
model         : 60
model name   : Intel(R) Core(TM) i7-4770 CPU @ 3.40GHz
stepping       : 3
microcode     : 0x28
cpu MHz       : 3693.327
cache size    : 8192 KB
physical id   : 0
siblings       : 8
core id        : 0
cpu cores     : 4
apicid         : 0
initial apicid: 0
fpu            : yes
fpu_exception  : yes
cpuid level   : 13
wp             : yes
flags          : fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov
pat pse36 clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdt
scp lm constant_tsc arch_perfmon pebs bts rep_good nopl xtopology nonstop_tsc cp
uid aperfmpf pnpi pclmulqdq dtes64 monitor ds_cpl vmx smx est tm2 ssse3 sdbg fm
a cx16 xtpr pdcm pcid sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_timer aes x
save avx f16c rdrand lahf_lm abm cpuid_fault epb invpcid_single pt1 ssbd ibrs ib
pb stibp tpr_shadow vnmi flexpriority ept vpid ept_ad fsgsbase tsc_adjust bmi1 a
vx2 smep bmi2 erms invpcid xsaveopt dtherm ida arat pln pts md_clear flush_l1d
bugs          : cpu_meltdown spectre_v1 spectre_v2 spec_store_bypass l1tf mds
swaps          : itlb_multihit srbds
bogomips      : 6784.38
clflush size   : 64
cache_alignment: 64
address sizes  : 39 bits physical, 48 bits virtual
power management:
```



# ⚠ Assume Caches Don't Exist For Now ⚠

Virtual Memory  
is much easier  
to understand if  
we assume no  
caches.

- We'll reintroduce caches along with *Translation Lookaside Buffers (TLBs)* soon.



# Paged Memory

- The Computer
- OS Basics: Context Switching
- Physical Memory and Disk Storage
- Virtual Memory and Virtual Addresses
- Paged Memory
- Page Table Details I

# OS Virtual Memory Management Responsibilities

- 1. Map virtual addresses to physical addresses.**
- 2. Use both memory and disk.**
  - Give illusion of larger memory by storing some content on disk.
  - Disk is usually much larger and slower than DRAM.
- 3. Protection:**
  - Isolate memory between processes.

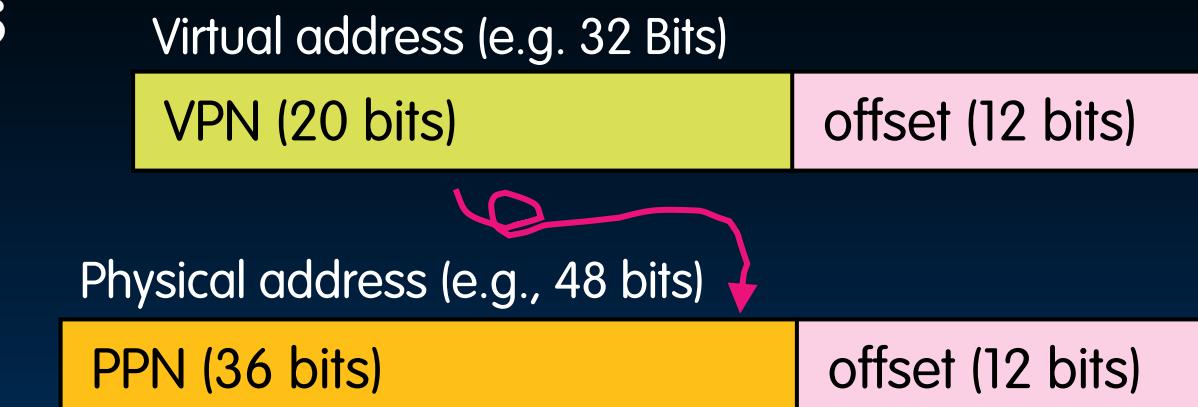
- The concept of “paged memory” dominates:
  - Physical memory (DRAM) is broken into *pages*.
  - A disk access loads an entire page into memory.
  - Typical page size: 4 KiB+ (on modern OSs)
    - Need 12 bits of *page offset* to address all 4 KiB bytes.

- If virtual and physical pages are the same size, then memory translation maps

Virtual Page Number (VPN)

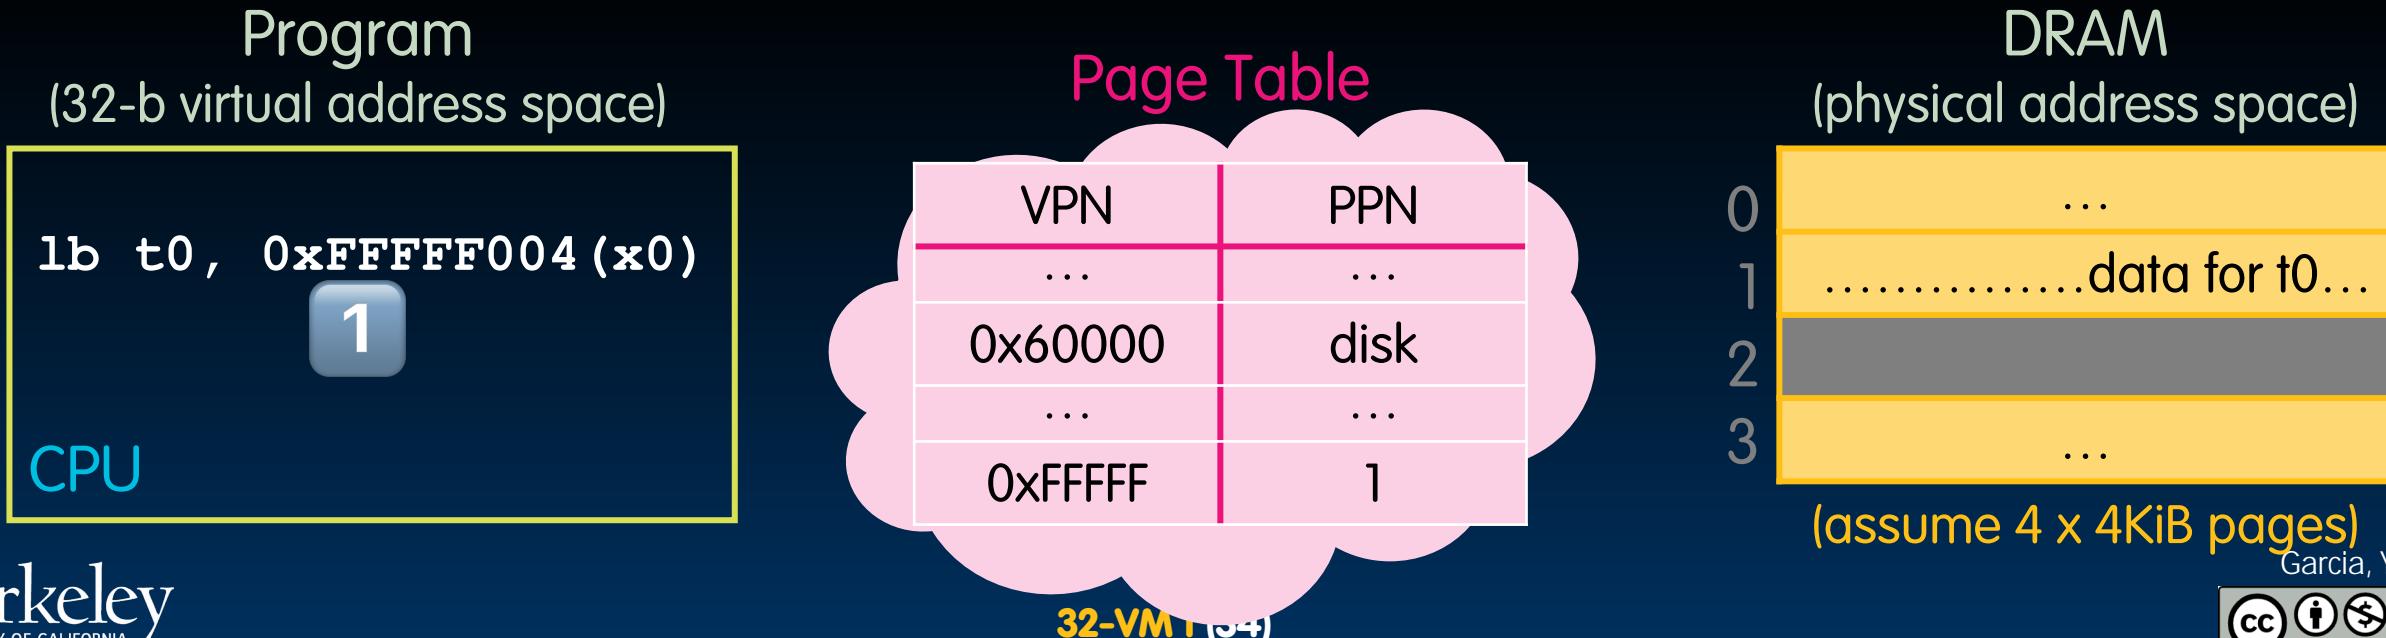
to a

Physical Page Number (PPN).



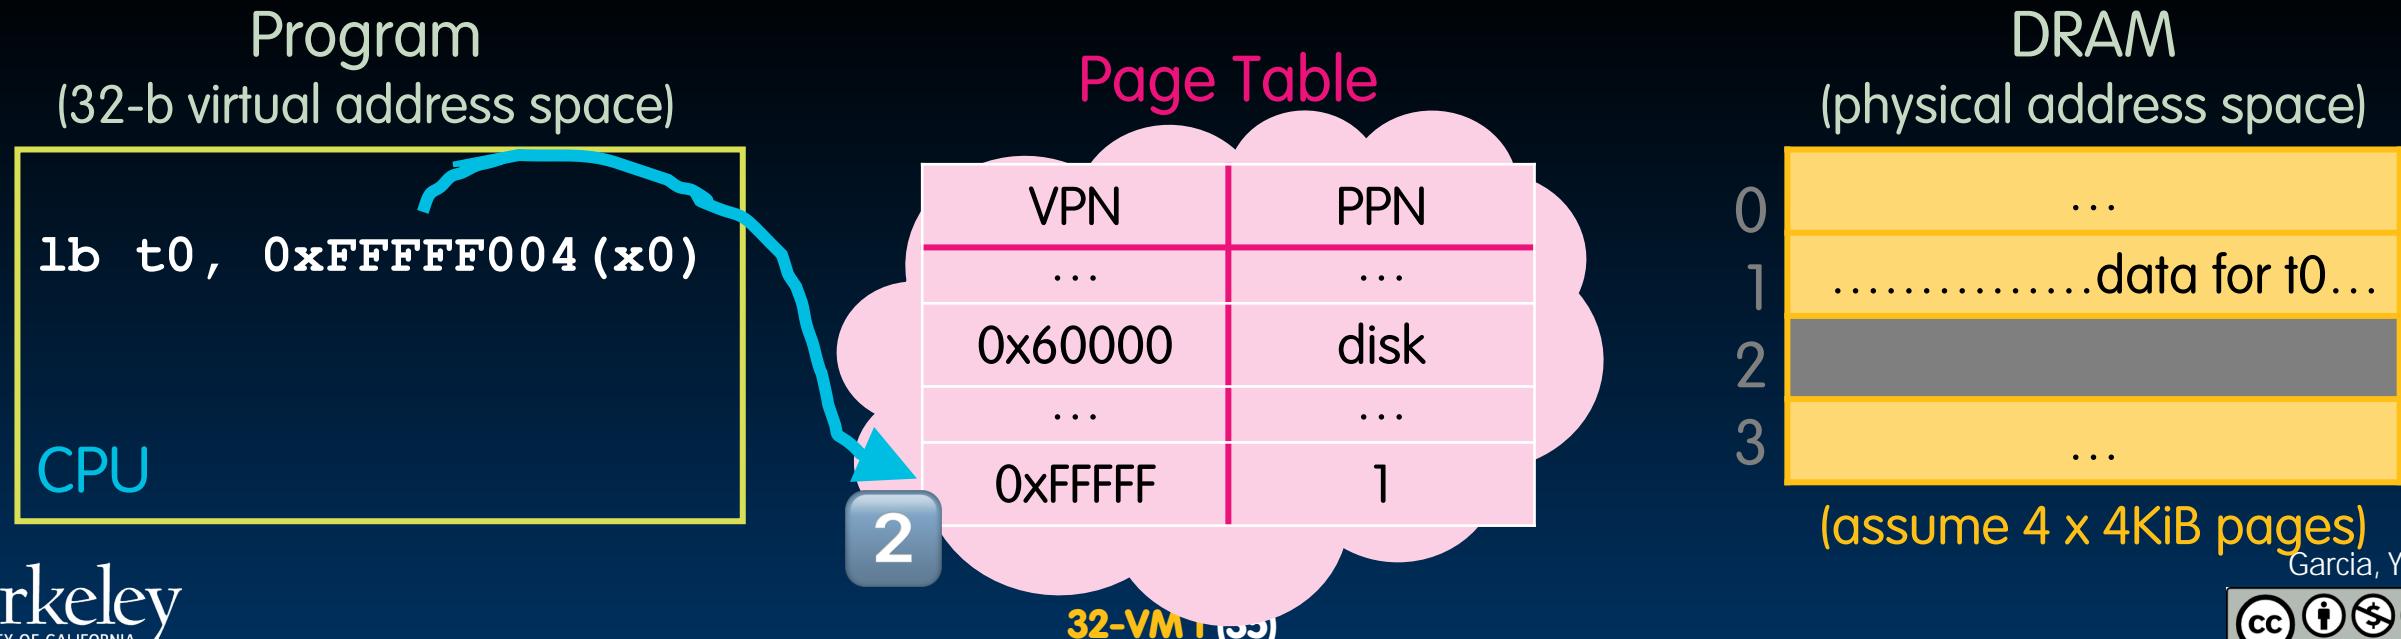
# Translation: How a Program Accesses Memory

1. Program executes a load specifying a virtual address (VA).
2. Computer translates VA to the physical address (PA) in memory.
  - Extract virtual page number (VPN) from VA, e.g., top 20 bits
  - Look up physical page number (PPN) in page table
  - Construct PA: physical page number + offset (from virtual address)
3. If the physical page is not in memory, then OS loads it in from disk.
4. The OS reads memory at the PA and returns the data to the program.



# Translation: How a Program Accesses Memory

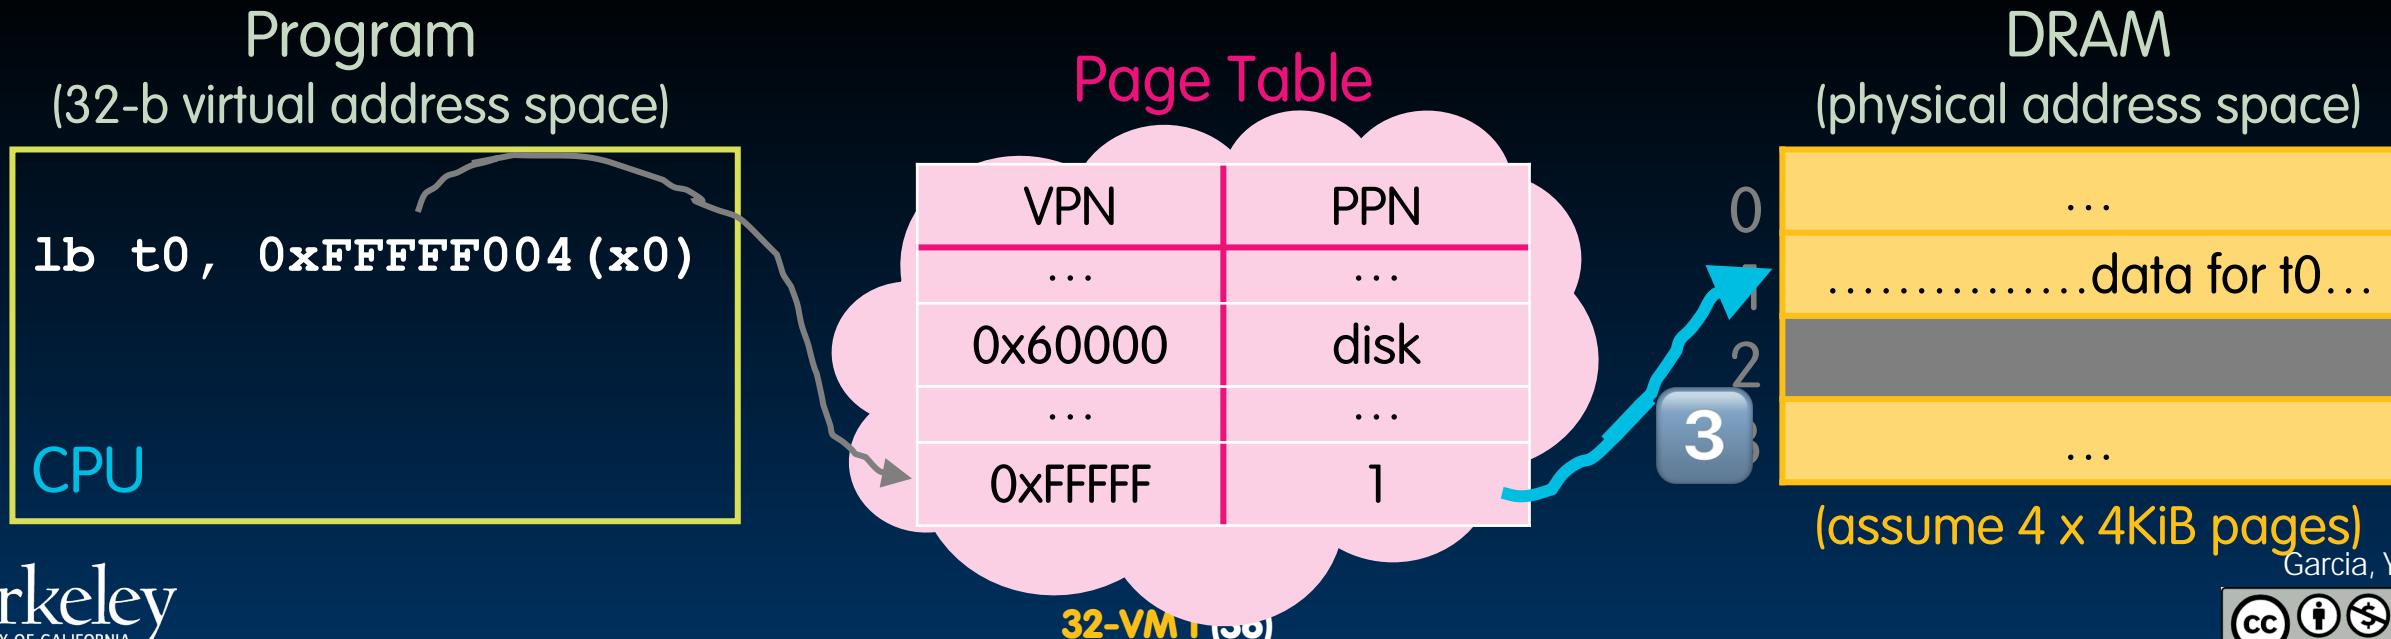
1. Program executes a load specifying a virtual address (VA).
2. Computer translates VA to the physical address (PA) in memory.
  - Extract virtual page number (VPN) from VA (e.g., top 20 bits if page size 4KiB =  $2^{12}$  B)
  - Look up physical page number (PPN) in page table
  - Construct PA: physical page number + offset (from virtual address)
3. If the physical page is not in memory, then OS loads it in from disk.
4. The OS reads memory at the PA and returns the data to the program.



Garcia, Yan

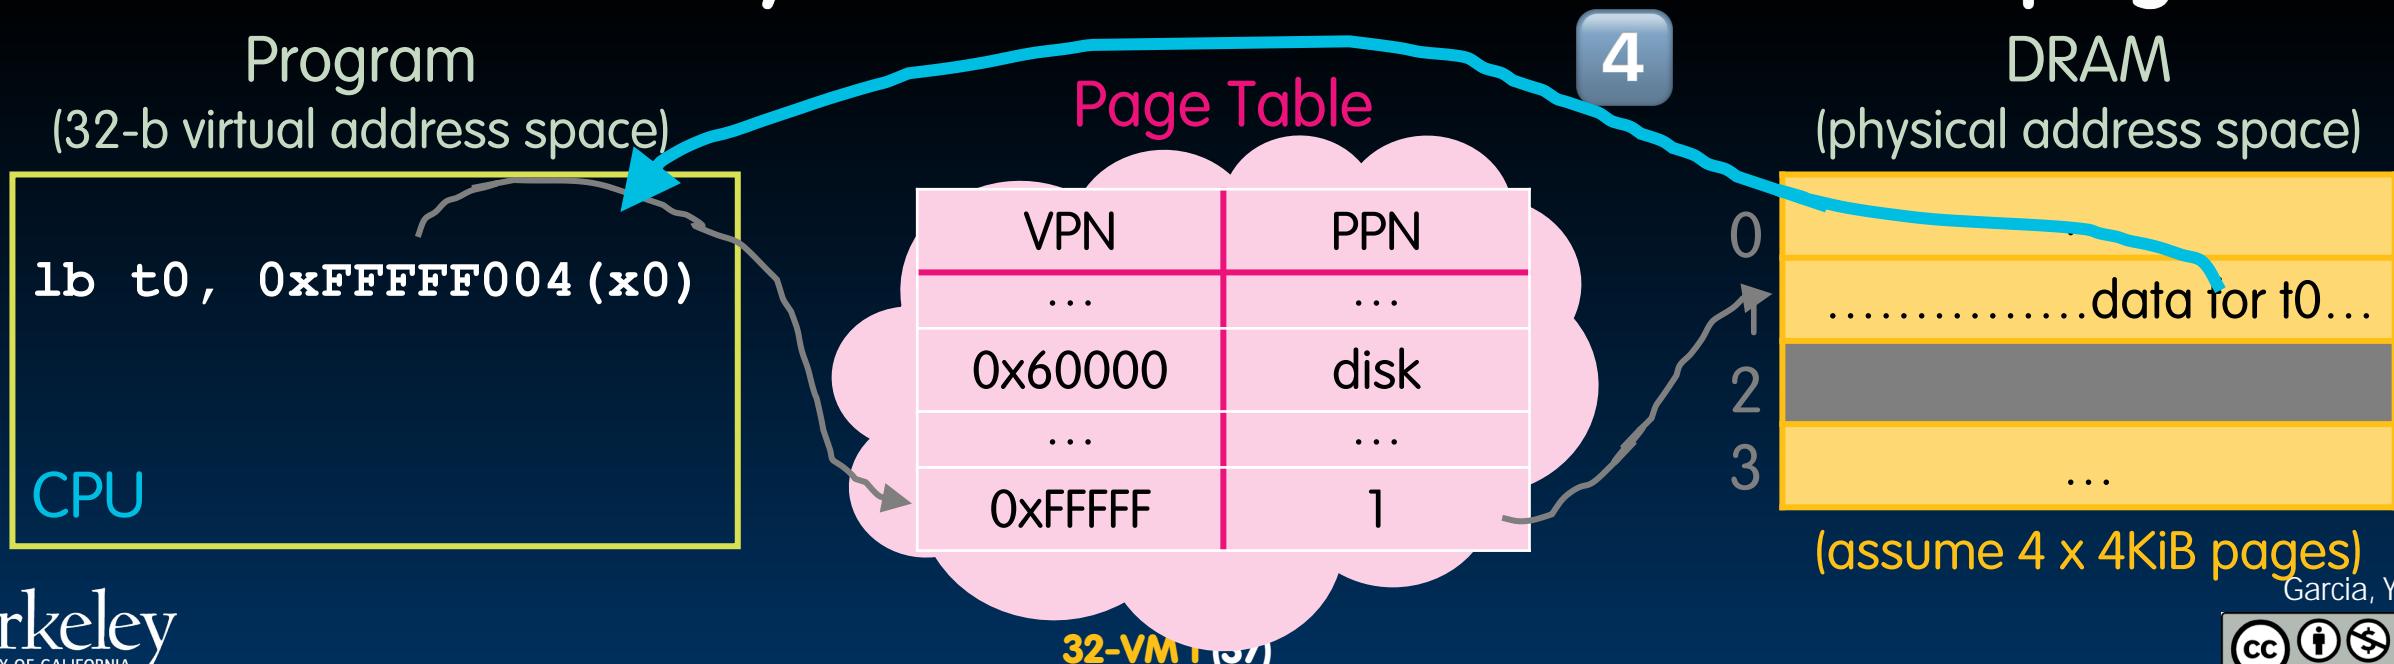
# Translation: How a Program Accesses Memory

1. Program executes a load specifying a virtual address (VA).
2. Computer translates VA to the physical address (PA) in memory.
  - Extract virtual page number (VPN) from VA, e.g., top 20 bits
  - Look up physical page number (PPN) in page table
  - Construct PA: physical page number + offset (from virtual address)
- 3. If the physical page is not in memory, then OS loads it in from disk.**
4. The OS reads memory at the PA and returns the data to the program.



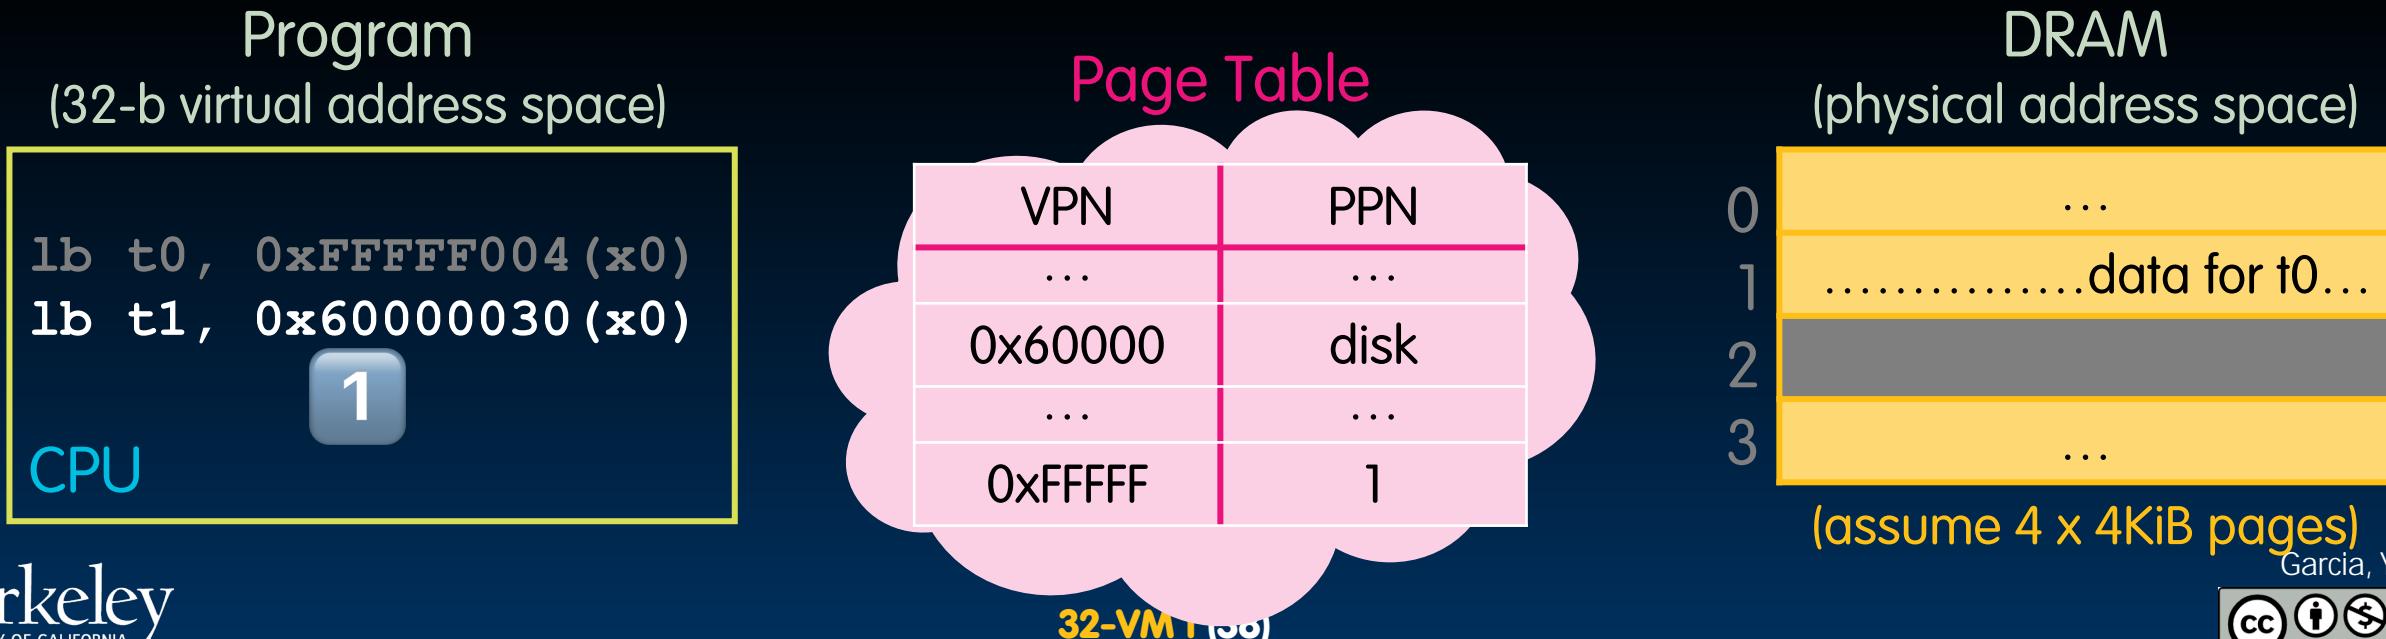
# Translation: How a Program Accesses Memory

1. Program executes a load specifying a virtual address (VA).
2. Computer translates VA to the physical address (PA) in memory.
  - Extract virtual page number (VPN) from VA, e.g., top 20 bits
  - Look up physical page number (PPN) in page table
  - Construct PA: physical page number + offset (from virtual address)
3. If the physical page is not in memory, then OS loads it in from disk.
4. The OS reads memory at the PA and returns the data to the program.



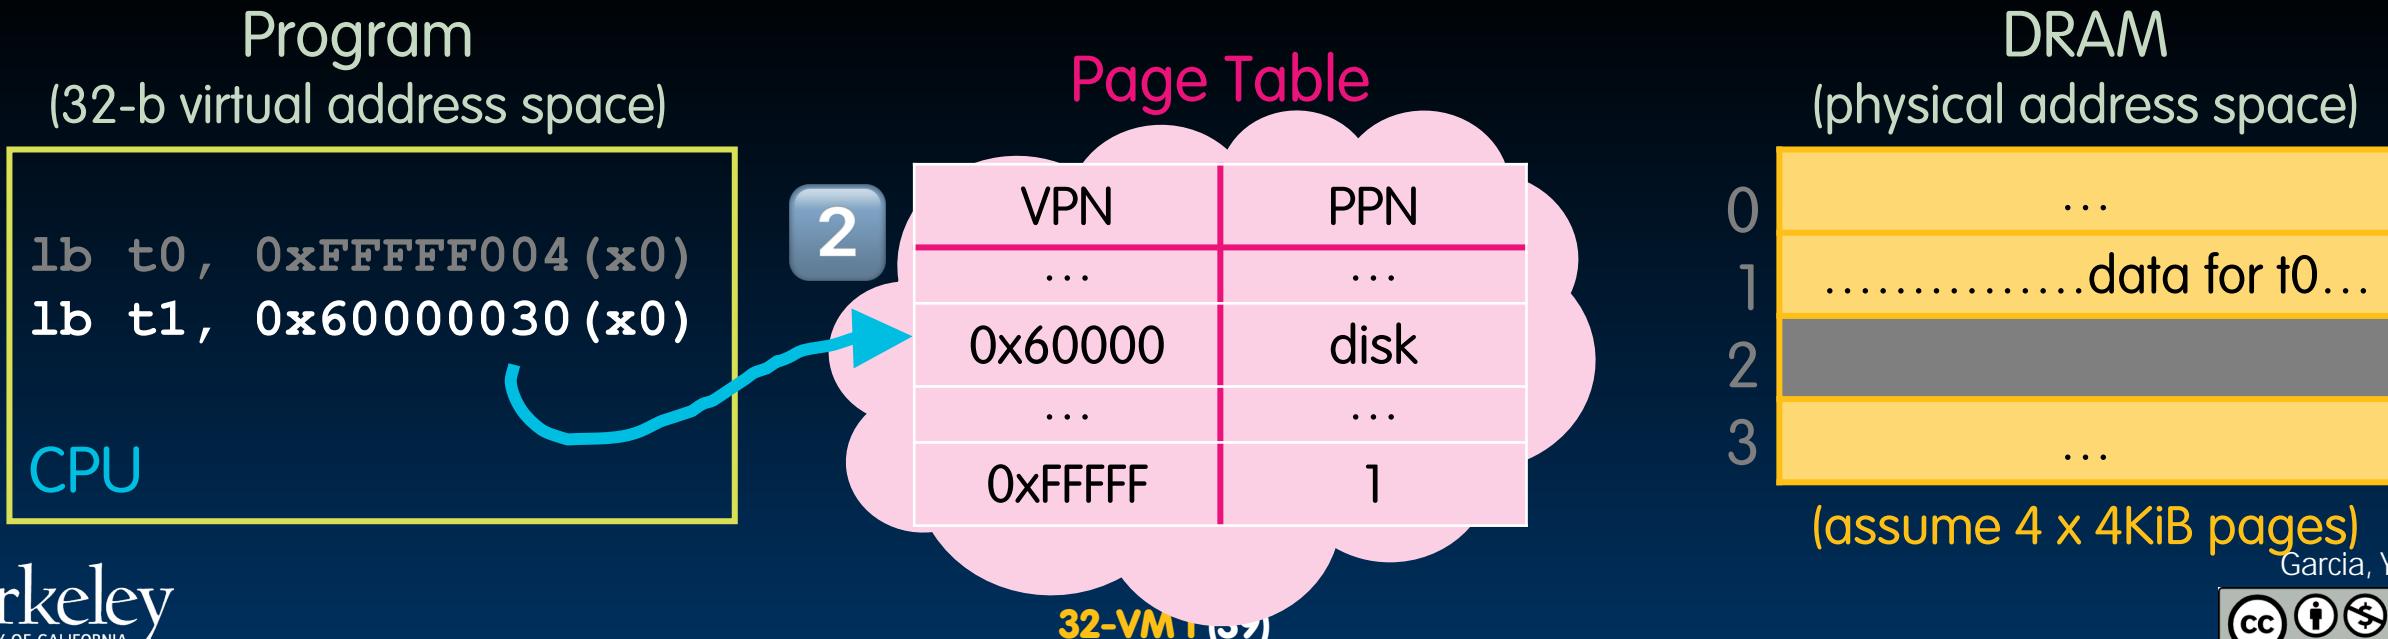
# Translation: How a Program Accesses Memory

1. Program executes a load specifying a virtual address (VA).
2. Computer translates VA to the physical address (PA) in memory.
  - Extract virtual page number (VPN) from VA, e.g., top 20 bits
  - Look up physical page number (PPN) in page table
  - Construct PA: physical page number + offset (from virtual address)
3. If the physical page is not in memory, then OS loads it in from disk.
4. The OS reads memory at the PA and returns the data to the program.



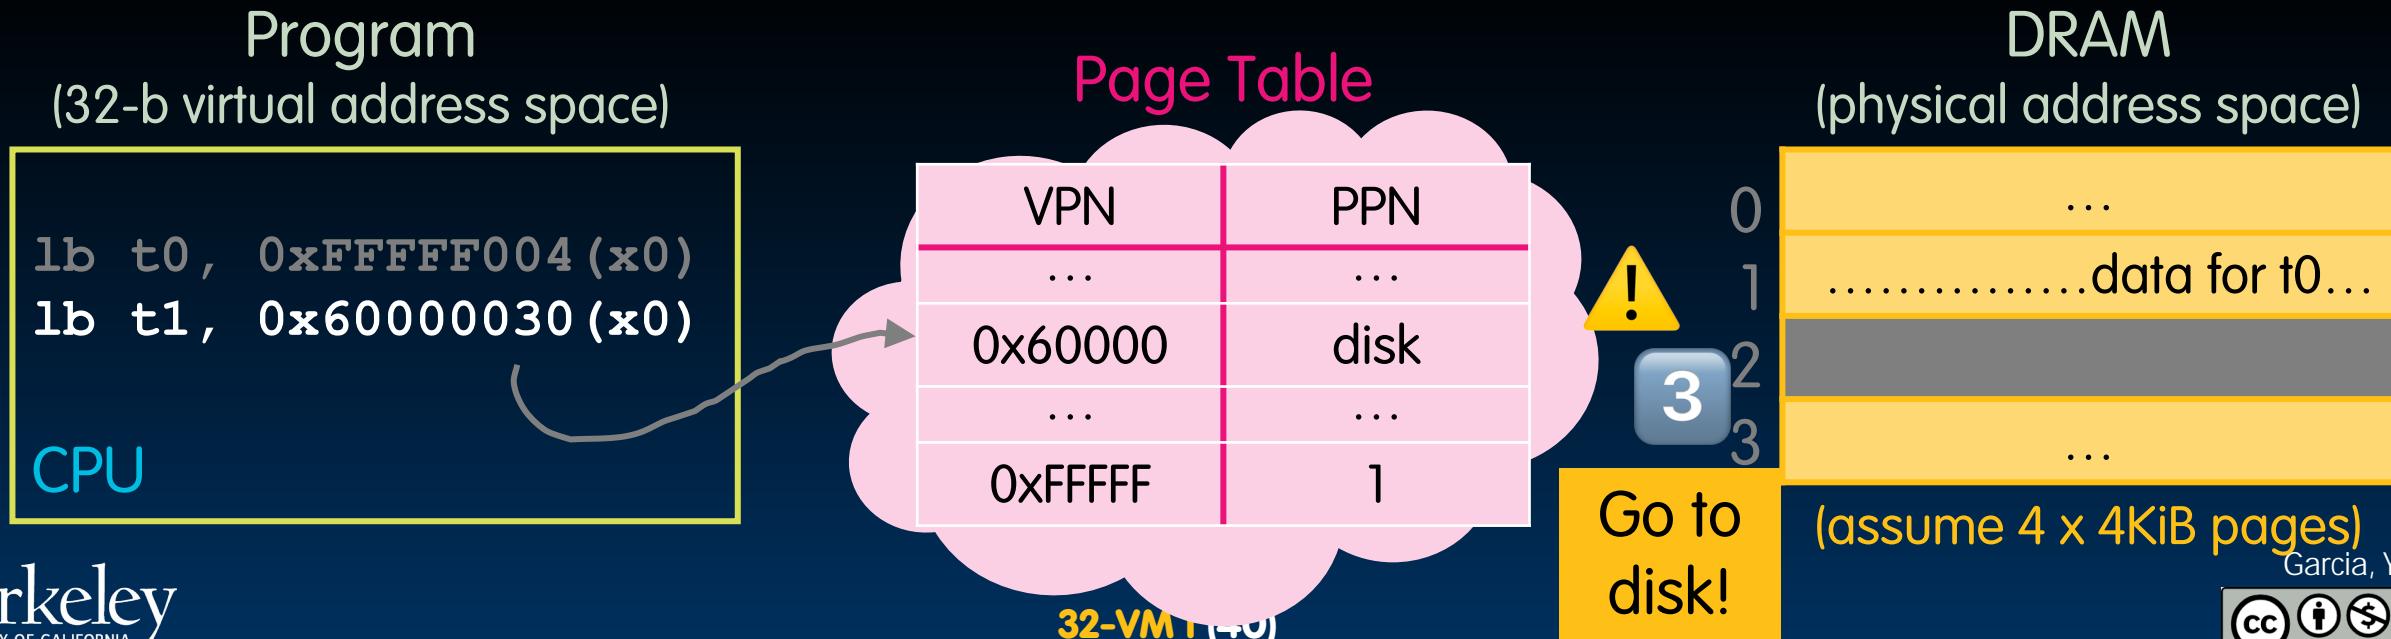
# Translation: How a Program Accesses Memory

1. Program executes a load specifying a virtual address (VA).
2. Computer translates VA to the physical address (PA) in memory.
  - Extract virtual page number (VPN) from VA (e.g., top 20 bits if page size 4KiB =  $2^{12}$  B)
  - Look up physical page number (PPN) in page table
  - Construct PA: physical page number + offset (from virtual address)
3. If the physical page is not in memory, then OS loads it in from disk.
4. The OS reads memory at the PA and returns the data to the program.



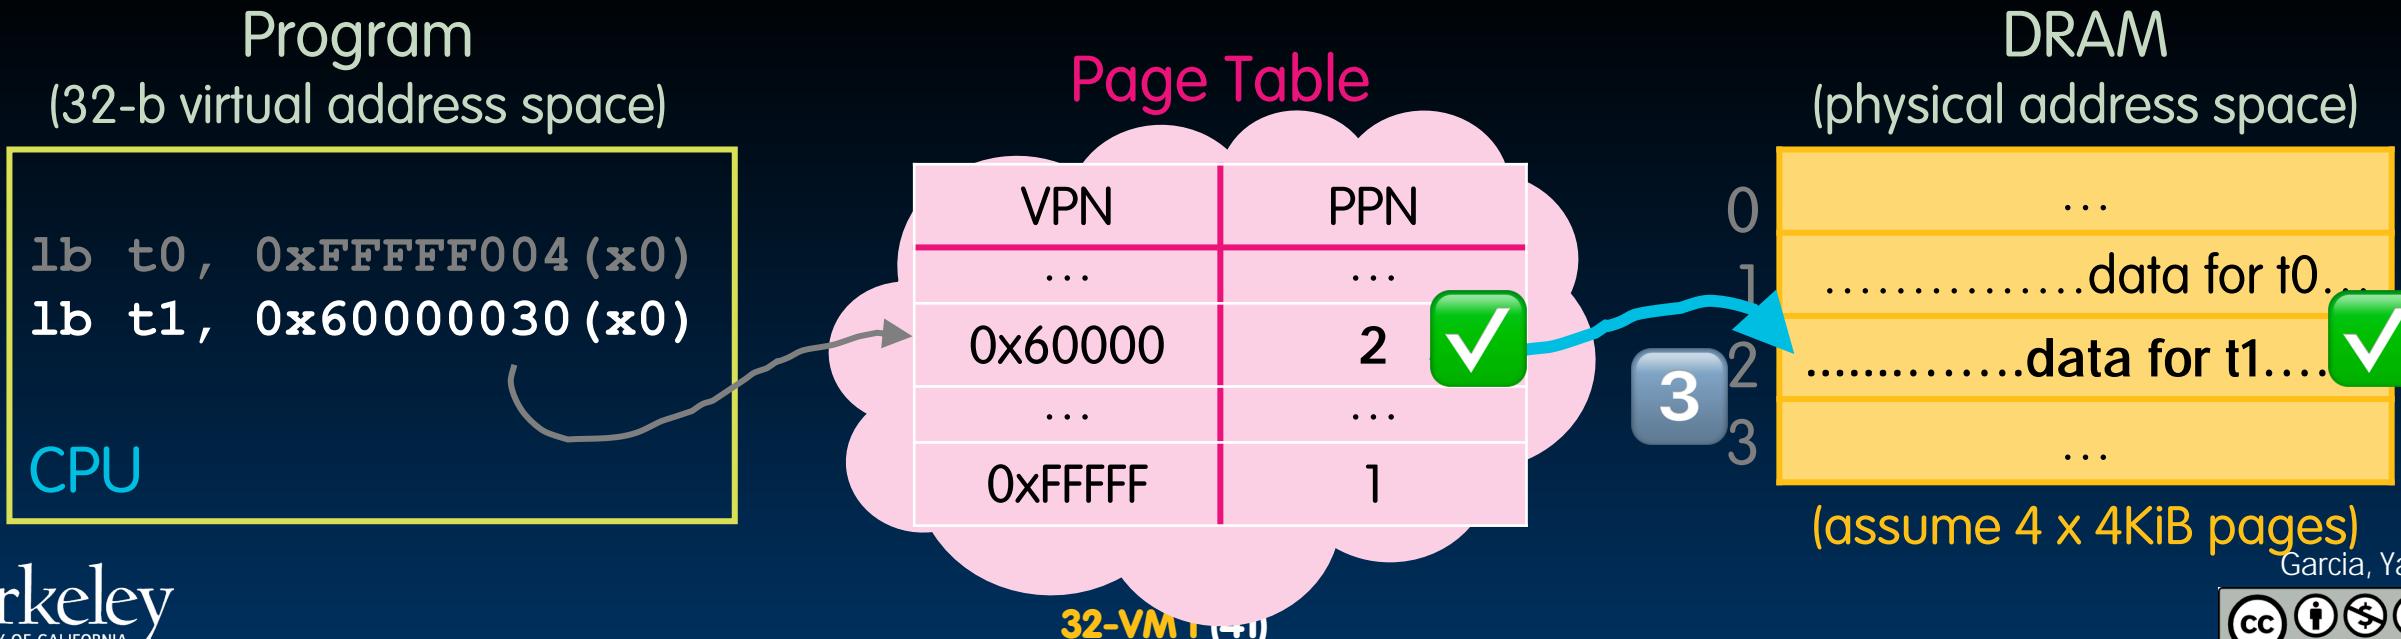
# Translation: How a Program Accesses Memory

1. Program executes a load specifying a virtual address (VA).
2. Computer translates VA to the physical address (PA) in memory.
  - Extract virtual page number (VPN) from VA, e.g., top 20 bits
  - Look up physical page number (PPN) in page table
  - Construct PA: physical page number + offset (from virtual address)
3. If the physical page is not in memory, then OS loads it in from disk.
4. The OS reads memory at the PA and returns the data to the program.



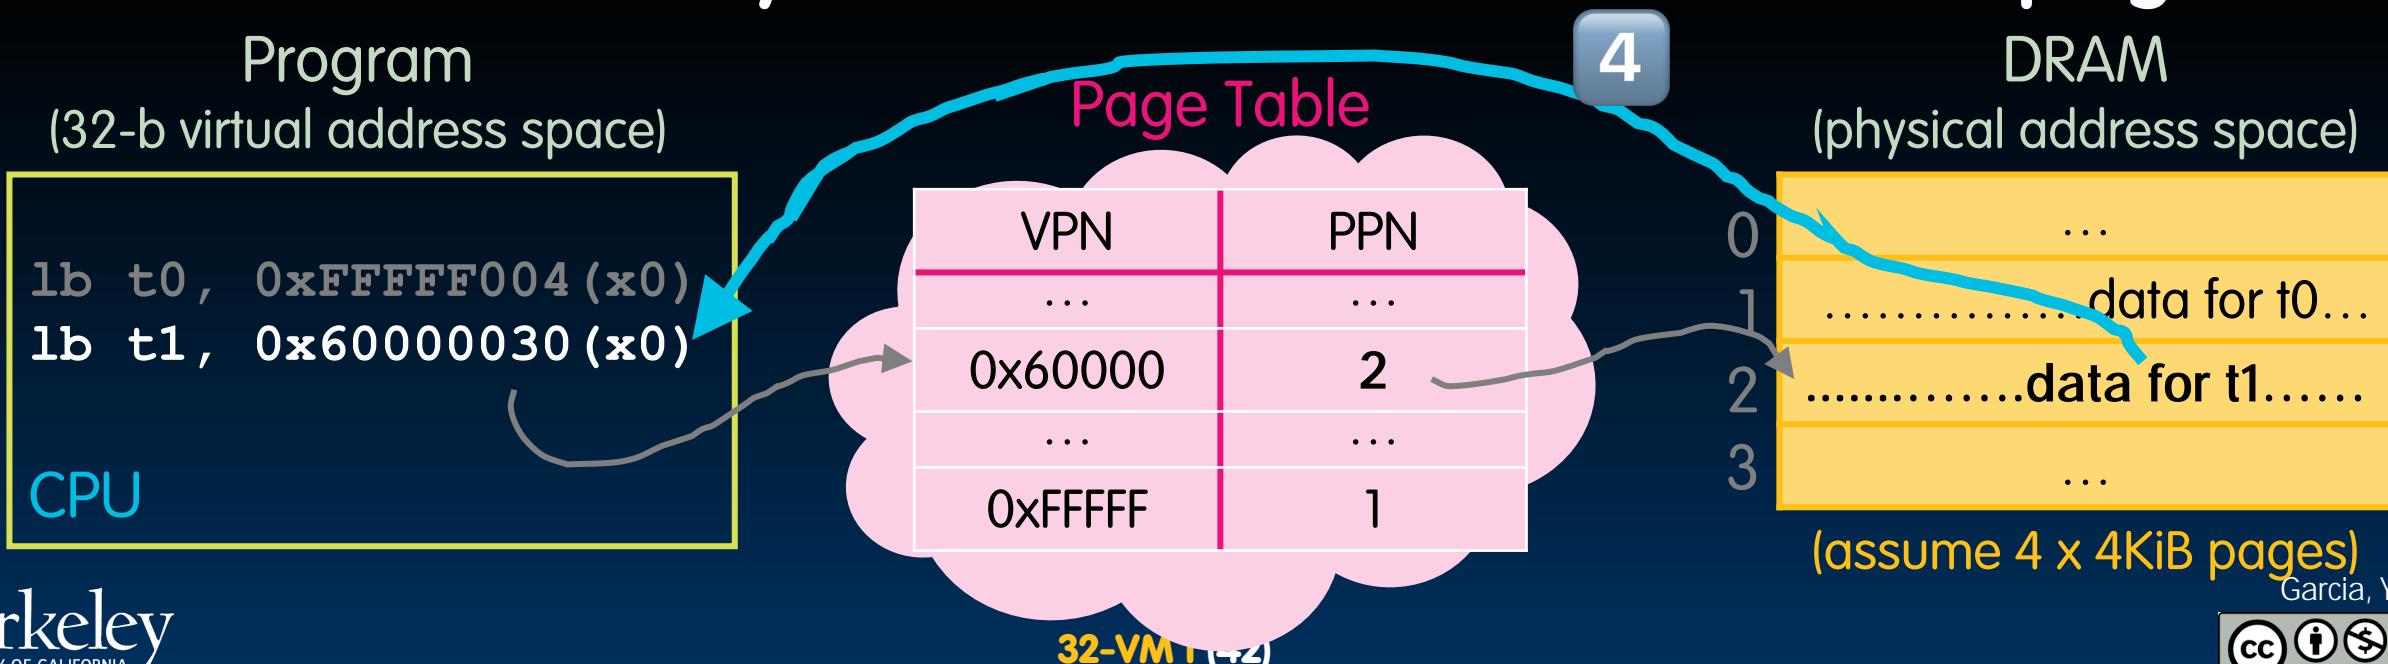
# Translation: How a Program Accesses Memory

1. Program executes a load specifying a virtual address (VA).
2. Computer translates VA to the physical address (PA) in memory.
  - Extract virtual page number (VPN) from VA, e.g., top 20 bits
  - Look up physical page number (PPN) in page table
  - Construct PA: physical page number + offset (from virtual address)
3. If the physical page is not in memory, then OS loads it in from disk.
4. The OS reads memory at the PA and returns the data to the program.



# Translation: How a Program Accesses Memory

1. Program executes a load specifying a virtual address (VA).
2. Computer translates VA to the physical address (PA) in memory.
  - Extract virtual page number (VPN) from VA, e.g., top 20 bits
  - Look up physical page number (PPN) in page table
  - Construct PA: physical page number + offset (from virtual address)
3. If the physical page is not in memory, then OS loads it in from disk.
4. The OS reads memory at the PA and returns the data to the program.



# What Do Page Tables Look Like?

- E.g., 32-bit virtual address space, 4-KiB pages

- 2<sup>32</sup> virtual addresses / (2<sup>12</sup> B/page)  
= 2<sup>20</sup> virtual page numbers

- One Page Table per process:

- One entry per virtual page number.
- Entry has physical page number (or disk address) as well as status bits.

Note: A Page Table is NOT a cache!!

- A Page Table does not have data!  
It is a lookup table.
- All VPNs have a valid entry.
  - But if it helps you, “no tags; index is VPN”

Page Table (2 <sup>20</sup> entries)			
0x00000			0
0x60000			2
			...
			disk
			...
0xFFFFF			1
	Status bits (more later)		Memory page/disk address

# Page Table Details I

- The Computer
- OS Basics: Context Switching
- Physical Memory and Disk Storage
- Virtual Memory and Virtual Addresses
- Paged Memory
- Page Table Details I



1. Map virtual addresses to physical addresses.



2. Use both memory and disk.

- Give illusion of larger memory by storing some content on disk.
- Disk is usually much larger and slower than DRAM.

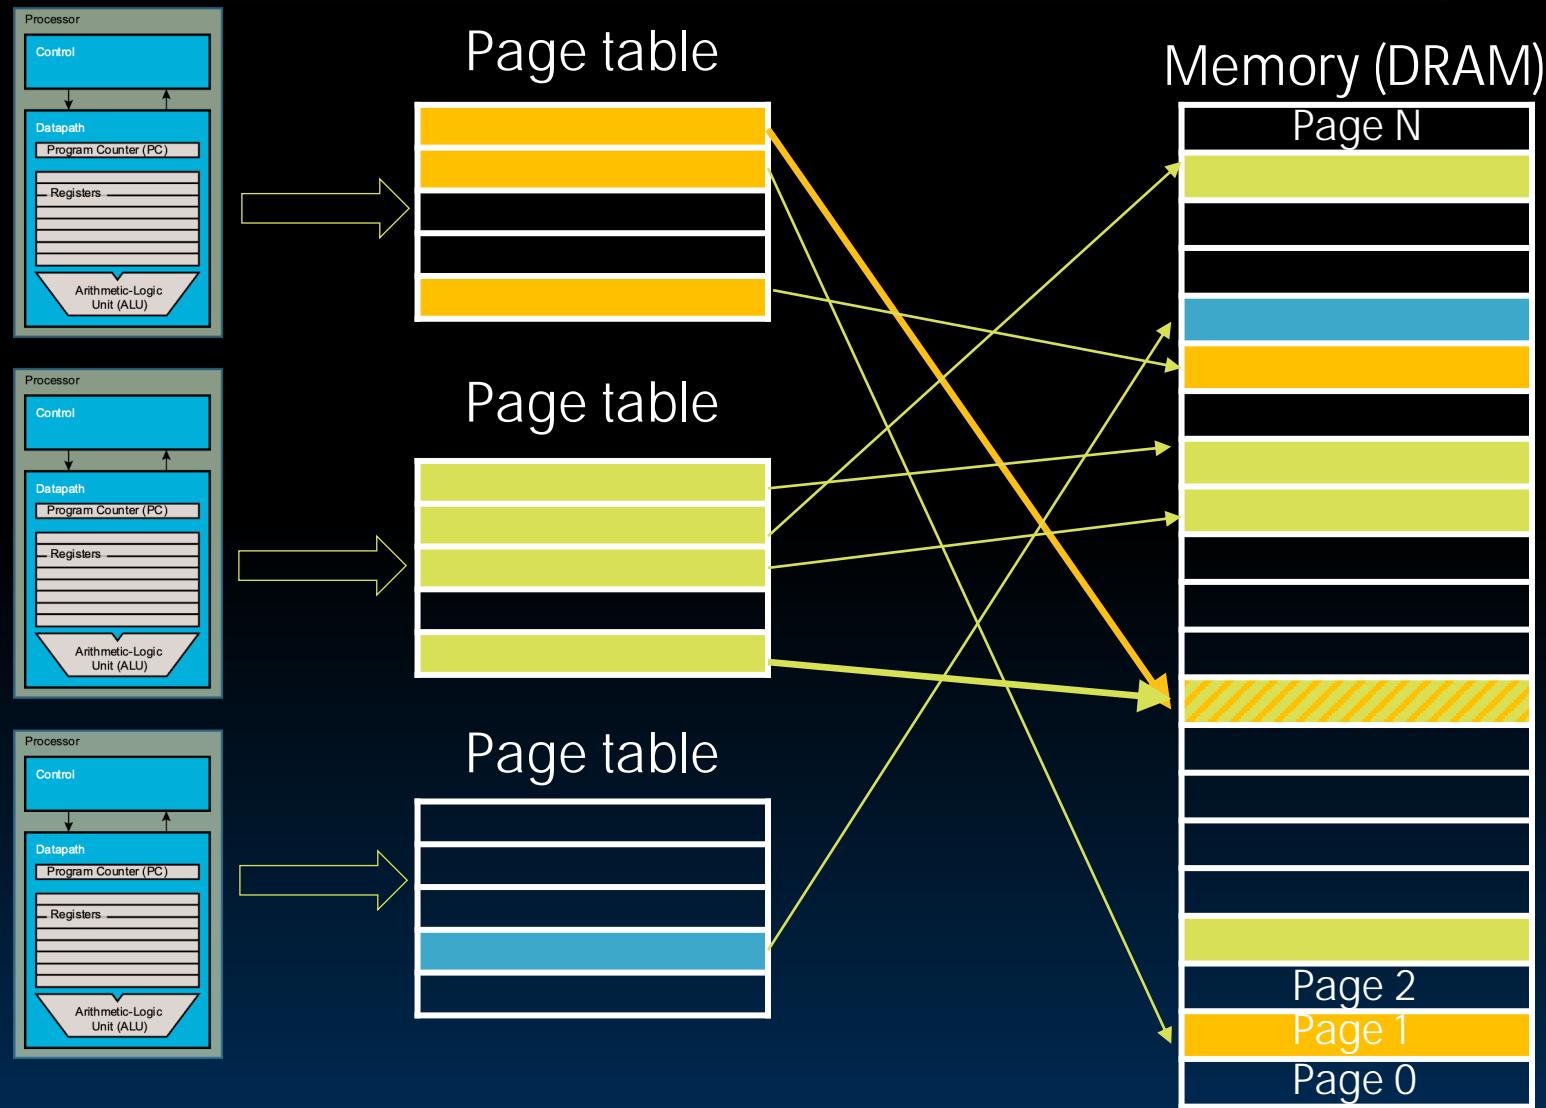
? 3. Protection:

- Isolate memory between processes.
- Each process gets dedicated “private” memory.
- Errors in one program won’t corrupt memory of other programs.
- Prevent user programs from messing with OS’s memory.

What if process tries to modify  
instructions or system data?

# Protection with Page Tables (1/2)

- **Each process has a dedicated page table.**
  - OS keeps track of which process is active.
- **Isolation: Assign processes different pages in DRAM**
  - Prevents accessing other processors' memory
  - Page tables managed by OS
- **Sharing is also possible:**
  - OS may assign same physical page to several processes, e.g., system data

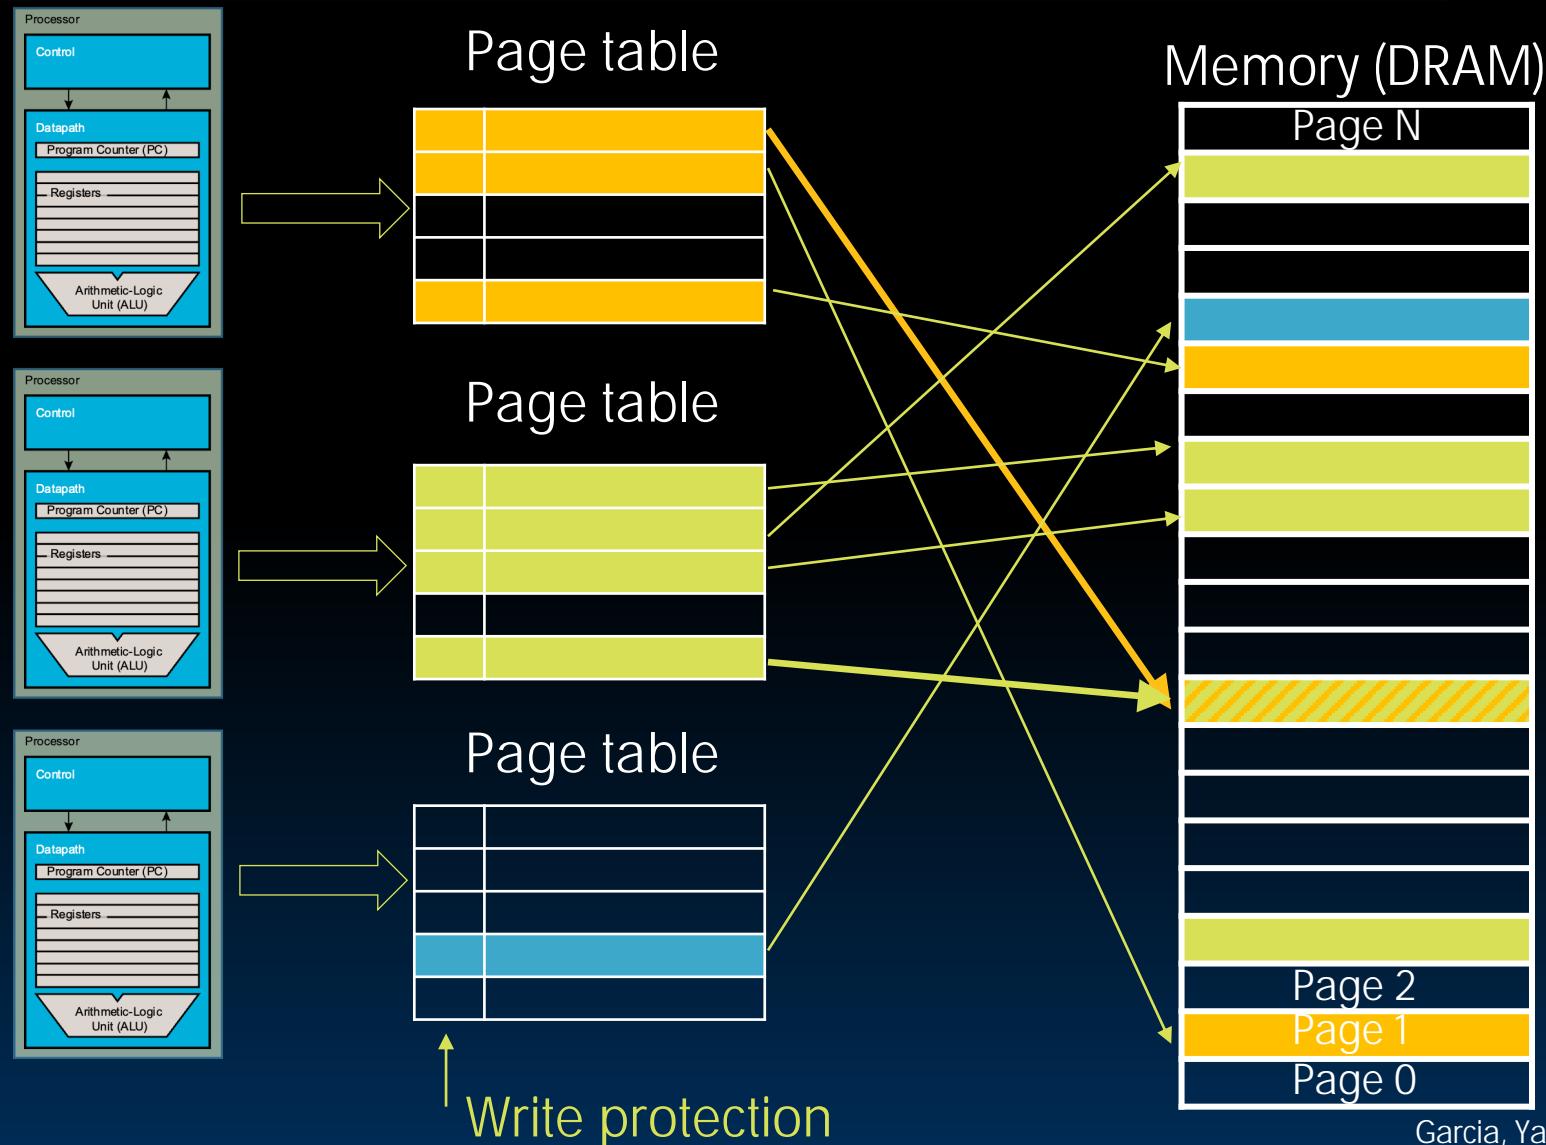


Garcia, Yan



# Protection with Page Tables (2/2)

- **Page Table Entry also includes a write protection bit.**
- **If on, then page is “protected”:**
  - e.g., program code, system data, etc.
  - Writing to a protected page triggers an exception.
  - Exceptions are handled by OS. (more later)



# Page Tables Are Stored in Memory (1/2)

(for next time)

- E.g., 32-Bit virtual address space, 4-KiB pages
  - Single page table size (suppose each entry is 4B, including status bits):
    - $4 \times 2^{20}$  Bytes = 4-MiB
    - 0.1% of 4-GiB memory. Not bad. But much too large for a cache!
- For now, store page tables in memory (DRAM).
  - Caveat: *Two* (slow) memory accesses per **lw/sw** on cache miss!

# Page Tables Are Stored in Memory (2/2)

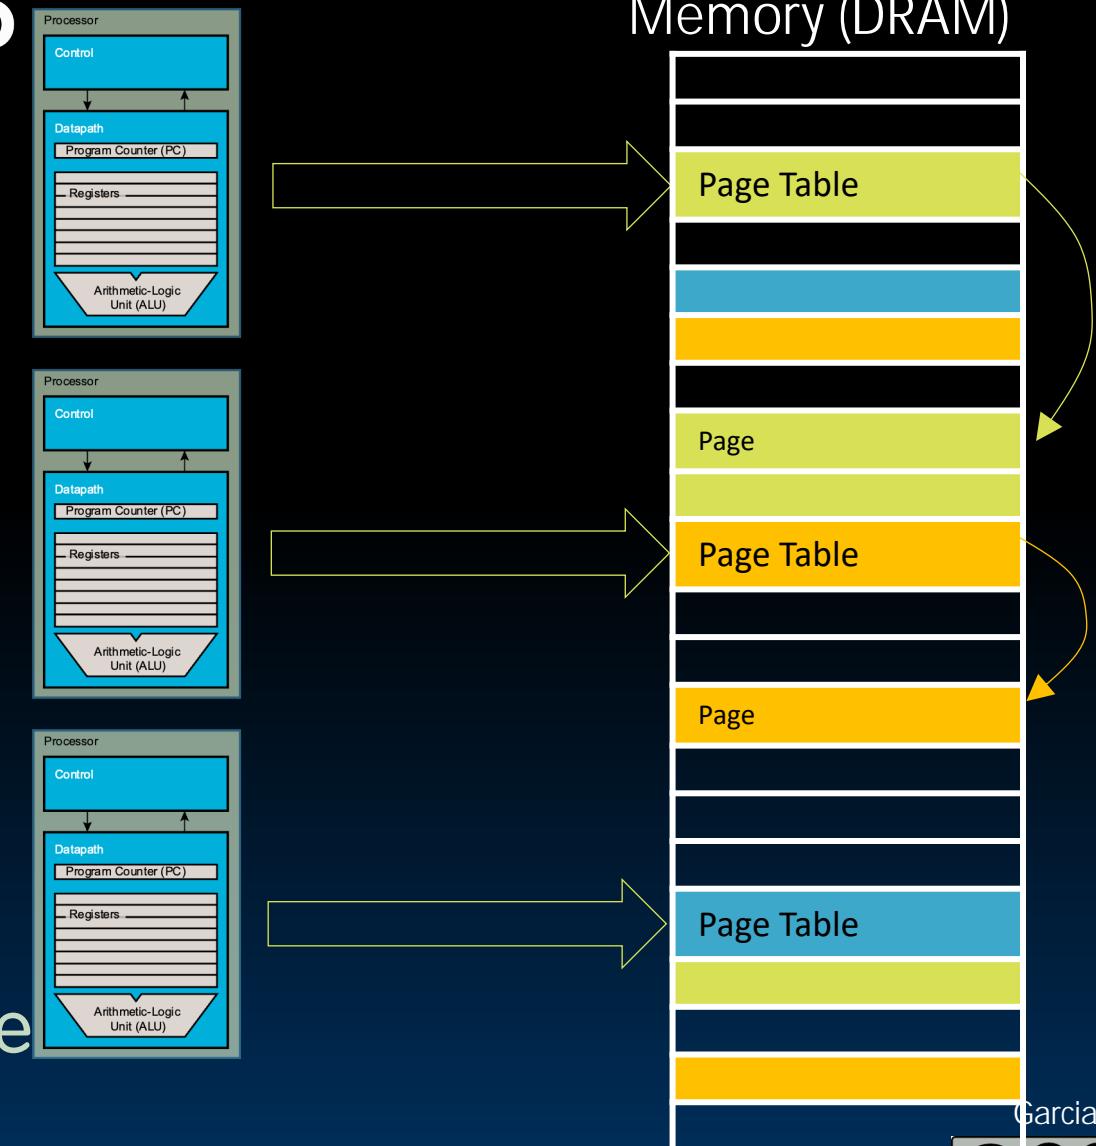
(for next time)

- **Caveat: `lw/sw` then requires two memory accesses:**

- Read page table (stored in main memory) to translate to physical address
  - Read physical page, also in main memory

- **To minimize the performance penalty:**

- Transfer blocks (not words) between DRAM and processor cache
  - Use a cache for frequently used page table entries ... (more later, TLB)



Garcia, Yan



# And in Conclusion...

- **The OS manages resources across multiple processes, all sharing the same CPU, memory, I/O devices, etc.**
- **Each process operates in virtual memory.**
  - For each process, the OS manages virtual $\leftrightarrow$ physical address translation via page tables.
- **Open questions:**
  - How does the OS “context switch”?
  - What if a page is not found in memory?
  - Write-back or write-through?
  - How to incorporate caches with virtual memory?