



UC Berkeley
Teaching Professor
Dan Garcia

CS61C

Great Ideas in **Computer Architecture** (a.k.a. Machine Structures)



UC Berkeley
Teaching Professor
Lisa Yan

C Pointers, Arrays, and Strings

指针

数组

字符串

Pointers

- Pointers
- Using Pointers Effectively
- Pointer Arithmetic
- Array Pitfalls
- Strings
- Word Alignment
- Function Pointer Example
(recorded)

Memory Is a Single Huge Array

- Consider memory to be a byte-addressed array.
 - Each cell of the array has an address associated with it.
 - Each cell also stores some value.
- Don't confuse the address referring to a memory location with the value stored in that location.
- For now, the abstraction lets us think we have access to ∞ memory, numbered from 0...

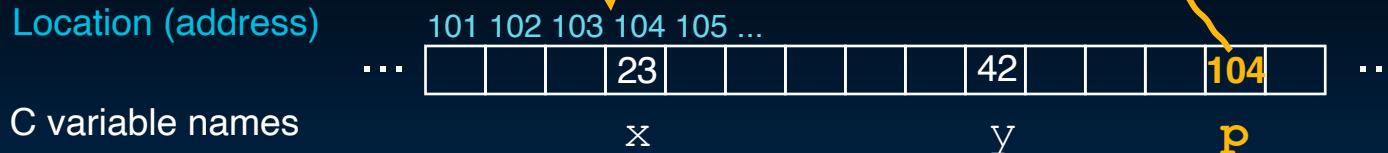


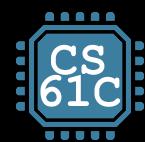
address help us locate the value

Memory Is a Single Huge Array

- An address refers to a particular memory location.
- In other words, it “points” to a memory location.
- Pointer:** A variable that contains the address of another variable.

pointer stores address of a variable





Pointer Syntax

0x100

p ???

0x104

x 3

```
1 int *p;  
2 int x = 3;
```

```
3 p = &x;
```

```
4 printf("p points to %d\n",  
        *p);
```

```
5 *p = 5;
```

- Declaration
- Tells compiler that variable **p** is address of an **int**

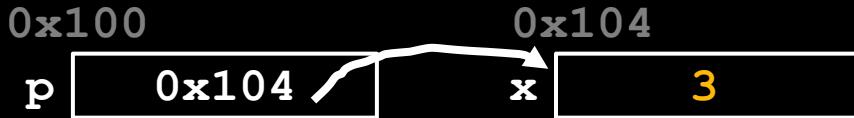


```
1 int *p;  
2 int x = 3;  
  
3 p = &x;  
  
4 printf("p points to %d\n",  
       *p);  
  
5 *p = 5;
```

- Declaration
- Tells compiler that variable **p** is address of an **int**
- Tells compiler to assign **address of x** to **p**
- **&:** “**address operator**” in this context

assign the address of x to p

& means to acquire the address of a variable



- ```
1 int *p;
2 int x = 3;

3 p = &x;

4 printf("p points to %d\n",
 *p);

5 *p = 5;
```
- Declaration
  - Tells compiler that variable **p** is address of an **int**
  
  - Tells compiler to assign address of **x** to **p**
  - **&:** “**address operator**” in this context
  
  - Gets **value pointed to by p**
  - **\***: “**dereference operator**” in this context
- \* mean acquire the value according to its address*

0x100                    0x104



```
1 int *p;
2 int x = 3;

3 p = &x;

4 printf("p points to %d\n",
 *p);

5 *p = 5;
```

- Declaration
- Tells compiler that variable **p** is address of an **int**
  
- Tells compiler to assign address of **x** to **p**
- **&:** “**address operator**” in this context
  
- Gets value pointed to by **p**
- **\***: “**dereference operator**” in this context
  
- Changes **variable pointed to by p**
- Use dereference operator **\*** on left of =

# Pointer Syntax



```
1 int *p;
2 int x = 3;
3
4 p = &x;
5
6 printf("p points to %d\n",
7 *p);
8
9 *p = 5;
```

The “\*” is used two ways here:

**Declaration** (L1): Indicate **p** is a pointer

**Dereference** (L4-5): Value pointed to by **p**

- Declaration
- Tells compiler that variable **p** is address of an **int**
  
- Tells compiler to assign address of **x** to **p**
- **&:** “**address operator**” in this context
  
- Gets **value pointed to by p**
- **\*:** “**dereference operator**” in this context
  
- Changes **variable pointed to by p**
- Use dereference operator **\*** on left of =

Garcia, Yan

# Pointers are Useful When Passing Parameters

- Java and C pass parameters “by value”: 值传递
  - A procedure/function/method gets a copy of the parameter.

Changing the function’s copy cannot change the original.

```
void addOne (int x)
{
 x = x + 1; x [3 4]
}

int y = 3; y [3]
addOne(y);
```

⚠️ y is still 3...

can't alter the origin functions will create a new local variable and the value is copied from parameter.

# Pointers are Useful When Passing Parameters

- Java and C pass parameters “by value”:
  - A procedure/function/method gets a copy of the parameter.

Changing the function’s copy cannot change the original.

```
void addOne (int x)
{
 x = x + 1; x [3 4]
}

int y = 3; y [3]
addOne(y);
```

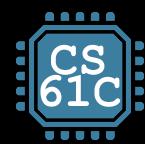
⚠️ y is still 3...

To get a function to change a value, **pass in a pointer**.

```
void addOne (int *p)
{
 *p = *p + 1; p [0x100]
}

int y = 3;
addOne(&y);
```

✓ y is now 4 !



# Pointers in C ... The Good, Bad, and the Ugly

Why use pointers? *pass a large struct or array to function*

- To pass a large struct or array to a function, it's easier/faster/etc. to pass a pointer.
  - Otherwise, we'd need to copy a huge amount of data!
- At the time C was invented (early 1970s), compilers didn't produce efficient code, so C was designed to give human programmer more flexibility.
  - Nowadays, computers are 100,000x faster; compilers are also way, way, way better.
- Still used for low-level system code, as well as implementation of "pass-by-reference" object paradigms in other languages.
- In general, pointers allow cleaner, more compact code.

## ⚠ So, what are the drawbacks?

- Pointers are probably the single largest source of bugs in C. **Be careful!**
  - Most problematic with dynamic memory management
  - Dangling references and memory leaks

} (more next time)

# Common C Bug: Garbage Addresses



- Declaring a pointer just allocates space to hold the pointer – it does not allocate something to be pointed to!
- Local variables in C are not initialized. They may contain *anything*.
- What does the following code do?

```
void f()
{
 int *ptr;
 *ptr = 5;
}
```

???



# Using Pointers Effectively

- Pointers
- Using Pointers Effectively
- Pointer Arithmetic
- Array Pitfalls
- Strings
- Word Alignment
- Function Pointer Example  
(recorded)

# Pointers to Different Data Types

- Pointers are used to point to any data type.
- Normally a pointer can only point to one type.
  - void \* is a type that can point to anything (generic pointer).
  - Use sparingly to help avoid program bugs... and security issues... and a lot of other bad things!

## You can even have pointers to functions...

- `int (*fn) (void *, void *) = &foo;`
  - `fn` is a function that accepts two `void *` pointers and returns an `int` and is initially pointing to the function `foo`.
- `(*fn) (x, y);` will then call the function

```
int *xptr;
char *str;
struct llist *foo_ptr;
```

void \*

Check out Dan's recorded function pointer example:

<https://www.youtube.com/watch?v=P1IYbp0fgY4>



# NULL pointers...

- **The pointer of all 0s is special.**
  - The "NULL" pointer, like in Java, Python, etc...
- **Since "0 is false", its very easy to do tests for null:**
  - `if(!p) { /* p is a null pointer */ }`
  - `if(q) { /* q is not a null pointer */ }`
- **If you write to or read from a null pointer, your program should crash.**

```
char *p = NULL;
p 0x00000000
```

# Structs, Revisited

```
typedef struct {
 int x;
 int y;
} Coord;

/* declarations */
Coord coord1, coord2;
Coord *ptr1, *ptr2;

/* instantiations
go here... */
```

```
/* dot notation */
int h = coord1.x;
coord2.y = coord1.y;

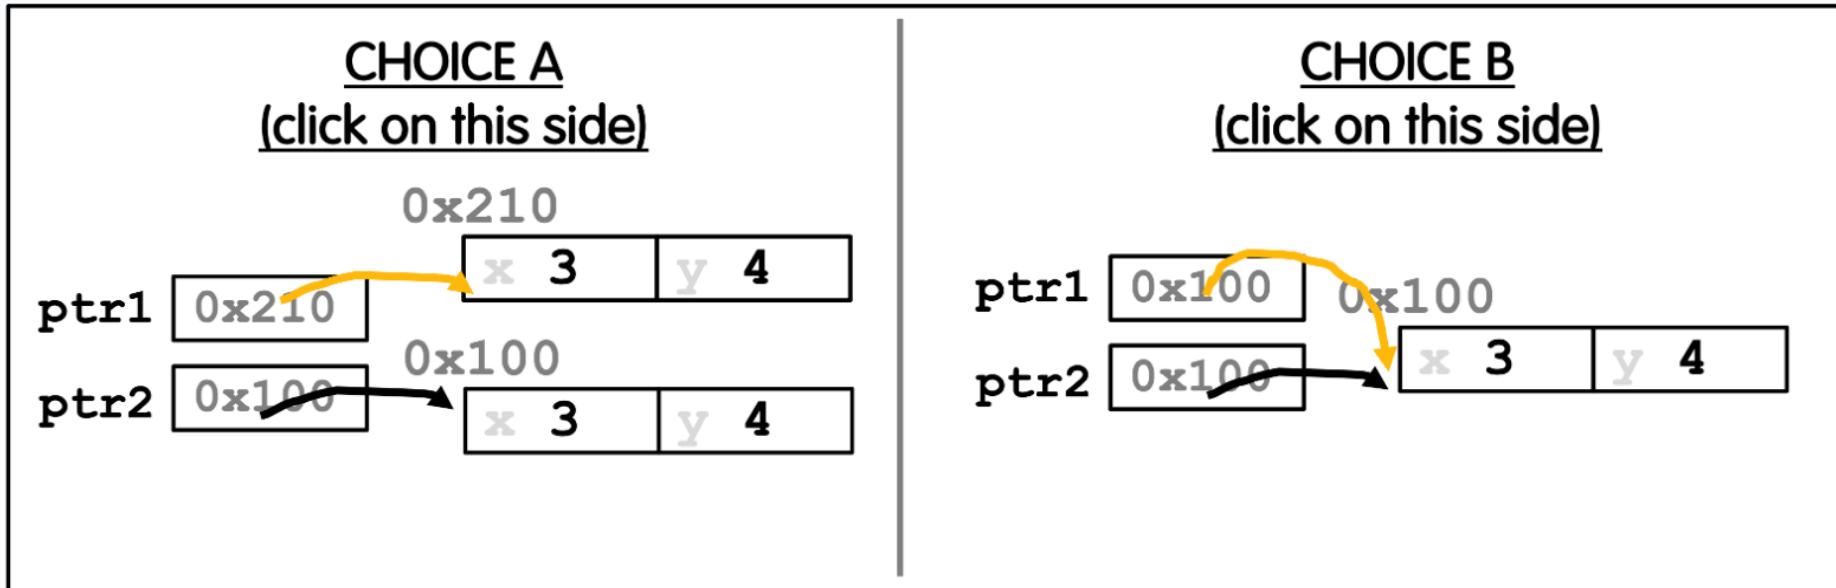
/* arrow notation */
int k;
k = ptr1->x;
k = (*ptr1).x; // equivalent

/* This compiles, but what does it do? */
ptr1 = ptr2;
```



# Assume ptr2 points to an initialized Coord struct {3, 4}.

## What does `ptr1 = ptr2;` do?



```
typedef struct {
 int x;
 int y;
} Coord;

/* declarations */
Coord coord1, coord2;
Coord *ptr1, *ptr2;

/* instantiations
go here... */
```

```
/* dot notation */
int h = coord1.x;
coord2.y = coord1.y;
```

```
/* arrow notation */
int k;
k = ptr1->x;
k = (*ptr1).x; // equivalent
```

/\* This compiles, but what does it do? \*/  
ptr1 = ptr2;



# Pointer Arithmetic

- Pointers
- Using Pointers Effectively
- Pointer Arithmetic
- Array Pitfalls
- Strings
- Word Alignment
- Function Pointer Example  
(recorded)



# A C array is really just a big block of memory

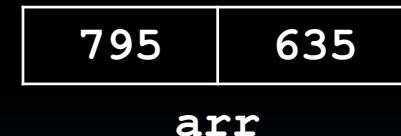
- Declaration:

- `int arr[2];`
  - ...declares a 2-element integer array



- Declaration and initialization

- `int arr[] = {795, 635};`
  - declares and fills a 2-elt integer array



- Accessing elements:

- `arr[num]`
  - returns the num<sup>th</sup> element.
  - This is shorthand for **pointer arithmetic**.

`arr[0]; // 795`

# Pointer Arithmetic

**sizeof**: compile-time op  
for # of bytes in object



**pointer + n**

Adds **n\*sizeof** ("whatever pointer is pointing to") to memory address.

**pointer - n**

Subtracts **n\*sizeof** ("whatever pointer is pointing to") from memory address.

$a[i] \equiv * (a+i)$

```
// 32-bit signed int array
int32_t arr[] = {50, 60, 70};

int32_t *q = arr;
```

pointer  $\pm$  integer

`sizeof`: compile-time op  
for # of bytes in object



### pointer + n

Adds `n*sizeof` ("whatever pointer is pointing to") to memory address

### pointer - n

Subtracts `n*sizeof` ("whatever pointer is pointing to") from memory address

$$a[i] \equiv * (a+i)$$

```
// 32-bit signed int array
int32_t arr[] = {50, 60, 70};

int32_t *q = arr;

printf(" *q: %d is %d\n", *q, q[0]);
printf("* (q+1): %d is %d\n", *(q+1), q[1]);
printf("* (q-1): %d is %d\n", *(q-1), q[-1]);
```

\*q: 50 is 50  
\*(q+1): 60 is 60  
\*(q-1): ??? is ???



## How to get a function to change a pointer?

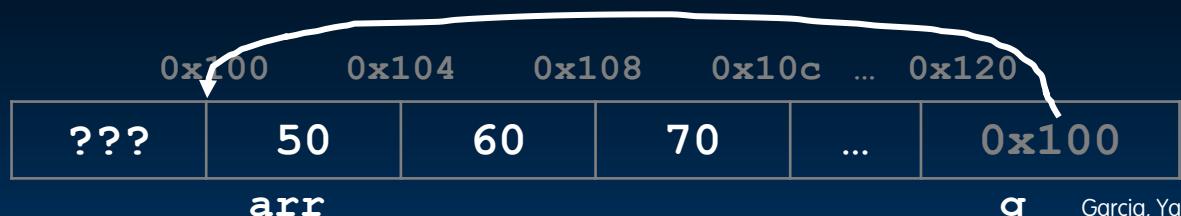
- Suppose we want `increment_ptr` to change where `q` points to.
- What gets printed?  
`*q is 50`

```
void increment_ptr(int32_t *p)
{ p = p + 1; }
```

0x100  
0x104

p

```
int32_t arr[3] = {50, 60, 70};
int32_t *q = arr;
increment_ptr(q);
printf("*q is %d\n", *q);
```





## Remember: C is pass-by-value!

- Instead, pass a **pointer to a pointer** ("handle").
- Declared as **data\_type \*\*h**.
- Now what gets printed?

\*q is 60

```
void increment_ptr(int32_t **h)
{ *h = *h + 1; }
```

0x120

h

```
int32_t arr[3] = {50, 60, 70};
int32_t *q = arr;
increment_ptr(&q);
printf("*q is %d\n", *q);
```



# Array Pitfalls

- Pointers
- Using Pointers Effectively
- Pointer Arithmetic
- Array Pitfalls
- Strings
- Word Alignment
- Function Pointer Example  
(recorded)



# Array Pitfall #1 of 123234983



- Declare array and initialize all elements of an array of known size n:
  - Wrong

```
int i, ar[10];
for(i = 0; i < 10; i++){ ... }
```

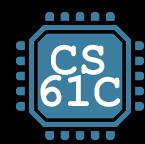
???

- Strongly encouraged

```
int ARRAY_SIZE = 10;
int i, a[ARRAY_SIZE];
for(i = 0; i < ARRAY_SIZE; i++){ ... }
```

- Why? SINGLE SOURCE OF TRUTH

- Utilize indirection and avoid maintaining two copies of the number 10!



# Arrays vs Pointers

- **Arrays are (almost) identical to pointers**
  - `char *string` and `char string[]` are nearly identical declarations
  - They differ in very subtle ways: incrementing, declaration of filled arrays... (more in a bit)

## ▪ Accessing Array Elements

- `arr` is an array variable, but it looks like a pointer in many respects (though not all).
- `arr[0]` is the same as `*arr`
- `arr[2]` is the same as `* (arr+2)`

An array variable is a “pointer” to the first (0-th) element.

# Arrays are not implemented as you'd think...

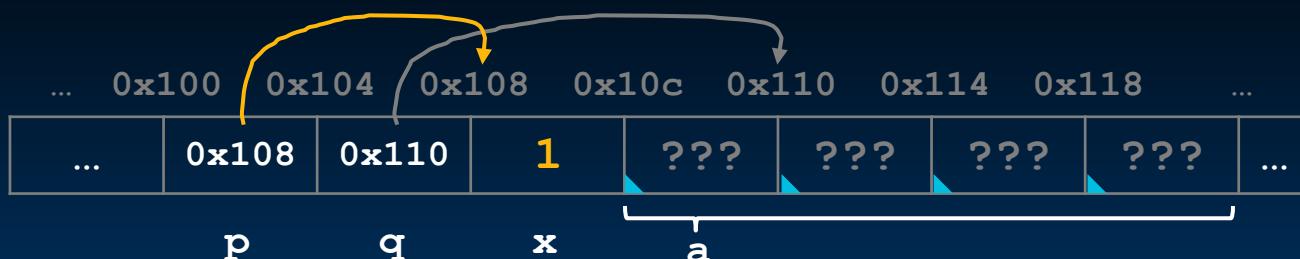
```
int *p, *q, x;
int a[4];
p = &x;
q = a + 1;
```

```
*p:1, p:108, &p:100
```

```
{ *p = 1;
printf("*p:%d, p:%x, &p:%x\n", *p, p, &p); // %d:signed decimal, %x:hex

*q = 2;
printf("*q:%d, q:%x, &q:%x\n", *q, q, &q);

*a = 3;
printf("*a:%d, a:%x, &a:%x\n", *a, a, &a);
```



# Arrays are not implemented as you'd think...

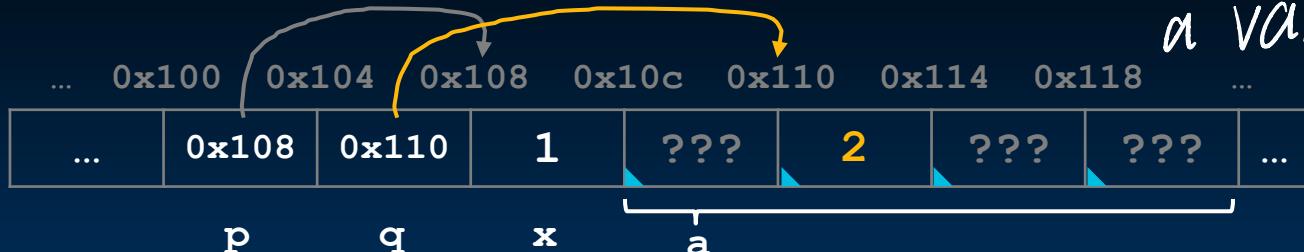
```
int *p, *q, x;
int a[4];
p = &x;
q = a + 1;

*p = 1;
printf("*p:%d, p:%x, &p:%x\n", *p, p, &p); // %d:signed decimal, %x:hex
```

```
*q = 2;
printf("*q:%d, q:%x, &q:%x\n", *q, q, &q);
```

```
*a = 3;
printf("*a:%d, a:%x, &a:%x\n", *a, a, &a);
```

An array name is not a variable.



# Arrays are not implemented as you'd think...

```
int *p, *q, x;
int a[4];
p = &x;
q = a + 1;

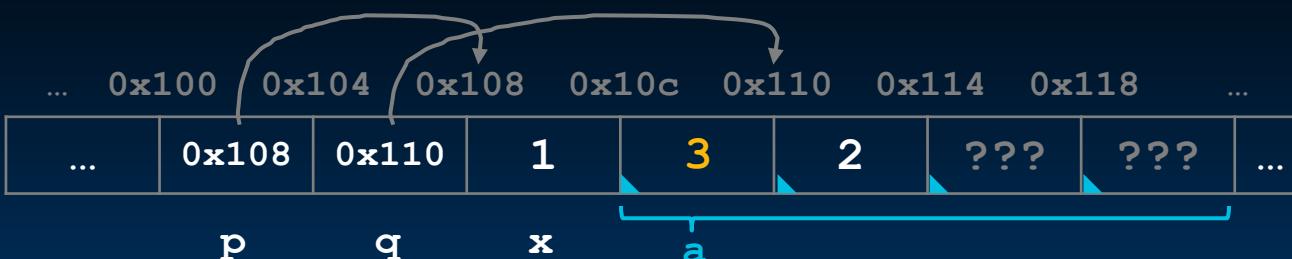
*p = 1;
printf("*p:%d, p:%x, &p:%x\n", *p, p, &p); // %d:signed decimal, %x:hex

*q = 2;
printf("*q:%d, q:%x, &q:%x\n", *q, q, &q);

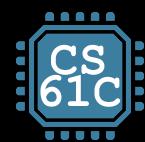
{*a = 3;
printf("*a:%d, a:%x, &a:%x\n", *a, a, &a);
```

\*p:1, p:108, &p:100  
\*q:2, q:110, &q:104  
\*a:3, a:10c, ~~&a:10c~~ !

K&R: "An array name is not a variable"



Lisa: "A C array is really just a big block of memory"



# Array Are Very Primitive (1/3)



## 1. Array bounds are not checked during element access.

- Consequence: We can accidentally access off the end of an array!

```
int ARRAY_SIZE = 100;
int foo[ARRAY_SIZE];
int i;
....
for(i = 0; i <= ARRAY_SIZE; ++i) {
 foo[i] = 0;
}
```

- Corrupts other parts of the program...
  - Including internal C data
- May cause crashes later

≡

### WhatsApp Security Advisories

CVE-2019-11933

A heap buffer overflow bug in libpl\_droidsonroids\_gif before 1.2.19, as used in WhatsApp for Android before version 2.19.291 could allow remote attackers to execute arbitrary code or cause a denial of service.

click a link preview from a specially crafted text message.

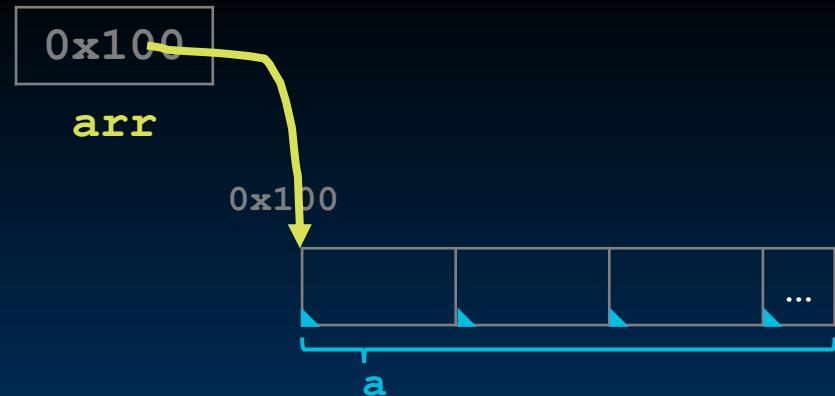
# Array Are Very Primitive (2/3)



1. Array bounds are not checked during element access.
  - Consequence: We can accidentally access off the end of an array!
2. An array is passed to a function as a pointer.
  - Consequence: The array size is lost!

same as: `int *arr`

```
int bar(int arr[], unsigned int size)
{
 ... arr[size - 1] ...
}
int main(void)
{
 int a[5], b[10];
 ...
 bar(a, 5);
 bar(b, 10);
 ...
}
```





# Array Are Very Primitive (3/3)

1. Array bounds are not checked during element access.
  - Consequence: We can accidentally access off the end of an array!
2. An array is passed to a function as a pointer.
  - Consequence: The array size is lost!
3. Declared arrays are only allocated while the scope is valid.

```
char *foo() {
 char string[32]; ...;
 return string;
} is incorrect
```

*the scope is not valid*

Fix: Dynamic memory allocation!  
(next time)

# Array Are Very Primitive, Summary

1. **Array bounds are not checked during element access.**
  - Consequence: We can accidentally access off the end of an array!
2. **An array is passed to a function as a pointer.**
  - Consequence: The array size is lost!
3. **Declared arrays are only allocated while the scope is valid.**

## Segmentation faults and bus errors:

- These are VERY difficult to find; be careful!
- You'll learn how to debug these in lab with **gdb**...

*gdb*

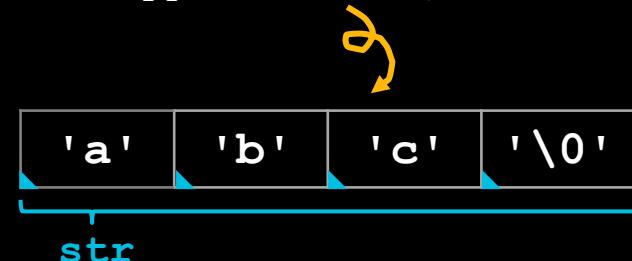
# Strings

- Pointers
- Using Pointers Effectively
- Pointer Arithmetic
- Array Pitfalls
- Strings
- Word Alignment
- Function Pointer Example  
(recorded)

String.h

- A C string is just an array of characters, followed by a null terminator.
  - Null terminator: the byte 0 (number) aka '\0' character)
- The standard C library `string.h` assumes **null-terminated strings**.
  - In particular, string length operation does **not** include the null terminator when you ask for length of a string!

```
char str[] = "abc";
```



```
int strlen(char s[])
{
 int n = 0;
 while (*s++ != 0) { n++; }
 return n;
}
```

... `strlen(str)` ... // 3 !

Garcia, Yan

# Word Alignment

- Pointers
- Using Pointers Effectively
- Pointer Arithmetic
- Array Pitfalls
- Strings
- Word Alignment
- Function Pointer Example  
(recorded)

# Memory and Addresses

1 byte = 8 bit

- Modern machines are “byte-addressable.”

- Hardware’s memory composed of 8-bit storage cells; each byte has a unique address

- We commonly think in terms of “word size”:



- aka number of bits in an address
  - A 32b architecture has 4-byte words.
- ```
sizeof(int *) ==  
sizeof(char *) == 4
```

| | |
|------------|-------------|
| 0xFFFFFFF8 | int32_t * |
| 0xFFFFFFF4 | short * |
| 0xFFFFFFF0 | char * |
| ... | |
| 0xFFFFFEC | xx xx xx xx |
| 0xFFFFFE8 | xx xx xx xx |
| ... | xx xx xx xx |
| 0x0C | xx xx xx xx |
| 0x08 | xx xx xx xx |
| 0x04 | xx xx xx xx |
| 0x00 | xx xx xx xx |

4 bytes

Garcia, Yan

Word Alignment and Endianess

4 byte

- A C pointer is just an abstracted memory address.
 - Pointer type declaration tells the compiler how many bytes to fetch on each dereference.
- But we often want “word alignment”:
 - Some processors will not allow you to address **32b** values without being on **4-byte** boundaries.
 - Others will just be very slow if you try to access “unaligned” memory.

| | | | | |
|------------|-------------------------------------|---------|--------|----|
| 0xFFFFFFF8 | int32_t * | | | |
| 0xFFFFFFF4 | | short * | | |
| 0xFFFFFFF0 | | | char * | |
| 0xFFFFFEC | 32-bit integer stored in four bytes | | | |
| 0xFFFFFE8 | 16-bit short | xx | xx | xx |
| ... | 8-bit char | xx | xx | xx |
| 0x0C | xx | xx | xx | xx |
| 0x08 | xx | xx | xx | xx |
| 0x04 | xx | xx | xx | xx |
| 0x00 | xx | xx | xx | xx |

4 bytes

Structures, Revisited, Again

- A "struct" is really just an instruction to C on how to arrange a bunch of bytes in a bucket.

- Structs provide enough space for the data.
- C will align the data with padding.

```
struct foo {  
    int32_t a;  
    char b;  
    struct foo *c;  
}
```



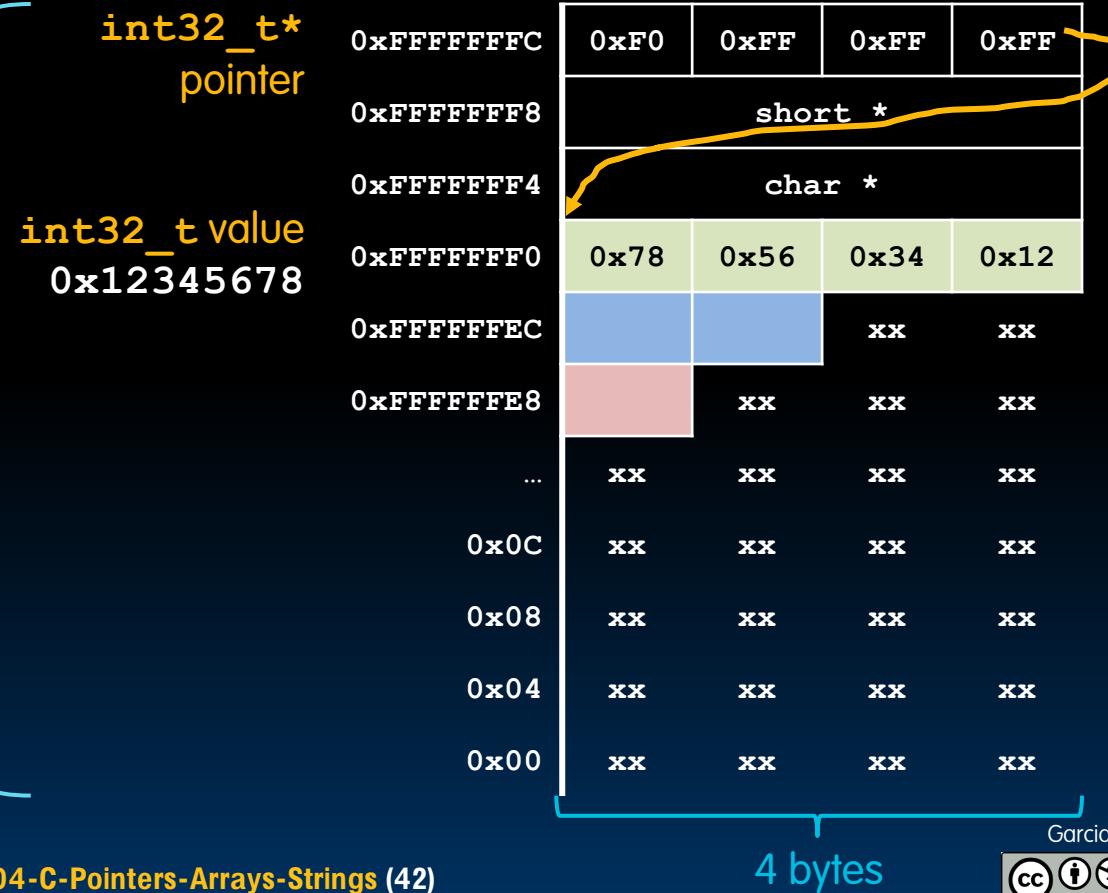
- For this struct, the actual layout on a 32b architecture would be as follows.
 - Note the 3 bytes of padding
 - `sizeof(struct foo) == 12`

| | | | |
|--------------|---------------|--------|--------|
| | 4 bytes for a | | |
| 1 byte for b | unused | unused | unused |
| | 4 bytes for c | | |

When debugging, watch out for endianness

- The hive machines are “little endian”.
 - The least significant byte of a value is stored first.
- (Contrast with “big endian”)
 - (Most significant byte is stored first)

Check out HW02 guide video,
to be released soon





And in Conclusion...

- Pointers and arrays are pretty much the same
- C knows how to increment pointers
- C is an efficient language, with little protection
 - Array bounds not checked
 - Variables not automatically initialized
- Use handles to change pointers
- (Beware) The cost of efficiency is more overhead for the programmer.
 - "C gives you a lot of extra rope, don't hang yourself with it!"



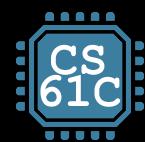
Garcia, Yan

Function Pointer Example (Recorded)

Recording:

<https://www.youtube.com/watch?v=P1IYbp0fqY4>

- Pointers
- Using Pointers Effectively
- Pointer Arithmetic
- Array Pitfalls
- Strings
- Word Alignment
- Function Pointer Example
(recorded)



map (actually mutate_map easier)

```
#include <stdio.h>

int x10(int), x2(int);
void mutate_map(int [], int n, int(*)(int));
void print_array(int [], int n);

int x2 (int n) { return 2*n; }
int x10(int n) { return 10*n; }

void mutate_map(int A[], int n, int(*fp)(int)) {
    for (int i = 0; i < n; i++)
        A[i] = (*fp)(A[i]);
}

void print_array(int A[], int n) {
    for (int i = 0; i < n; i++)
        printf("%d ",A[i]);
    printf("\n");
}
```

```
% ./map
3 1 4
6 2 8
60 20 80
```

```
int main(void)
{
    int A[] = {3,1,4}, n = 3;
    print_array(A, n);
    mutate_map (A, n, &x2);
    print_array(A, n);
    mutate_map (A, n, &x10);
    print_array(A, n);
}
```