



UC Berkeley  
Teaching Professor  
Dan Garcia

# CS61C

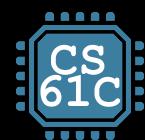
## Great Ideas in **Computer Architecture** (a.k.a. Machine Structures)



UC Berkeley  
Teaching Professor  
Lisa Yan

## Memory (Mis)Management

内存管理



# What gets printed?

*sizeof( )*

**sizeof()**: compile-time operator; gives size in bytes (of type or variable).

```
// for this exercise, assume
// shorts are 16b, pointers are 32b
void mystery(short arr[], int len) {
    printf("%d ", len);
    printf("%d\n", sizeof(arr));
}
```

```
int main() {
    short nums[] = {1, 2, 3, 99, 100};
    printf("%d ", sizeof(nums));           total size = 10
    mystery(nums, sizeof(nums)/sizeof(short));
    return 0;
}
```

*字节*

- A. 10 5 10
- B. 10 5 4
- C. 80 5 80
- D. 80 5 32
- E. Other

# What gets printed?

**sizeof()**: compile-time operator; gives size in bytes (of type or variable).

```
// for this exercise, assume
// shorts are 16b, pointers are 32b
void mystery(short arr[], int len) {
    printf("%d ", len);
    printf("%d\n", sizeof(arr));
}
```

- A. 10 5 10
- B. 10 5 4
- C. 80 5 80
- D. 80 5 32
- E. Other

Array has decayed  
to a pointer

```
int main() {
    short nums[] = {1, 2, 3, 99, 100};
    printf("%d ", sizeof(nums));
    mystery(nums, sizeof(nums)/sizeof(short));
    return 0;
}
```

In array's declared scope,  
array size total array size.  
# elements in array.

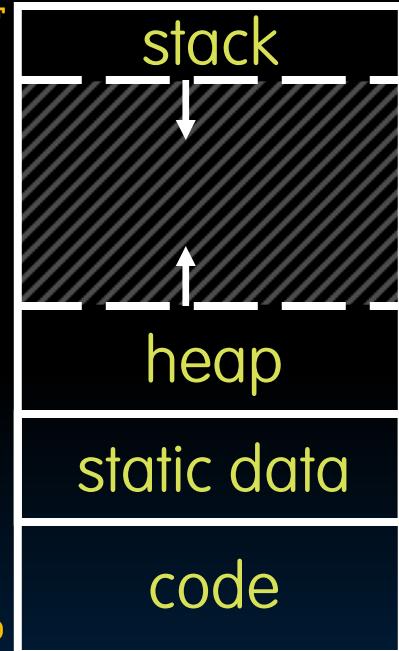
# Memory Locations

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management

# C Program Address Space

地址空间

- A program's address space contains 4 regions:
  - Stack**: local variables inside functions, grows downward
  - Heap**: space requested for pointers via `malloc()`; resizes dynamically, grows upward *动态调整*
  - Static Data**: variables declared outside `main`, does not grow or shrink
  - Code (aka Text)**: loaded when program starts, does not change
  - 0x00000000 chunk is unwriteable/unreadable** so you that crash on `NULL` pointer access
- Programming in C requires knowing where objects are in memory, otherwise things don't work as expected.
  - By contrast, Java hides location of objects.



For now, OS somehow prevents accesses between stack and heap (gray hash lines). Wait for virtual memory!

Garcia, Yan

# Where are variables allocated?

- If declared **outside** a function (**global**) allocated in “**static**” storage.
- If declared **inside** function (**local**) allocated on the “**stack**” and freed when function returns.
  - NB: `main()` is a function
- For both these memory types, the management is automatic.
  - You don’t need to worry about deallocated when you are no longer using them.
  - But a variable **does not exist anymore** once a function ends!

```
int myGlobal;  
main() {  
    int myTemp;  
}
```

static Data  
Stack

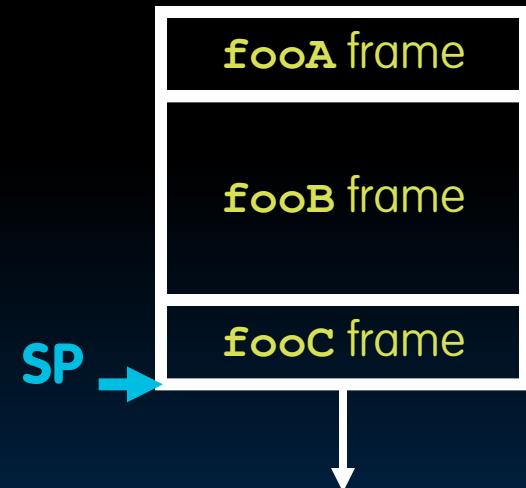
# The Stack

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management



- Every time a function is called, a new “stack frame” is allocated on the stack.
- Stack frame includes:
  - Return “instruction” address (who called me?)
  - Arguments
  - Space for other local variables
- Stack frames contiguous blocks of memory; stack pointer indicates start of stack frame.
- When function ends, stack frame is tossed off the stack; frees memory for future stack frames.
- Later, we'll cover details for RISC-V processor.

```
fooA() { fooB(); }  
fooB() { fooC(); }  
fooC() { ... }
```

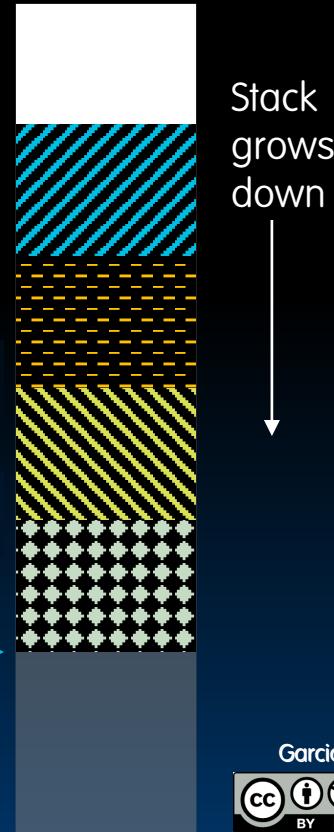


# The Stack is Last In, First Out (LIFO)

```
int main ()  
{ a(0); ...  
}  
void a (int m)  
{ b(1);  
}  
void b (int n)  
{ c(2);  
}  
void c (int o)  
{ d(3);  
}  
void d (int p)  
{  
}
```

Last in First Out

Stack Pointer →



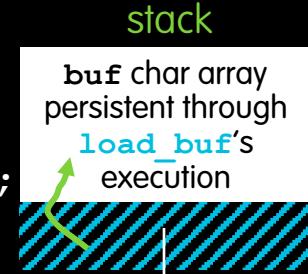
# Passing Pointers into the Stack



It is fine to pass a pointer to stack space further down.

```
#define BUFSIZE 256
int main() {
    ...
    char buf[BUFSIZE];
    load_buf(buf, BUFSIZE);
    ...
}
```

buf is pointer

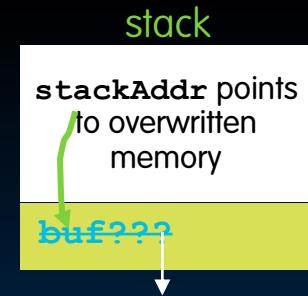


However, it is catastrophically bad to return a pointer to something in the stack!

- Memory will be overwritten when other functions are called!
- So your data would no longer exist...and writes can overwrite key pointers, causing crashes!

```
char *make_buf() {
    char buf[50];
    return buf;
}

int main() {
    ...
    char *stackAddr = \
        make_buf();
    foo();
    ...
}
```



# The Heap

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management

动态

# What is the Heap?

动态内存

- The heap is **dynamic memory** – memory that can be allocated, resized, and freed during program runtime.
  - Where Java `new` command allocates memory
  - Useful for persistent memory across function calls
  - But biggest source of pointer bugs, memory leaks, ...
- Large pool of memory, not allocated in contiguous order**
  - Back-to-back requests for heap memory *could* result in blocks very far apart
- In C, specify number of bytes of memory explicitly to allocate item**
  - `malloc()`: Allocates raw, uninitialized memory from heap
  - `free()`: Frees memory on heap
  - `realloc()`: Resizes previously allocated heap blocks to new size
  - Unlike the stack, memory gets reused only when programmer **explicitly** cleans up

# void \*malloc(size\_t n)

- Allocates a block of uninitialized memory:

- size\_t n is an unsigned integer type big enough to "count" memory bytes.
- Returns `void *` pointer to block of memory on heap.
- A return of `NULL` indicates no more memory (**always** check for it!!!)

- To allocate a struct:

```
typedef struct { ... } TreeNode;  
TreeNode *tp = (TreeNode *) malloc(sizeof(TreeNode));
```

Typecast casts return value  
from type `(void *)` to `(TreeNode *)`

give size in bytes

`sizeof(type)` gives size in bytes.

- To allocate an array of 20 ints:

```
int *ptr = (int *) malloc(20*sizeof(int));
```

- Many years ago ints used to be 16b. Now, 32b or 64b...

Assuming size of objects can lead to misleading, unportable code. Use `sizeof()`!

a pointer pointing at one integer which is the first one!



# void free(void \*ptr)

- **Dynamically frees heap memory**
  - `ptr` is a pointer containing an address originally returned by `malloc()`/`realloc()`.

```
int *ptr = (int *) malloc (sizeof(int)*20);  
...  
free(ptr); // implicit typecast to (void *)
```

- **When you free memory, be sure to pass the original address returned from `malloc()`. Otherwise, crash (or worse!)**

# void \*realloc(void \*ptr, size\_t size)

- Resize a previously allocated block at ptr to a new size.
  - Returns new address of the memory block.
    - In doing so, it may need to copy all data to a new location.
  - `realloc(NULL, size); // behaves like malloc` from 0 to 5th
  - `realloc(ptr, 0); // behaves like free, deallocated heap block`
- Remember: Always check for return `NULL`, which would mean you've run out of memory!

```
int *ip; ip = (int *) malloc(10*sizeof(int));  
... ... ... /* check for NULL */  
ip = (int *) realloc(ip, 20*sizeof(int));  
/* contents of first 10 elements retained */  
... ... ... /* check for NULL */  
realloc(ip, 0); /* equivalent to free(ip); */
```

expand the  
memory

# Linked List Example

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management

# Linked List Example

```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}
void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```

```
struct Node {
    char *data;
    struct Node *next;
};
```

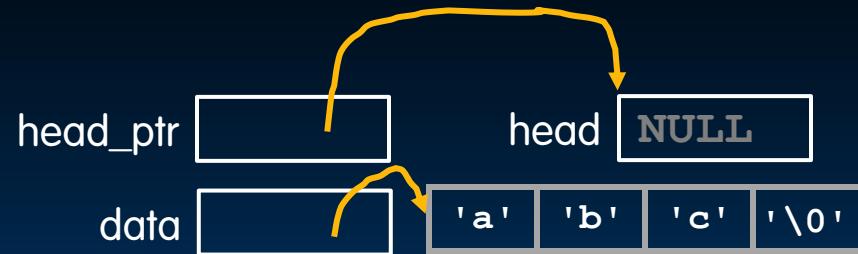
*the address of pointer*

*strcpy( , )*

*String copy function*

# Linked List Example

```
# include <string.h>
int main() {
    1   struct Node *head = NULL;
    2   add_to_front(&head, "abc");
    ... // free nodes, strings here...
}
void add_to_front(struct Node **head_ptr, char *data) {
    3   struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    4   node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    5   strcpy(node->data, data); // strcpy also copies null terminator
    6   node->next = *head_ptr;
    7   *head_ptr = node;
}
```



# Linked List Example

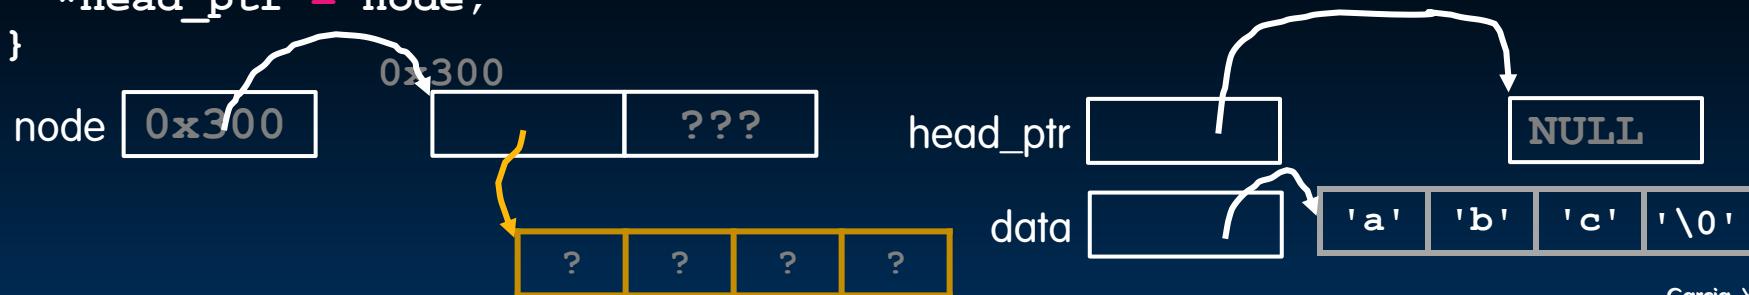
```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```



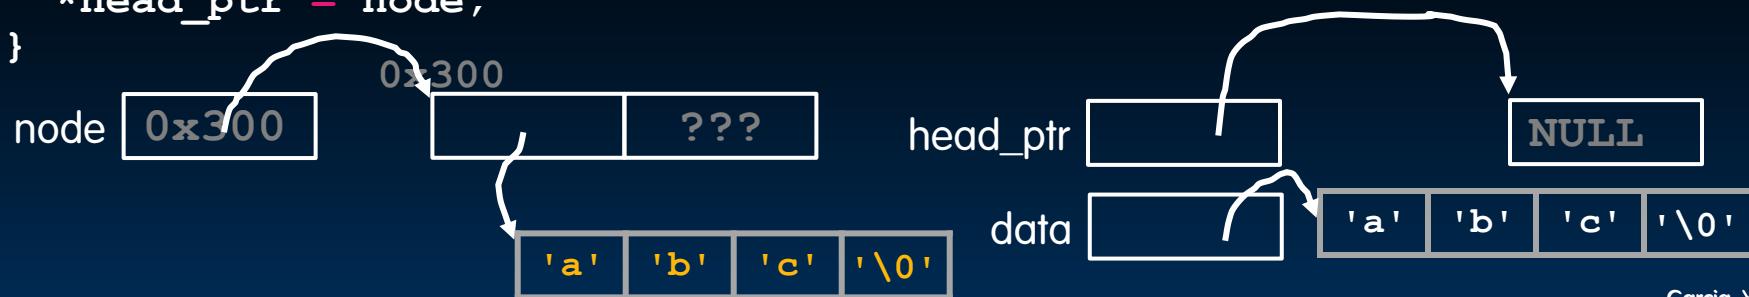
# Linked List Example

```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}
void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```



# Linked List Example

```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}
void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```



# Linked List Example

```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}

void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```



# Linked List Example

```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}
void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```



# Linked List Example

```
# include <string.h>
int main() {
    struct Node *head = NULL;
    add_to_front(&head, "abc");
    ... // free nodes, strings here...
}
void add_to_front(struct Node **head_ptr, char *data) {
    struct Node *node = (struct Node *) malloc(sizeof(struct Node));
    node->data = (char *) malloc(strlen(data) + 1); // extra byte!
    strcpy(node->data, data); // strcpy also copies null terminator
    node->next = *head_ptr;
    *head_ptr = node;
}
```

Check out the lecture code in the drive!



# When Memory Goes Bad

- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management

# Working with Memory

*Code and static data never change*

## Code, Static storage are easy:

- They never grow or shrink.

## Stack space is also easy:

- Stack frames are created and destroyed in LIFO order.
- Just avoid "**dangling references**": pointers to deallocated variables (e.g., from old stack frames).

## ⚠ Working with the heap is tricky:

- Memory can be allocated / deallocated at any time!

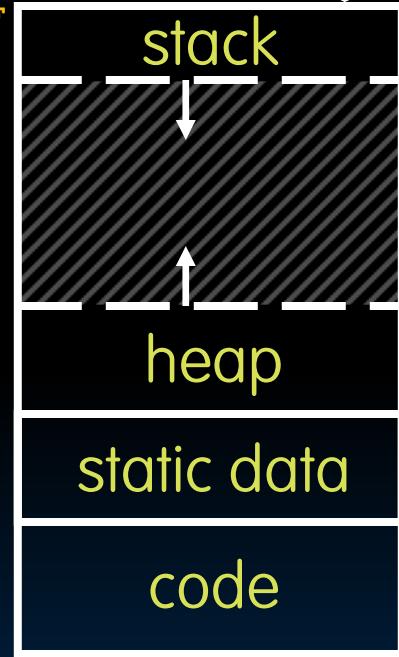
*overflow*

- Memory leak**: If you forget to deallocate memory } Your program will eventually run out of memory

- "Use after free"**: If you use data after calling free }

- "Double free"**: If you call free 2x on same memory }

0xFFFF FFFF



0x0000 0000

Possible crash or  
exploitable vulnerability

Garcia, Yan





- The runtime **does not check for the programmer's failure to manage memory.**
  - Memory is so performance-critical that there just isn't time to do this.
  - Usual result: you corrupt the memory allocator's internal structure, and you find out much later in a totally unrelated part of your code!
- **Memory leak: Failure to free () allocated memory**
  - ⚠ Initial symptoms. Nothing...
    - Until you hit a critical point, memory leaks aren't actually a problem
  - ⚠ ...Later symptoms: performance drops off a cliff...
    - Memory hierarchy behavior tends to be great just up until it isn't, then it hits several cliffs
  - ⚠ ...and then your program is killed off!
    - Because the operating system (OS) says "no" when you ask for more memory



- “Dangling reference”

**When you keep using a pointer,  
even after it has been deallocated**

- **Reads after the free may be corrupted!**

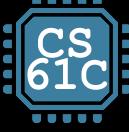
- If something else takes over that memory, your program will probably read the wrong information!

- **Writes corrupt other data!**

- Uh oh... Your program crashes later!

```
struct foo *f;  
...  
f = malloc(sizeof(struct foo));  
...  
free(f);  
...  
bar(f->a); // !!!
```

*mother-fucker*



# Double-Free...



```
struct foo *f = (struct foo *) malloc(10*sizeof(struct foo));  
...  
free(f);  
...  
free(f); // !!!
```

*f doesn't exist any more.*

- May cause either a **use-after-free** (because something else called **malloc()** and got that data) or **corrupt heap data** (because you are no longer freeing a pointer tracked by **malloc**)

# Forgetting realloc() Can Move Data



- “Dangling reference”
- Remember, when you realloc it can copy data to a different part of the heap.

```
int *nums;  
nums = malloc(10*sizeof(int));  
...  
int *g = nums;  
...  
nums = realloc(nums,  
               20*sizeof(int));  
  
// g could now point  
// to invalid memory
```

```
int *nums;  
nums = malloc(10*sizeof(int));  
...  
// forget to update nums  
// on realloc call  
realloc(nums, 20*sizeof(int));  
  
// nums could now point  
// to invalid memory,  
// and we could have potentially  
lost a pointer to a new block
```

- Reads may be corrupted, and writes may corrupt other pieces of memory.

We have to define the new pointer

## How many memory management errors are in this code?

```
void free_mem_x() {  
    int fnh[3];  
    ...  
    free(fnh); X  
}  
  
void free_mem_y() {  
    int *fum = malloc(4*sizeof(int));  
    free(fum+1); X  
    ...  
    free(fum);  
    ...  
    free(fum); X  
}
```

- A. 1
- B. 2
- C. 3
- D. 0
- E. Other

When poll is active, respond at [pollev.com/yaml](http://pollev.com/yaml)  
Text **YANL** to 22333 once to join

Respond at [pollev.com/yanl](http://pollev.com/yanl)

Text **YANL** to **22333** once to join, then **A, B, C, D, or E**

# How many memory management errors are in this code?



# Faulty Heap Management

How many memory management errors are in this code?

```
void free_mem_x() {  
    int fnh[3];  
    ...  
    free(fnh);  
}
```

*fnh is not allocated in heap*

*} free() on stack-allocated memory*

- A. 1
- B. 2
- C. 3
- D. 0
- E. Other

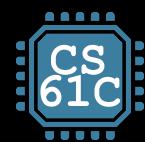
```
void free_mem_y() {  
    int *fum = malloc(4*sizeof(int));  
    free(fum+1);  
    ...  
    free(fum);  
    ...  
    free(fum);  
}
```

*} free() on memory that isn't the  
pointer from malloc*

*fum+1* is not allocated by  
*malloc()*

*} Double free()*

*free() for twice*



# Valgrind to the rescue!!!

- **Valgrind slows down your program by an order of magnitude, but...**
  - It adds a tons of checks designed to catch most (but not all) memory errors
    - Memory leaks
    - Misuse of free
    - Writing over the end of arrays
- **Tools like Valgrind are absolutely essential for debugging C code.**

Check out Lab 02!

# Implementing Memory Management

Material not tested. Recording:

<https://www.youtube.com/watch?v=Sq5tSeWfnGY>

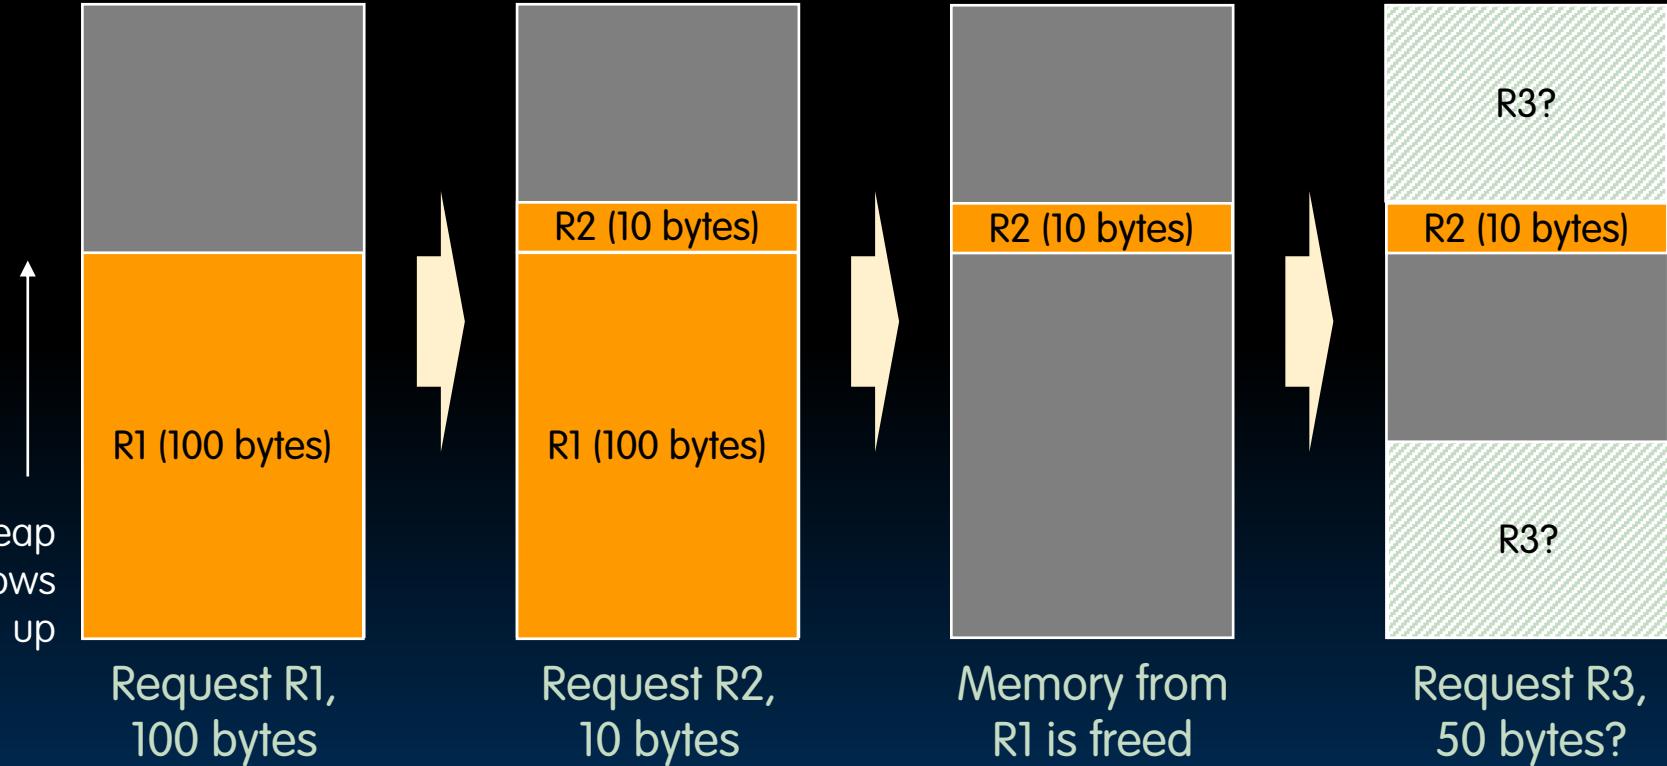
- Memory Locations
- The Stack
- The Heap
- Linked List Example
- When Memory Goes Bad
- Implementing Memory Management

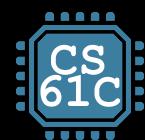
# Heap Management Requirements

- Want `malloc()` and `free()` to run quickly
- Want minimal memory overhead
- Want to avoid fragmentation\*,   
when most of our free memory is in many small chunks
  - In this case, we might have many free bytes but not be able to satisfy a large request since the free bytes are not contiguous in memory.

\* This is technically external fragmentation

# Heap Management Example





# K&R Malloc/Free Implementation

- **From Section 8.7 of K&R** ↗ 简介  
▫ Code in the book uses some C language features we haven't discussed and is written in a very terse/style; don't worry if you can't decipher the code
- **Each block of memory is preceded by a header that has two fields:**
  - size of the block, and
  - a pointer to the next block
- **All free blocks are kept in a circular linked list.**
- **In an allocated block, the header's pointer field is unused.**

# K&R Malloc/Free Implementation

- **malloc()** searches the free list for a block that is big enough. If none is found, more memory is requested from the operating system. If what it gets can't satisfy the request, it fails.
- **free()** checks if the blocks adjacent to the freed block are also free.
  - If so, adjacent free blocks are merged (**coalesced**) into a single, larger free block.
  - Otherwise, freed block is just added to the free list.

# Choosing a block in malloc()

- If there are multiple free blocks of memory that are big enough for some request, how do we choose which one to use?
  - **best-fit**: choose the smallest block that is big enough for the request.
  - **first-fit**: choose the first block we see that is big enough.
  - **next-fit**: like first-fit, but remember where we finished searching and resume searching from there.



# And in Conclusion...

- **C has 3 pools of memory**
  - Static storage: global variable storage, basically permanent, entire program run
  - The Stack: local variable storage, parameters, return address
  - The Heap (dynamic storage): malloc() grabs space from here, free() returns it.
- **Heap data is biggest source of bugs in C code**
- **malloc () handles free heap space with freelist. Take CS162 for more!**



Garcia, Yan