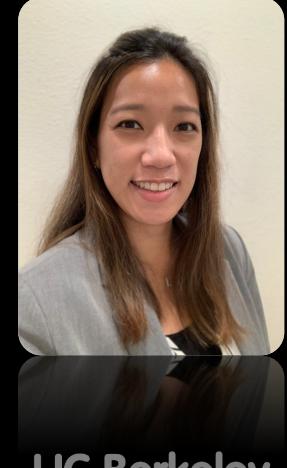




UC Berkeley  
Teaching Professor  
Dan Garcia

# CS61C

## Great Ideas in Computer Architecture (a.k.a. Machine Structures)

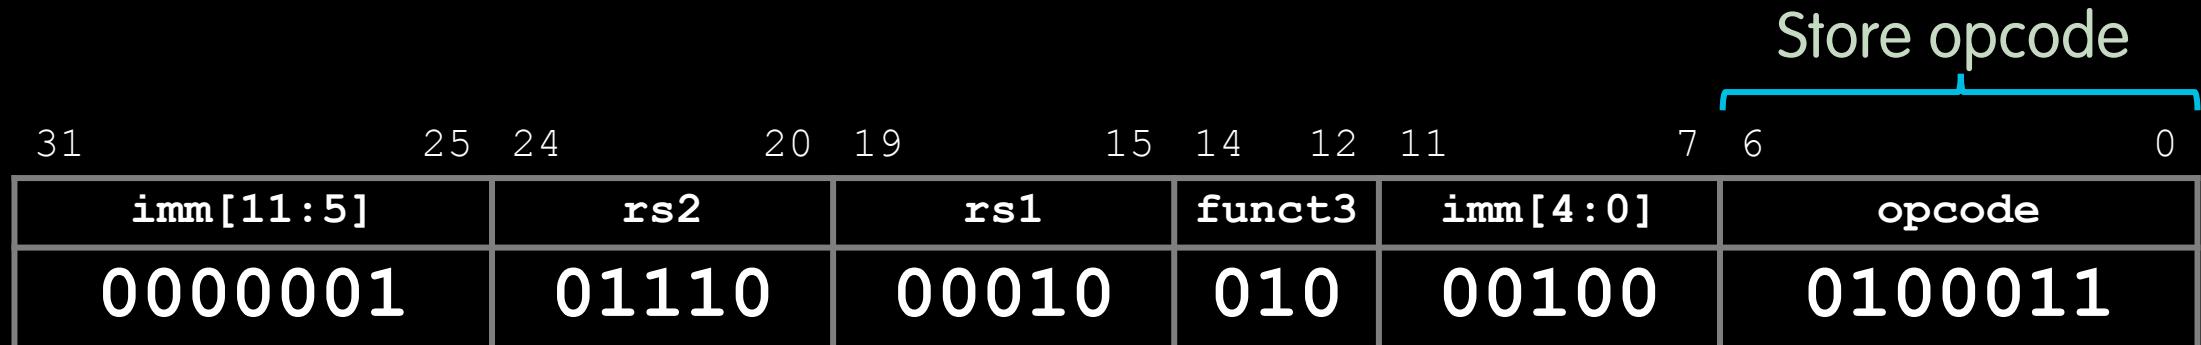


UC Berkeley  
Teaching Professor  
Lisa Yan

## RISC-V Instruction Formats, Part II

# Disassembling the S-Format

- What is the corresponding assembly instruction?

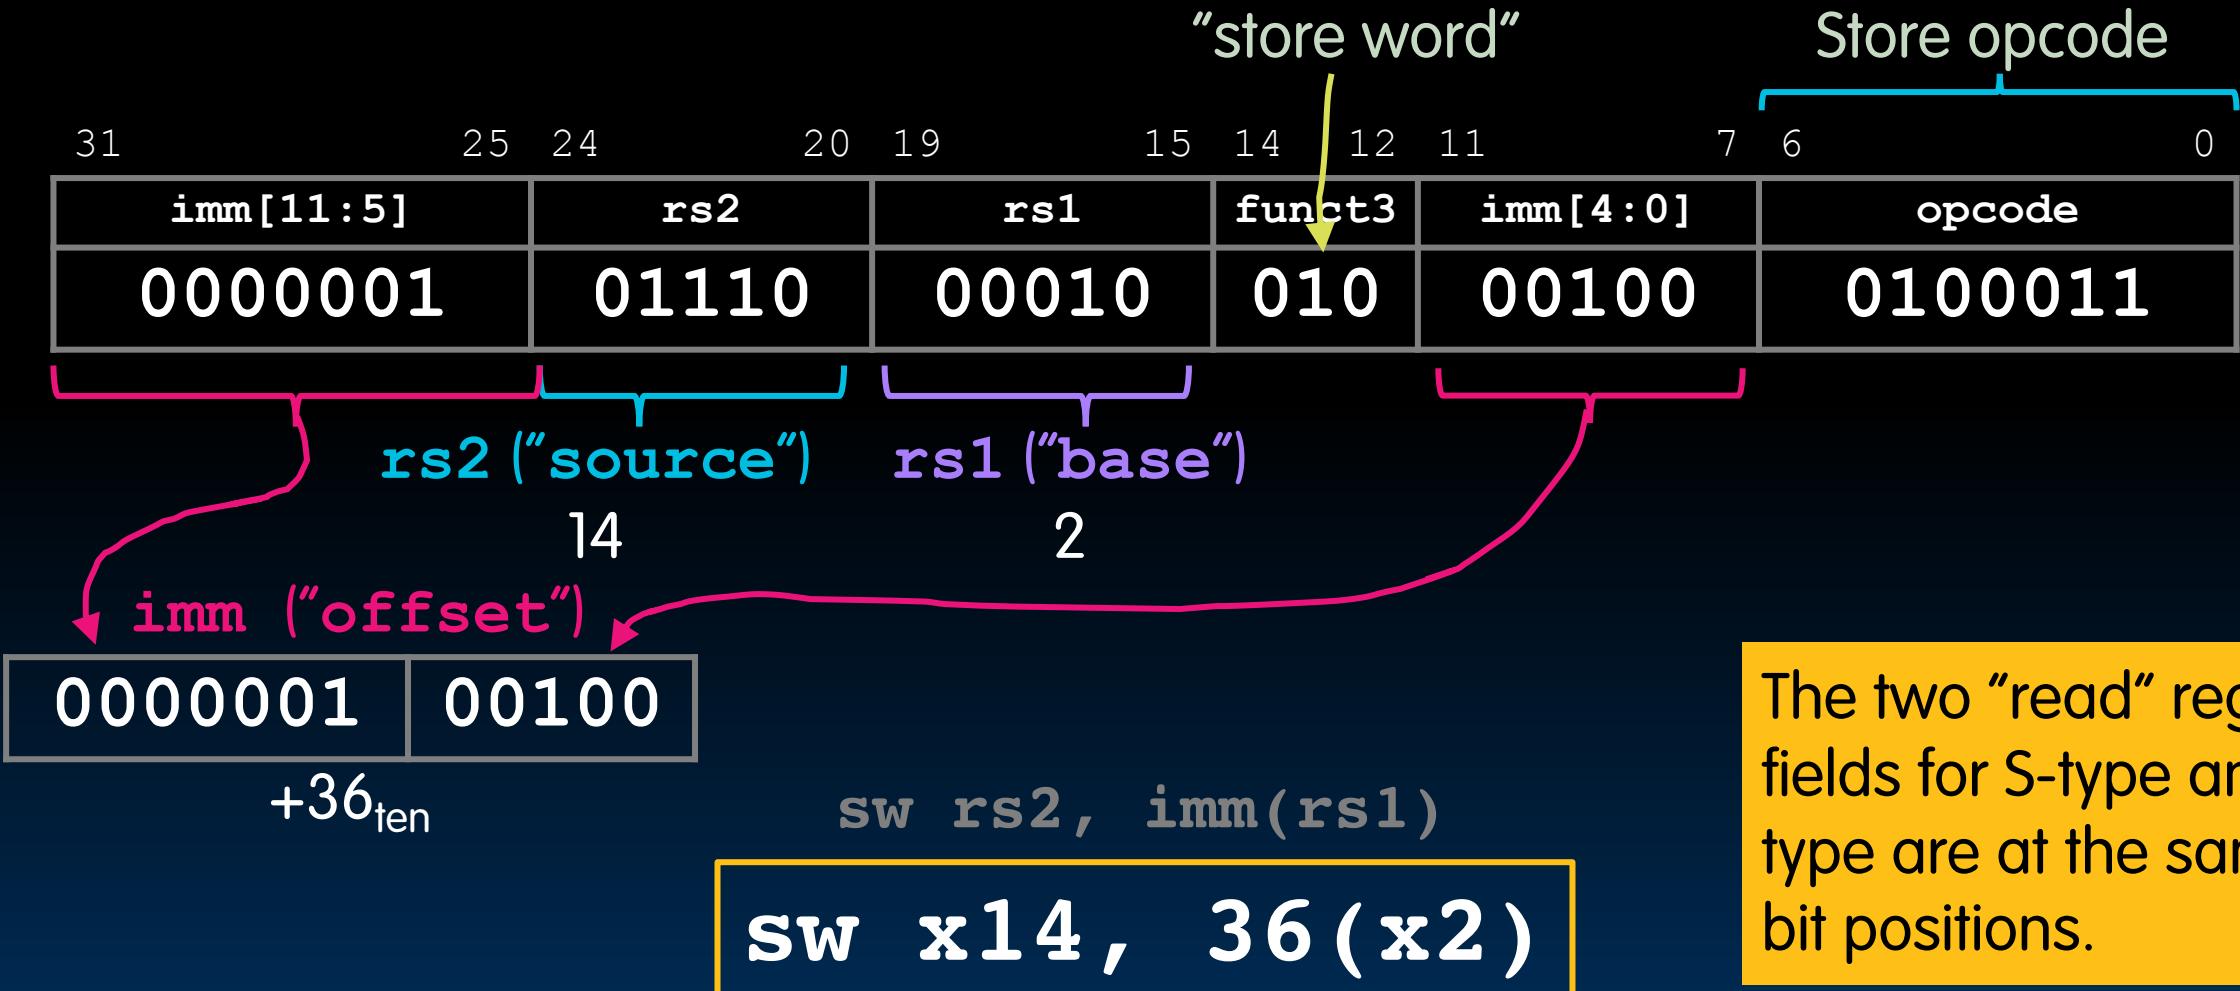


S-Format Lookup table:

imm[11:5]	rs2	rs1	funct3	opcode	
imm[11:5]	rs2	rs1	000	imm[4:0]	0100011 sb
imm[11:5]	rs2	rs1	001	imm[4:0]	0100011 sh
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011 sw

# Disassembling the S-Format

- What is the corresponding assembly instruction?



# **VIEW FROM THE TOP**

***Conversations with Global Leaders***

## **Driving Moore's Law:**

A Fireside Chat with Ann Kelleher and  
Dean Tsu-Jae King Liu

**ANN KELLEHER**

*EXECUTIVE VICE PRESIDENT &  
GENERAL MANAGER OF TECHNOLOGY DEVELOPMENT  
INTEL*

**Wednesday, Sept. 21 | 12-1 p.m.  
Banatao Auditorium, Sutardja Dai Hall**

Event is free and open to the Berkeley community.  
Lunch will be provided.

Co-hosted with the Society of Women Engineers  
[engineering.berkeley.edu/events](http://engineering.berkeley.edu/events)



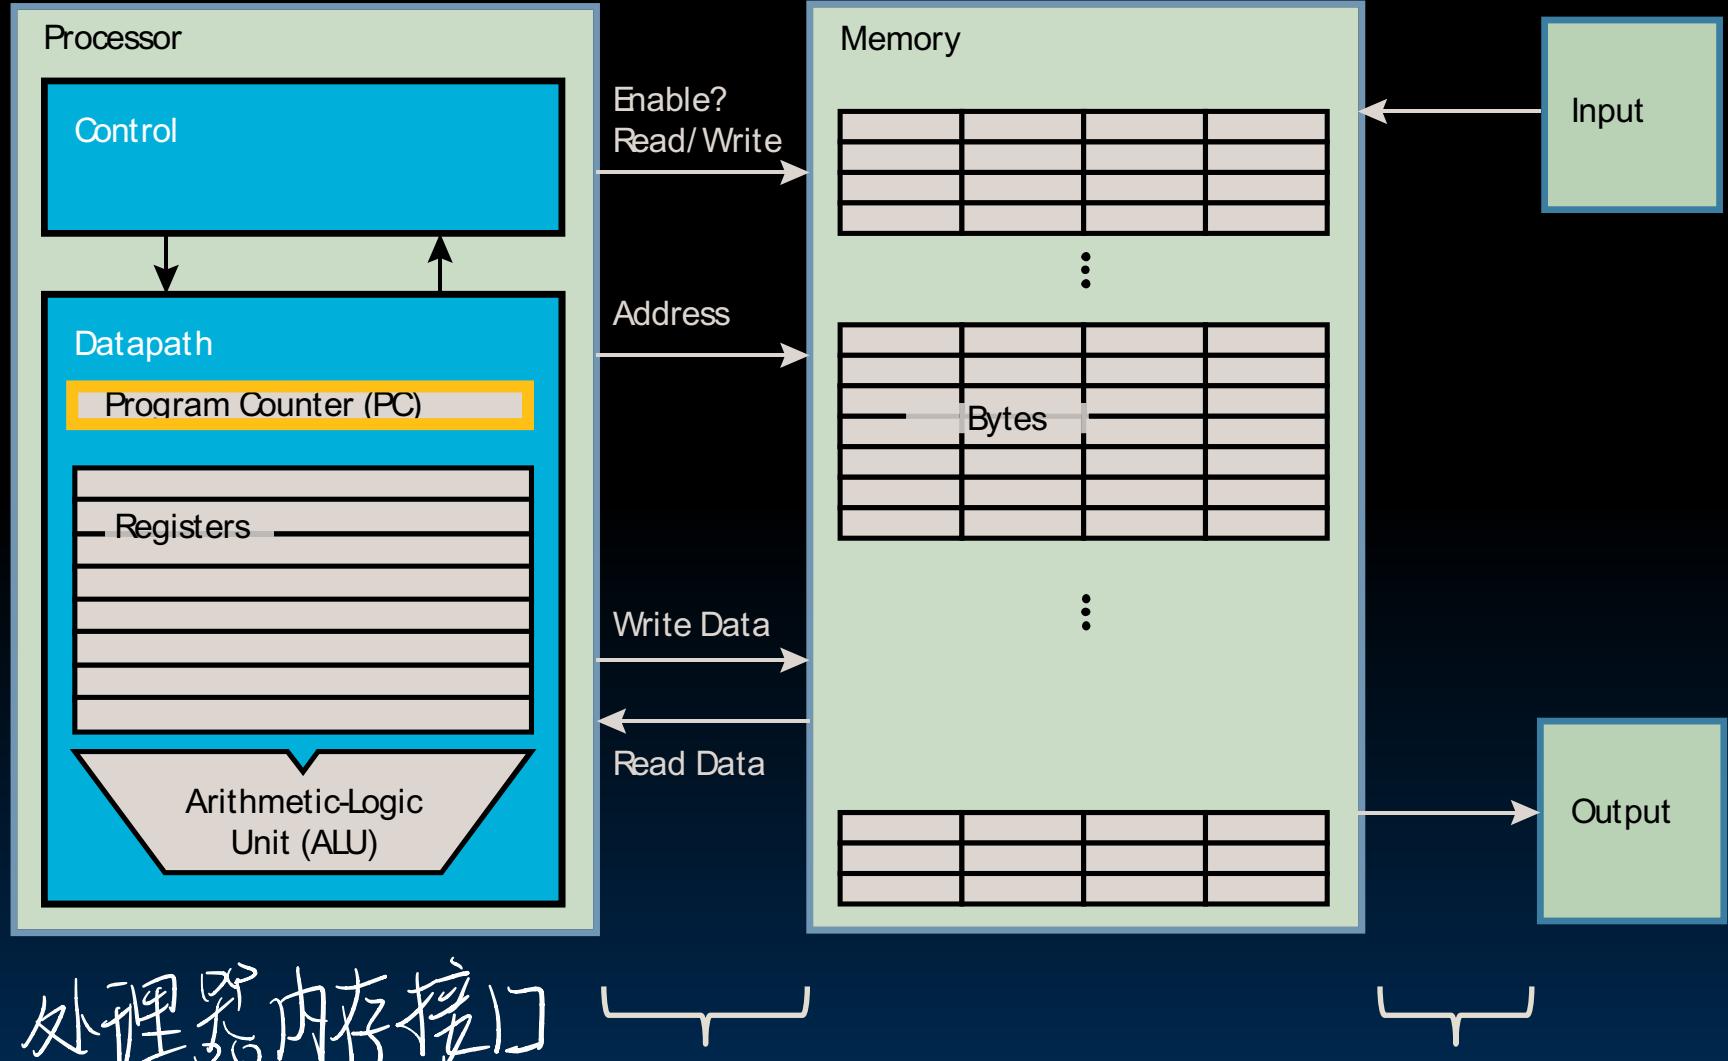
**Berkeley**  
**Engineering**

# Updating the Program Counter (PC)

- Updating the Program Counter (PC)
- B-Format Layout
- J-Format Layout
- U-Format: Long Immediates
- `jalr`: I-Format

# The Program Counter (PC)

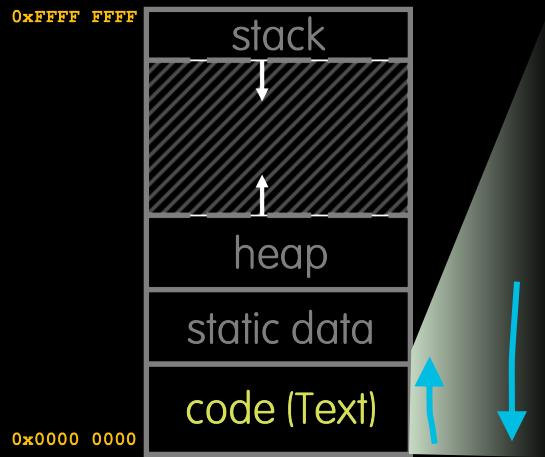
程序计数器



**The Program Counter (PC)** is a special internal register inside processor holding byte address of the instruction being executed.

It is separate from the 32-registers Register File!

# PC = PC + 4 for most instructions



Program Address Space  
程序地址空间

程序地址空间

Address	Machine Code	Original Assembly Code
0x0c	00A98863	Loop: beq x19,x10,End
0x10	00A90933	add x18,x18,x10
0x14	FFF98993	addi x19,x19,-1
0x18	FF5FF06F	j Loop
0x1c	00A98863	End:
		...

machine code correspond to assembly code

Address	Machine Code	Original Assembly Code
0x0c	00A98863	Loop: beq x19,x10,End
0x10	00A90933	add x18,x18,x10
0x14	FFF98993	addi x19,x19,-1
0x18	FF5FF06F	j Loop
0x1c	00A98863	End:
		...

Registers before execution

PC	0000 0010
RegFile	
x10	0000 0001
x18	0000 0000
x19	0000 0002

+4

PC	0000 0014
RegFile	
x10	0000 0001
x18	0000 0001
x19	0000 0002

Garcia, Yan



# Branches Update PC to "Jump"

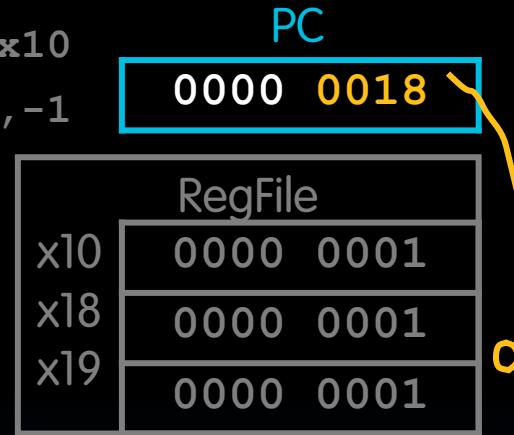


Kris Kross - Jump (Official Video) -

## Unconditional Branches

```

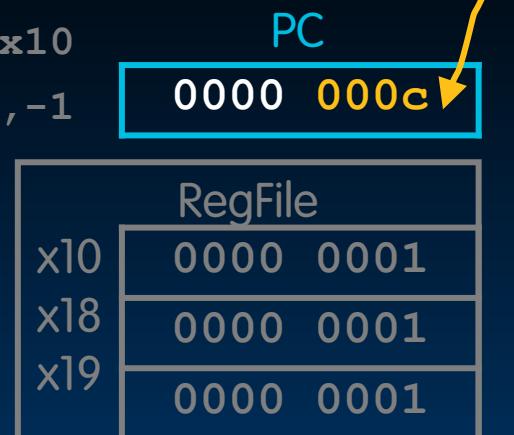
0x0c | Loop: beq x19,x10,End
      add x18 x18,x10
      addi x19,x19,-1
      j Loop
      End: ...
0x10
0x14
0x18
0x1c
  
```



**0x18** → **0x0c**

```

0x0c | Loop: beq x19,x10,End
      add x18,x18,x10
      addi x19,x19,-1
      j Loop
      End: ...
0x10
0x14
0x18
0x1c
  
```



-12

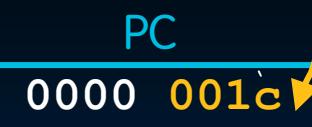
## Conditional Branch (if branch taken)

→ 0x0c | Loop: beq x19,x10,End  
0x10  
0x14  
0x18  
0x1c | End: ...  
j Loop



+16

0x0c | Loop: beq x19,x10,End  
0x10  
0x14  
0x18  
0x1c | End: ...  
j Loop



If branch not taken,  
then  $PC = PC + 4$

# PC-Relative Addressing

相对地址

- PC-Relative Addressing: Supply a signed offset to update PC.
  - $PC = PC + \text{byte\_offset}$
  - “Position-Independent Code”: If all of code moves, *relative offsets* don’t change!

0x0c	Loop:	<code>beq x19, x10, End</code>
0x10		<code>add x18, x18, x10</code>
0x14		<code>addi x19, x19, -1</code>
0x18		<code>j Loop</code>
0x1c	End:	...

Take branch:  
 $PC = PC + 16$   
 Don't take branch:  
 $PC = PC + 4$

0x400	Loop:	<code>beq x19, x10, End</code>
0x404		<code>add x18, x18, x10</code>
0x408		<code>addi x19, x19, -1</code>
0x40c		<code>j Loop</code>
0x410	End:	...

- Branches generally change the PC by a small amount, therefore in RISC-V instructions we encode relative offsets as *signed immediates*.

- Contrast with: Absolute Addressing 绝对地址

- Supply new address to overwrite PC.  $PC = \text{new\_address}$  手写
- Use sparingly: Brittle to code movement/need to build 32-bit immediate (more later)

# B-Format Layout

- Updating the Program Counter (PC)
- B-Format Layout
- J-Format Layout
- U-Format: Long Immediates
- `jalr`: I-Format

# RISC-V Conditional Branches

- Example: `beq rs1, rs2, Label`
- Conditional branches need:
  - Two **read** registers, **no destination** register (like S-Format)
  - Some way to encode label (i.e., where to jump to)
- RISC-V uses PC-relative addressing to encode labels.
  - If we *don't* branch:  $\text{PC} = \text{PC} + 4 \text{ bytes}$
  - If we *do* branch:  $\text{PC} = \text{PC} + (\text{relative byte offset to Label})$
- B-Format Fields:
  - 7b (**opcode**) + 3b (**funct3**) + 5b (**rs1**) + 5b (**rs2**) = 20b
  - 12b immediate field to represent PC's relative offset. 偏移量
    - Signed two's complement means we can represent  $\pm 2^{11}$  "units" from the PC.

二进制补码

The 12-bit immediate field for conditional branches is:

**A. In units of 1 byte.**

- One branch reaches  $\pm 2^{11}$  bytes from PC  $\rightarrow \pm 2^9 \times 32\text{-bit instructions}$ .
  - Never branch into middle of instruction.

**B. In units of 2 bytes (16-bit “half-words”).**

- One branch reaches  $\pm 2^{10} \times 32\text{-b instructions}$  from PC.
  - Multiply the offset by 2 before adding to the PC.

**C. In units of 4 bytes (32-bit words).**

- One branch reaches  $\pm 2^{11} \times 32\text{-b instructions}$  from PC.
  - Multiply the offset by 4 before adding to the PC.

[pollev.com/yan1](http://pollev.com/yan1)



# How do branches encode PC relative offset?



## Scaling the Offset to Maximize Branch Range

The 12-bit immediate field for conditional branches is:

A. In units of 1 byte.

- One branch reaches  $\pm 2^{11}$  bytes from PC  $\rightarrow \pm 2^9 \times 32\text{-bit instructions}$ .
  - Never branch into middle of instruction.

B. In units of 2 bytes.

- One branch reaches  $\pm 2^{10} \times 32\text{-b instructions}$  from PC.
  - Multiply the offset by 2 bytes before adding to the PC.

C. In units of 4 bytes.

- One branch reaches  $\pm 2^{11} \times 32\text{-b instructions}$  from PC.
  - Multiply the offset by 4 bytes before adding to the PC.

None of the above

A

B

C

The 12-bit immediate field for conditional branches is:

✗ In units of 1 byte.

- One branch reaches  $\pm 2^{11}$  bytes from PC  $\rightarrow \pm 2^9 \times 32\text{-bit instructions}$ .
  - Never branch into middle of instruction.
- *Wasteful!* Lower 2 bits would always be wasted (always 00).

✓ In units of 2 bytes (16-bit “half-words”). *half word*

- One branch reaches  $\pm 2^{10} \times 32\text{-b instructions}$  from PC.
  - Multiply the offset by 2 before adding to the PC.

✗ In units of 4 bytes (32-bit words).

- One branch reaches  $\pm 2^{11} \times 32\text{-b instructions}$  from
  - Multiply the offset by 4 before adding to the PC.
- Does not support 16-bit compressed instructions!

Branch immediate range  
[-2048, +2047]  
represents PC relative  
offsets [-4096, +4094] in  
2-byte increments.

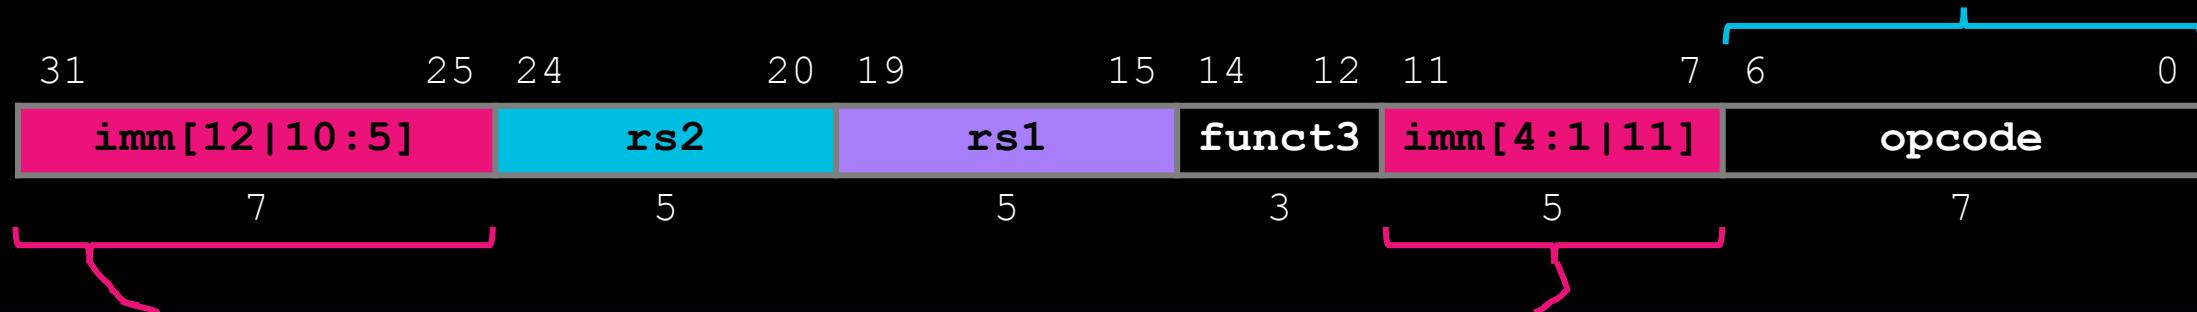
# RISC-V Feature: 16-bit instructions

- RISC-V Base ISA for RV32, RV64, RV128 all have 32-bit wide instructions, but it supports extensions: 32位指令可扩展
  - 16b compressed instructions
  - Variable-length instructions that are multiples of 16b in length
- RISC-V therefore *scales the branch offset by 2 bytes*, even when there are no 16b instructions!
- Implications for this class (which uses RISC-V processors that only support 32-bit instructions):
  - Half of the possible branch targets will be errors.
  - RISC-V Conditional Branches can only reach  $\pm 2^{10}$  32-bit instructions on either side of PC.

# B-Format Instruction Layout

- B-Format (textbook: SB-Type) close to S-Format:  
**opname      rs1 , rs2 , Label**

All *conditional/branch* instructions have opcode **1100011**.



Immediate represents relative offset in increments of 2 bytes ("half-words")

- To compute new PC:  $PC = PC + \text{byte\_offset}$ .
- 12 immediate bits imply  $\pm 2^{10}$  32-bit instructions reachable:
  - 1 bit: 2's complement (allow  $+\/-$  offset)
  - 1 bit: half-word/16-b instruction support

(Lowest bit of offset is always zero, so no need to store in instruction.)

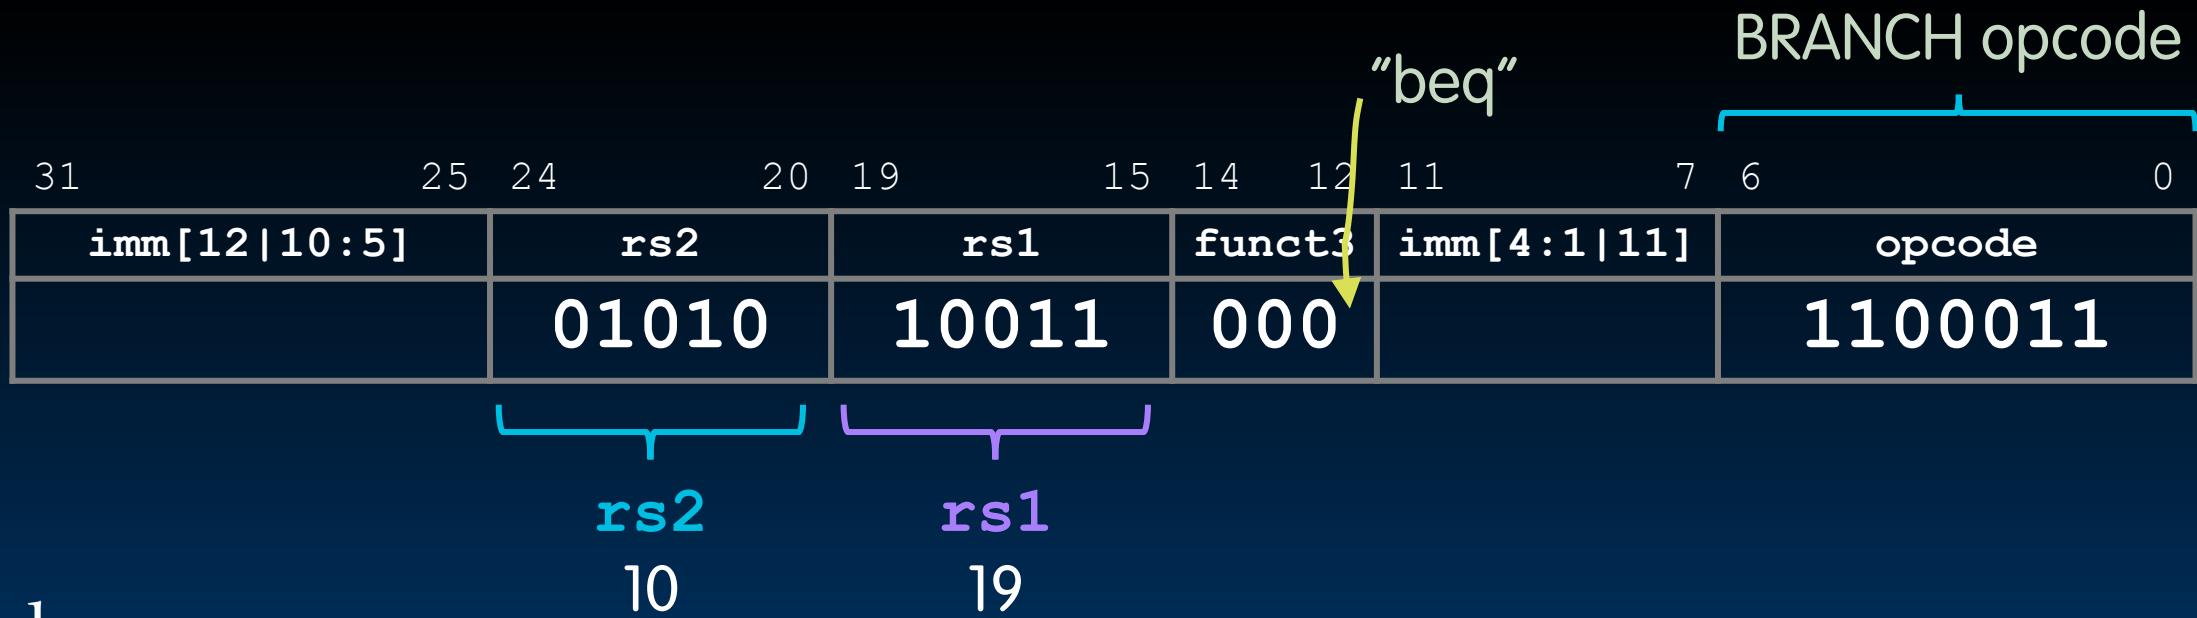


# B-Format Example (1/2)

```
Loop: beq x19,x10,End  
      add x18,x18,x10  
      addi x19,x19,-1  
      j Loop  
  
End: ... # target instr
```

How do we encode the  
beq instruction in this  
RISC-V code?

1. Fill in fields (other than immediate).



# B-Format Example (2/2)

+4  
instruct-  
ions  
Loop:

```
beq x19,x10,End
add x18,x18,x10
addi x19,x19,-1
j Loop
```

End: ... # target instr

How do we encode the  
beq instruction in this  
RISC-V code?

2. Construct branch offset (in bytes), then encode immediate.

31	25	24	20	19	15	14	12	11	7	6	0
imm[12 10:5]	rs2	rs1		funct3	imm[4:1 11]				opcode		
0000000	01010	10011		000	10000				1100011		

- +4 32-bit instructions = +16 byte offset

12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	0	0	0	0	0	1	0	0	0	0

13-bit signed **byte\_offset** ("imm")

# Why Swirl Immediate Bits Around? HW Design!

- Recall: RISC-V register field positions consistent across instr. formats.
- RISC-V also tries to keep bit positions of immediates consistent:

	31	25	24	20	19	15	14	12	11	7	6	0
I-type	imm[11:5]	imm[4:0]		rs1		funct3		rd		opcode		
	YXXXXXXX	WWWWV										
S-type	imm[11:5]		rs2		rs1	funct3	imm[4:0]		opcode			
	YXXXXXXX						WWWWV					
B-type	imm[12 10:5]		rs2		rs1	funct3	imm[4:1 11]		opcode			
	ZXXXXXXX						WWWWY					

- ☞
- Instruction Bit 31 is always the *sign bit* (highest bit to sign extend in immediate)
  - B-type has a *13-bit* immediate encoding b/c of implicit zero in Imm bit 0.

☞

Between S and B, only two bits change meaning:  
 Instruction Bit 7 → S: Imm Bit 0; B: Imm Bit 11  
 Instruction Bit 31 → S: Imm Bit 11; B: Imm Bit 12

RISC-V immediate bit encoding is optimized to reduce hardware cost.

# All six RV32 B-Format Instructions

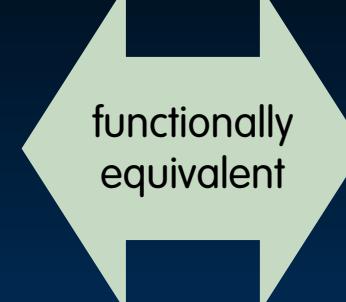
			funct3		opcode	
imm[12 10:5]	rs2	rs1	000	imm[4:1 11]	1100011	beq
imm[12 10:5]	rs2	rs1	001	imm[4:1 11]	1100011	bne
imm[12 10:5]	rs2	rs1	100	imm[4:1 11]	1100011	blt
imm[12 10:5]	rs2	rs1	101	imm[4:1 11]	1100011	bge
imm[12 10:5]	rs2	rs1	110	imm[4:1 11]	1100011	bltu
imm[12 10:5]	rs2	rs1	111	imm[4:1 11]	1100011	bgeu

# Conditionally Branching Even Further

- Conditional Branches are typically used for if-else statements, for/while loops.
  - Usually pretty small (<50 lines of C code) *no more than 50 lines*
- B-Format has limited range:
  - $\pm 2^{10}$  32-bit instructions from current instruction
- What if destination is further away?
  - Enter the *unconditional jump!*

**beq** x10, x0, far  
# next instr

if equal



if not equal  
next:  
**bne** x10, x0, next  
**j** far  
# next instr

Garcia, Yan

jump

# J-Format Layout

- Updating the Program Counter (PC)
- B-Format Layout
- J-Format Layout
- U-Format: Long Immediates
- `jalr`: I-Format

# The jal instruction and j pseudoinstruction

- **jal does a Jump And Link (i.e., “link and jump”):** Link and Jump  
 $\text{jal } \underline{\text{rd}}, \underline{\text{Label}}$ 
  - Jump to Label
  - Write address of the *following* instruction to rd.
- **Two use cases:**
  1. Call a function (and simultaneously save return address in named register)
  2. Unconditional branch
    - **j pseudoinstruction:** jump and discard return address

$\text{jal } \underline{\text{ra}}, \underline{\text{FuncLabel}}$

$\text{j } \underline{\text{Label}}$



$\text{jal } \underline{x0}, \underline{\text{Label}}$

*Jump discards the return address*

# The jal instruction and j pseudoinstruction

- jal does a Jump And Link (i.e., “link and jump”):

**jal rd, Label**

- Jump to **Label** (i.e., add relative offset to PC)
- Write address of the *following* instruction to **rd**.

$$\left. \begin{array}{l} \text{PC} = \text{PC} + \text{offset} \\ \text{rd} = \text{PC} + 4 \end{array} \right\}$$

- Two use cases:

1. Call a function (and simultaneously save return address in named register)

**jal ra, FuncLabel**

2. Unconditional branch

- j pseudoinstruction: jump and discard return address
- Often used to conditionally branch further than B-Types

**j Label**



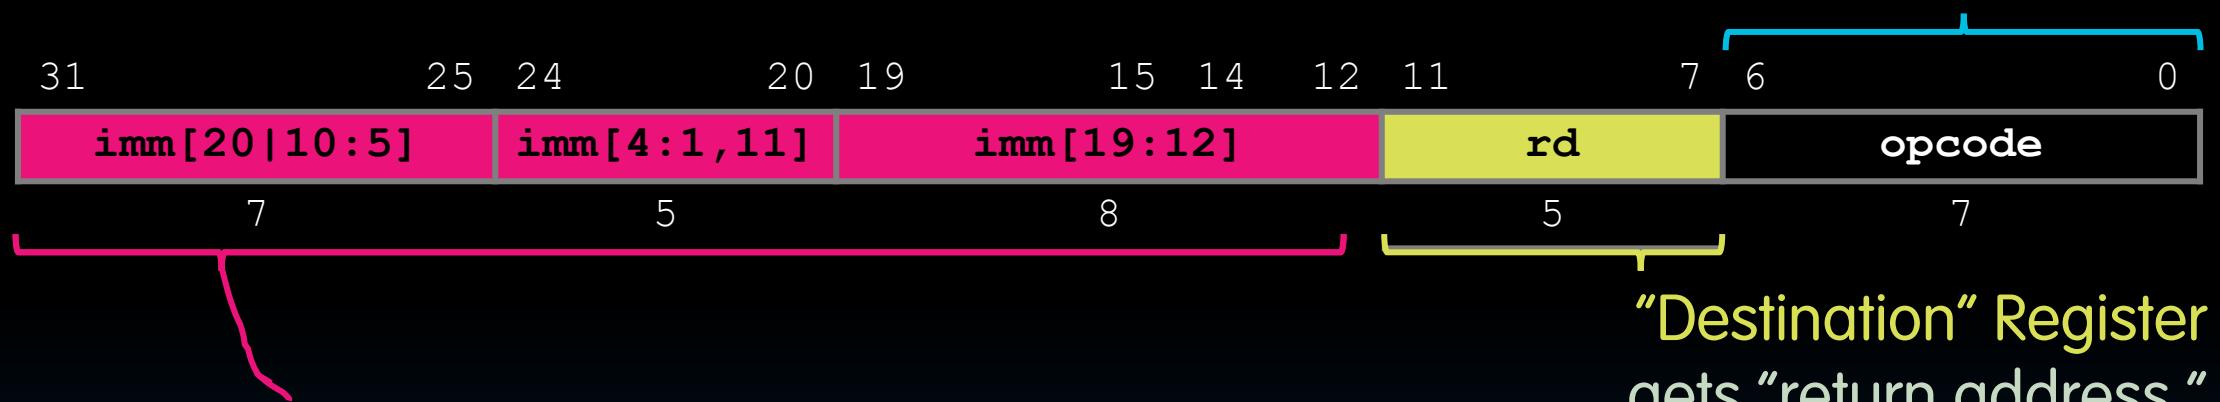
**jal x0, Label**

# J-Format Instruction Layout

- J-Format (textbook: UJ-Type) used only for jal:

# jal rd, Label

opcode 1101111



**Immediate** represents relative offset in increments of **2 bytes**.

“Destination” Register  
gets “return address.”

- To compute new PC:  $PC = PC + \text{byte\_offset}$ .
  - 20 immediate bits imply  $\pm 2^{18}$  32-bit instructions reachable:
    - 1 bit: 2's complement (allow  $+$ / $-$  offset), 1 bit: half-word/16-b instruction
  - Immediate bit encoding optimized to reduce HW cost

What about jumping *further*?  
Next: How to load 32-bit  
immediate into a register!

# U-Format: Long Immediates

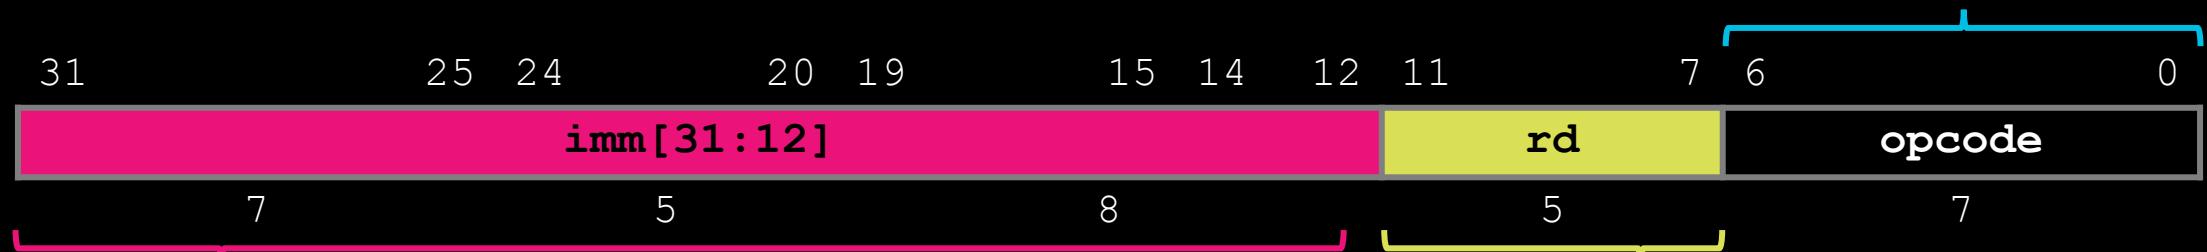
- Updating the Program Counter (PC)
- B-Format Layout
- J-Format Layout
- U-Format: Long Immediates
- `jalr`: I-Format

# U-Format Instruction Layout

- “Upper Immediate” instructions:

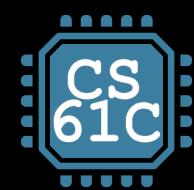
**opname      rd , immed**

U-Format opcodes:  
lui    0110111  
auipc 0010111



Immediate represents upper 20 bits of a 32-bit immediate  
operand **imm** = **immed** << 12.





# lui and addi creates Long Immediates

**lui rd, immed**

- The **lui instruction**, Load Upper Immediate:

- Write a 20-bit immediate value into the upper 20-bits of register rd.
- Clear the lower 12 bits.

$$\left. \begin{array}{l} \text{rd} = \text{immed} \ll 12 \\ \\ \text{li} = \text{lui} + \text{addi} \end{array} \right\}$$

- lui together with an addi (to set lower 12 bits) can create any 32-bit value in a register:**

```
lui x10, 0x87654      # x10 = 0x87654000  
addi x10, x10, 0x321  # x10 = 0x87654321
```

The li pseudoinstruction (Load immediate) resolves to **lui + addi** as needed, e.g., li x10, 0x87654321.

# lui Edge Case: addi Extends Sign

- How should we set the immediate 0xB0BACAFE?



- Unfortunately:

`lui x10, 0xB0BAC # x10 = 0xB0BAC000`

`addi x10, x10, 0xAF # x10 = 0xB0BABAFE`

- Recall: `addi` sign-extends the 12-bit immediate.
- If “sign bit” set, subtracts 1 from the upper 20-bits!

$$\begin{array}{r} 0xB0BAC000 \\ + 0xFFFFFAFE \\ \hline \end{array}$$

Diagram illustrating the calculation:

$$\begin{array}{r} 0xB0BAC \quad 0x000 \\ + \cancel{0xFFFF} - 1 \quad + 0xAF \\ \hline 0xB0BA\cancel{B} \quad 0xAF \\ \text{concat} \\ \hline 0xB0BABA\cancel{FE} \end{array}$$

0xB0BABA~~FE~~

# lui Edge Case: addi Extends Sign

- How should we set the immediate 0xB0BACAFE?



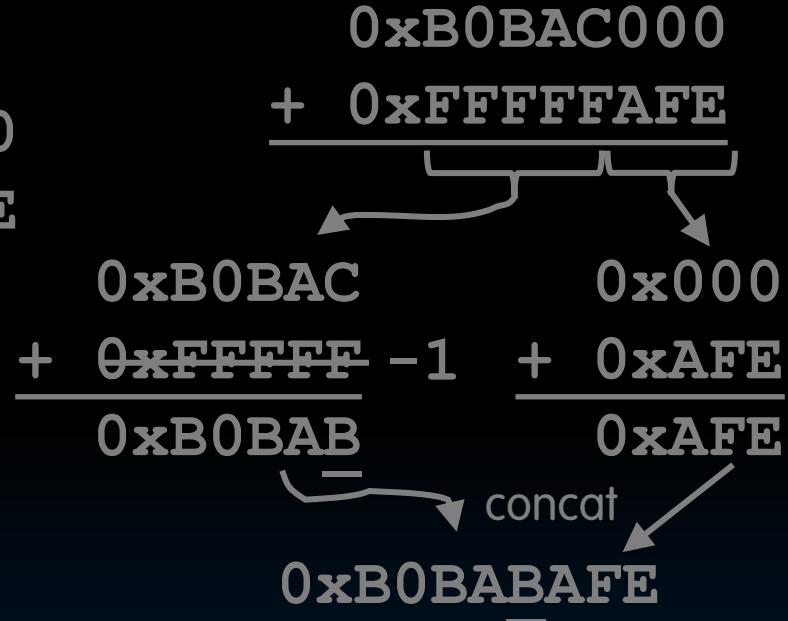
- Unfortunately:

```
lui x10, 0xB0BAC          # x10 = 0xB0BAC000
```

```
addi x10, x10, 0xAF         # x10 = 0xB0BABAAFE
```

- Recall: `addi` sign-extends the 12-bit immediate.
- If “sign bit” set, subtracts 1 from the upper 20-bits!

*lower down a bit*



✓ Solution: If 12-bit immediate is negative, add 1 to the upper 20-bit load.

```
lui x10, 0xB0BAD          # x10 = 0xB0BAD000
```

```
addi x10, x10, 0xAF        # x10 = 0xB0BACAAFE
```

Use pseudoinstructions!  
`li` automatically handles this edge case.

# auipc loads the PC into the Register File

auipc rd, immed

- The **auipc instruction**

Adds an Upper Immediate to the PC.

- Example:

auipc x5, 0xABCD # x5 = PC + 0xABCD

- In Practice:

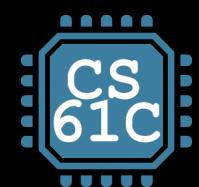
Label: auipc x5, 0 # puts address of Label in x5

- Loads the PC into a register accessible by other instructions.

auipc is most often used together with jalr to do PC-relative addressing with super large offsets.

# jalr: I-Format

- Updating the Program Counter (PC)
- B-Format Layout
- J-Format Layout
- U-Format: Long Immediates
- **jalr: I-Format**



# The jalr instruction and jr pseudoinstruction

- **jalr does a Jump And Link Register:**

jalr rd, rs1, imm

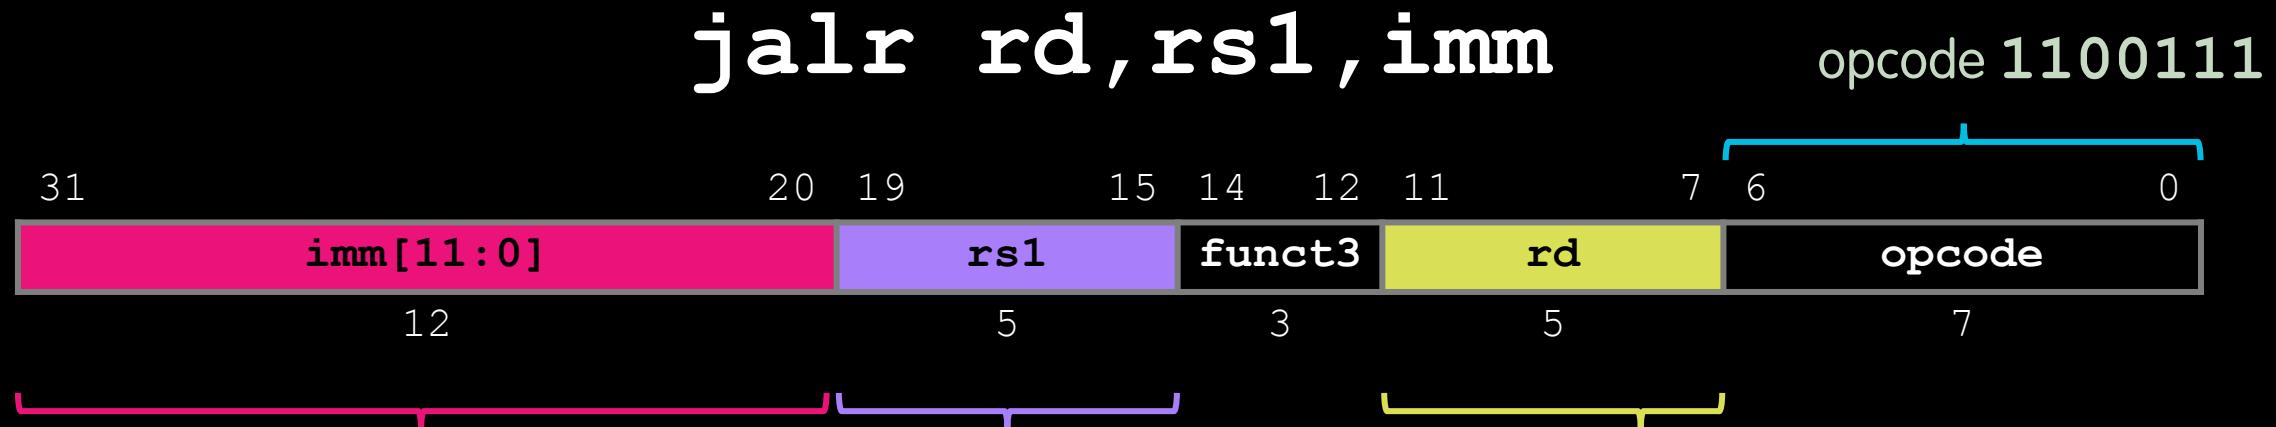
- Jump to **rs1 + imm.**
  - Write address of the *following* instruction to **rd.**
- PC = **rs1 + imm**  
**rd** = PC + 4
- **Use cases (so far):**

1. Return to caller

**jr ra**  **jalr x0, ra, 0**

# I-Format Instruction Layout: jalr

- jalr uses I-Format:



Immediate

**imm** and **rs1** are added together to update PC.

$$PC = rs1 + imm$$

⚠ Note I-Type!

Unlike B-type, J-type, **jalr** will not multiply **imm** by 2.

◦ Immediate offset therefore must be written in *units of bytes*.

"Destination" Register  
gets "return address."

$$rd = PC + 4$$

# The jalr instruction and jr pseudoinstruction

- **jalr** does a Jump And Link Register:

**jalr rd, rs1, imm**

- Jump to **rs1 + imm**.
  - Write address of the *following* instruction to **rd**.
- }  $\begin{array}{l} \text{PC} = \text{rs1} + \text{imm} \\ \text{rd} = \text{PC} + 4 \end{array}$

- Use cases:

1. Return to caller

**jr ra**  $\leftrightarrow$  **jalr x0, ra, 0**

Unlike **jal** (relative to PC), **jalr** addresses are relative to **rs1**, which is modifiable by arithmetic instructions. We can do bigger jumps!

2. Call a function (and save return address) at any 32-bit *absolute address*.
3. Jump *PC-relative* with a *32-bit offset*.

}  $\begin{array}{l} \text{lui x1 <hi20bits>} \\ \text{jalr ra, x1, <lo12bits>} \\ \text{auipc x1 <hi20bits>} \\ \text{jalr x0, x1, <lo12bits>} \end{array}$

# “And in Conclusion...”

- We've covered (almost) the entire RV32 ISA!
- Practice assembling and disassembling!

imm[31:12]		rd	0110111	LUI	0000000	shamt	rs1	001	rd	0010011	SLLI
imm[31:12]		rd	0010111	AUIPC	0000000	shamt	rs1	101	rd	0010011	SRLI
imm[20:10:1 11 19:12]		rd	1101111	JAL	0100000	shamt	rs1	101	rd	0010011	SRAI
imm[11:0]	rs1	000	rd	JALR	0000000	rs2	rs1	000	rd	0110011	ADD
imm[12 10:5]	rs2	rs1	000	BEQ	0100000	rs2	rs1	000	rd	0110011	SUB
imm[12 10:5]	rs2	rs1	001	BNE	0000000	rs2	rs1	001	rd	0110011	SLL
imm[12 10:5]	rs2	rs1	100	BLT	0000000	rs2	rs1	010	rd	0110011	SLT
imm[12 10:5]	rs2	rs1	101	BGE	0000000	rs2	rs1	011	rd	0110011	SLTU
imm[12 10:5]	rs2	rs1	110	BLTU	0000000	rs2	rs1	100	rd	0110011	XOR
imm[12 10:5]	rs2	rs1	111	BGEU	0000000	rs2	rs1	101	rd	0110011	SRL
imm[11:0]		rs1	000	LB	0100000	rs2	rs1	101	rd	0110011	SRA
imm[11:0]		rs1	001	LH	0000000	rs2	rs1	110	rd	0110011	OR
imm[11:0]		rs1	010	LW	0000000	rs2	rs1	111	rd	0110011	AND
imm[11:0]		rs1	100	LBU	0000000	pred	succ	000	00000	0001111	FENCE
imm[11:0]		rs1	101	LHU	0000	0000	0000	001	00000	0001111	FENCE.I
imm[11:5]	rs2	rs1	000	SB	0000000000000000			00000	000	00000	ECALL
imm[11:5]	rs2	rs1	001	SH	00000000000001			00000	000	00000	EBREAK
imm[11:5]	rs2	rs1	010	SW	csr	rs1	001	rd	1110011	CSRWR	
imm[11:0]		rs1	000	ADDI	csr	rs1	010	rd	1110011	CSRRS	
imm[11:0]		rs1	010	SLTI	csr	rs1	011	rd	1110011	CSRRC	
imm[11:0]		rs1	011	SLTIU	csr	zimm	101	rd	1110011	CSRRWI	
imm[11:0]		rs1	100	XORI	csr	zimm	110	rd	1110011	CSRRSI	
imm[11:0]		rs1	110	ORI	csr	zimm	111	rd	1110011	CSRRCI	
imm[11:0]		rs1	111	ANDI							

LUI	0000000	shamt	rs1	001	rd	0010011	SLLI
AUIPC	0000000	shamt	rs1	101	rd	0010011	SRLI
JAL	0100000	shamt	rs1	101	rd	0010011	SRAI
JALR	0000000	rs2	rs1	000	rd	0110011	ADD
BEQ	0100000	rs2	rs1	000	rd	0110011	SUB
BNE	0000000	rs2	rs1	001	rd	0110011	SLL
BLT	0000000	rs2	rs1	010	rd	0110011	SLT
BGE	0000000	rs2	rs1	011	rd	0110011	SLTU
BLTU	0000000	rs2	rs1	100	rd	0110011	XOR
BGEU	0000000	rs2	rs1	101	rd	0110011	SRL
LB	0100000	rs2	rs1	101	rd	0110011	SRA
LH	0000000	rs2	rs1	110	rd	0110011	OR
LW	0000000	rs2	rs1	111	rd	0110011	AND
LBU	0000000	pred	succ	000	00000	0001111	FENCE
LHU	0000	0000	0000	001	00000	0001111	FENCE.I
SB	0000000000000000			00000	000	00000	ECALL
SH	00000000000001			00000	000	00000	EBREAK
SW	csr	rs1	001	rd	1110011	CSRWR	
ADDI	csr	rs1	010	rd	1110011	CSRRS	
SLTI	csr	rs1	011	rd	1110011	CSRRC	
SLTIU	csr	zimm	101	rd	1110011	CSRRWI	
XORI	csr	zimm	110	rd	1110011	CSRRSI	
ORI	csr	zimm	111	rd	1110011	CSRRCI	
ANDI							

Not in CS61C

<https://riscv.org/technical/specifications/>

Garcia, Yan

